

Distributed Systems: Distributed File Systems – Part 2

Dr Soumyabrata DEV
<https://soumyabrata.dev/>

School of Computer Science and Informatics
University College Dublin
Ireland



Distributed Systems: Andrew File System

Introduction

- Developed at Carnegie Mellon University as part of an IBM-CMU collaboration.
- The key objective was to provide **scalability and security** to support file sharing between up to 5000 concurrent users.
- 2 key design features:
 - Whole-file serving: Entire files and directories are transmitted to client computers
 - Whole-file caching: Once a copy or chunk of a file has been transferred it is stored in a permanent local cache.

Typical Scenario

- When the client issues an open command for a shared file it checks the cache.
 - A copy of the file is downloaded to the cache if it is not present.
- The copy is then opened and the resulting UNIX file descriptor is returned.
- Subsequent read, write and other operations are applied to the local copy.
- While the client issues a close command, the local copy is checked.
 - Changes cause the file to be transmitted back to the server.
 - The client copy is maintained.


Why go for upload/download?

- Based on observations about typical workloads in academic and other environments:
 - Files are small (less than 10k in size)
 - Reads are about 6 times more common than writes.
 - Sequential access is common, random access is rare.
 - Most files are accessed by only one user
 - Most shared files are modified by one user
 - Recently used files will most probably be used again
- Today, file sizes two magnitudes larger – but then so are network transmission speeds!

Design Issues


- Andrew employs a large local cache (say 100Mb):
 - It is used to store local copies of the files that the user most commonly accesses.
 - This includes both shared code libraries and that user's personal files.
 - This means that, after the initial access, most shared files can be accessed at the same speed as local files.
- Unfortunately, this does not work for files such as **databases** which are frequently accessed and updated by multiple users.

AFS Design




AFS File Space

Local



Shared



AFS Implementation

● Vice:

- Implements a Flat File Service.
- Groups files into volumes (e.g. system binaries, documentation, a users home directory) for ease of location and movement.
- Each file is identified by a unique 96-bit file identifier (fid).

Volume number

File Handle

Uniquifier

● Venus:

- Updates the local cache, downloading files as necessary.
- Manages shared directory structure.
- Resolves directories issued by clients to fids using a step-by-step lookup.
- Maintains cache coherence via callback promises

AFS Cache Coherence

- Vice supplies a token, known as a callback promise, for each copy of a file that is transmitted to a Venus process.
 - It represents a promise that Vice will contact that Venus process should any other client modify that file.
 - Callback promises are stored locally, initially with a valid state.
- When a file is updated, the Vice server performs a callback.
 - It contacts each client that was issued a token and informs it that the file was updated.
 - This causes the callback promise to be changed to cancelled.
- When Venus receives an open request, it checks the callback promise.
 - If it is set to cancelled, then a fresh copy of the file is retrieved.

AFS Cache Coherence

- When a workstation is restarted, Venus imposes an additional check on all files in the cache:
- Upon first access of a file, it submits a cache validation request to the server containing the files' last modification timestamp.
- If the timestamp is correct, then the callback promise is validated.
- This validation also occurs if time T has elapsed since the communication with the server.

AFS Update Semantics

- Attempts to approximate to one-copy semantics for UNIX file access.
- Full one-copy semantics requires that a write results in all cached copies of a file being updated before another operation is applied to that file.
- For a client C operating on a file F the following guarantees are maintained:
 - After successful open: $\text{latest}(F, S)$
 - Current value of file F at client C is the same as the value at server S
 - After a failed open/close: $\text{failure}(S)$
 - Open/close not performed at server
 - After successful close: $\text{updated}(F, S)$
 - Client's value of F has been successfully propagated to S

Other Aspects

- Threads

- Venus and Vice are multi-threaded

- Read-only replicas

- Multiple read-only copies of volumes containing frequently used (read) files (e.g. /bin) are maintained on different servers.

- Bulk transfers

- The 64k limit on file transfer size reduces network overheads

AFS Review

	FFS	AFS
Objective	Reference Architecture	Scalable and Secure
Security	Directory Service/RPC	RPC-based
	Encryption	Private key algorithm
Naming	Location Transparency	Location Indepence
File Access	Remote Access Model + Cache	Upload / Download + Cache
Cache-Update Policy	Write-Through	Callback Promise (token)
	N/A	Write-on-close
Consistency	N/A	Timestamp-based
Service Type	Stateless	Stateful
Replication	N/A	Read-only

Peer to Peer Networks

What is P2P?

- Peer-to-Peer (P2P) Systems:
 - A set of networked devices that have **equal responsibility** and which **share resources and workload** as necessary.
 - Each device is both a **client** and a **server**.
 - There is no (or little) centralised control.
- Examples of Peer-to-Peer Systems:
 - File Sharing Applications e.g. Gnutella, BitTorrent, Kazaa, ...
 - Instant Messengers and Telephony e.g. Skype, Yahoo, MSN
 - Data Processing Services e.g. [Seti@Home](#)
- Peer-to-Peer Systems **most effective** when used to store very large collections of immutable data
- Their design **diminishes their effectiveness** for applications that store and update mutable data

Characteristics of P2P Systems

- Aim of P2P Systems:
 - To deliver a service that is fully decentralised and self-organising, dynamically balancing the storage and processing loads between all the participating computers as computers join and leave the service.
- Key Characteristics:
 - Every user contributes resources to the service.
 - While resources differ, all nodes in a P2P system have the same functional capabilities and responsibilities.
 - Their correct operation does not depend on the existence of any centrally-administered systems.
 - They can be designed to offer a limited degree of anonymity to the providers and users of resources.
- Key Problem:
 - How to deal with the placement and accessing of shared resources in a manner that balances the workload and ensures availability, but does not add undue overheads?

Placement of Resources

- Deciding at which node(s) a resource should be located
 - Can be managed by the P2P application, or left to nature (the users)
- Examples:
 - Andrew File System
 - Consists of a set of file servers that decide how many copies of a file are needed to meet demand and where those copies should be located.
 - Users are not aware of the existence of multiple copies of the file (location independence).
 - Napster, Gnutella
 - Each user runs the application (peer) and decides which file(s) will be located at their node.
 - Copies are created in an ad-hoc way - each user downloads the files that they want (and then shares them).
 - Users are aware of the number of copies that they have access to.



Accessing (Locating) Resources

- In (centralised) managed environments: searching is not a problem:
 - We know where all the copies of a file are
 - We can develop algorithms that employ knowledge of the placement scheme that has been used.
- In P2P environments: resources are distributed over a diverse set of nodes.
 - How do we find the resource we are looking for?
 - We need to perform some kind of search!
- In such ad-hoc environments we have a problem:
 - We do not know where the resource is located
 - We may not even know what nodes exist!
 - While a number techniques have been developed to handle this, locating resources in a P2P system is still a hot topic of research...

Types of P2P Systems

- Three main types of P2P system:

- Centralized systems: peer connects to server which coordinates and manages communication

- e.g. SETI@home

centralised discovery
centralised communication

- Decentralized systems: peers run independently with no central services.

- Discovery is decentralized and communication takes place between the peers. e.g. Gnutella, Pastry

decentralised discovery
decentralised communication

- Hybrid systems: peers connect to a server to discover other peers

- peers manage the communication themselves (e.g. Napster).

- This is also called 'Brokered P2P'.

centralised discovery
decentralised communication



Classifying P2P File Systems

- Generation 1 (e.g. Napster)
 - Centralised file list (a list of IDs that map onto peers)
 - Files are located by searching the file list
 - Once located, the file is downloaded directly from the source peer.
 - Suffers from many drawbacks (not least the legal one) and is not widely used anymore.
- Generation 2 (e.g. Gnutella, Freenet, Kazaa & BitTorrent)
 - File list is distributed between the various peers using a Distributed Hash Table (DHT).
 - Network overhead is reduced through the election of certain nodes as “ultrapeers” to reduce network overheads.
 - This is the type of system that is currently in use
- Generation 3: P2P middleware
 - Third generation is characterized by middleware layers for the application independent management of distributed resources on a global scale
 - Still in research phase. Several research teams have developed and evaluated middleware platforms and deployed them in a range of application services
 - Best Examples: Pastry (2001), Tapestry (2004), CAN (2001), Chord (2001) and Kademlia (2002)

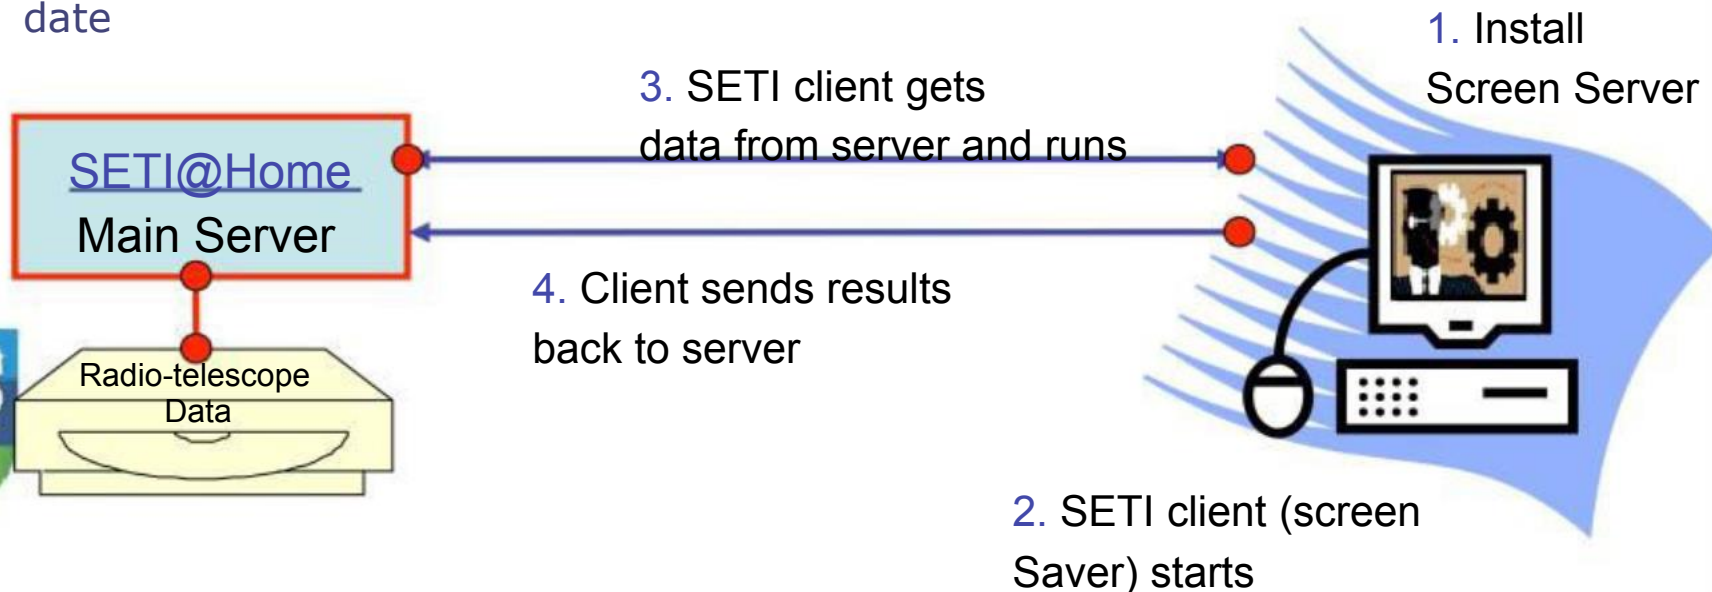


First Generation P2P Systems

Distributed Computation

Example 1: SETI@HOME (Centralized P2P)

- Launched In 1996
- Scientific experiment - uses Internet-connected computers in the Search for Extraterrestrial Intelligence (SETI)
- Distributes a screen saver-based application to users
- Applies signal analysis algorithms different data sets to process radio-telescope data.
- Has more than 3 million users - used over a million years of CPU time to date



First Generation P2P Systems

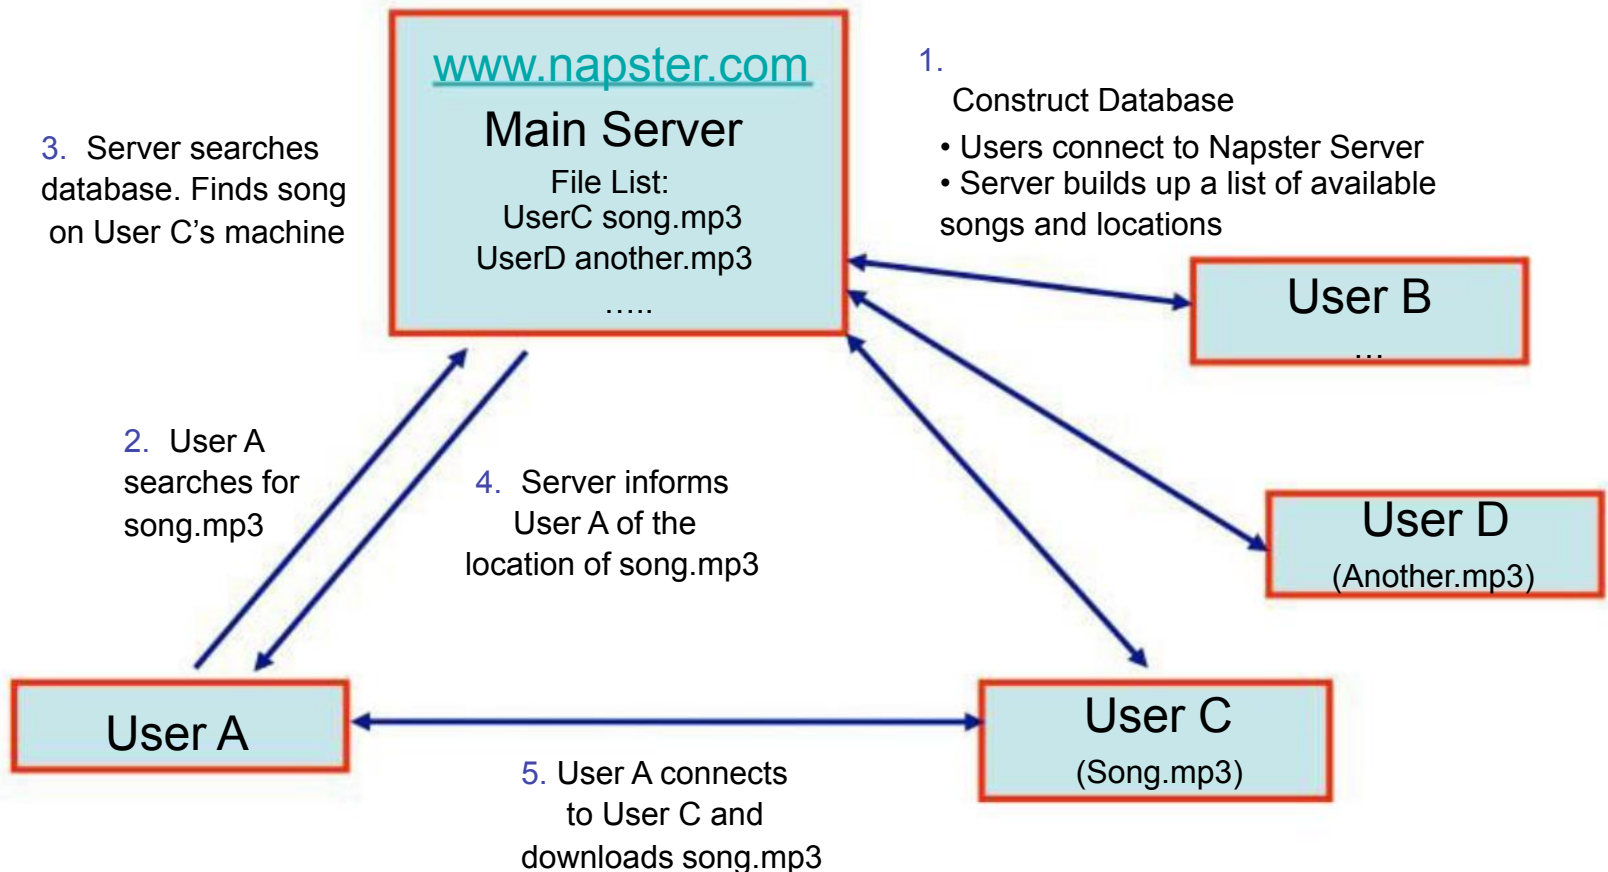
Example 2: Napster (Hybrid P2P)

- Centralised P2P file distribution system
- Timeline:
 - Launched in May 1999, by Shawn Fanning (19) and Sean Parker (20) while attending Northeastern University, Boston
 - Allowed Users to share MP3 Files - compression format, good quality but 1/12th original size
 - April 2000 - Metallica starts law suit - Huge and long court case
 - November 2000 - Napster has 38 Million members
 - July 2001 - Napster ordered offline, June 2002 bankrupt



First Generation P2P Systems

Example 2: Napster (Hybrid P2P)



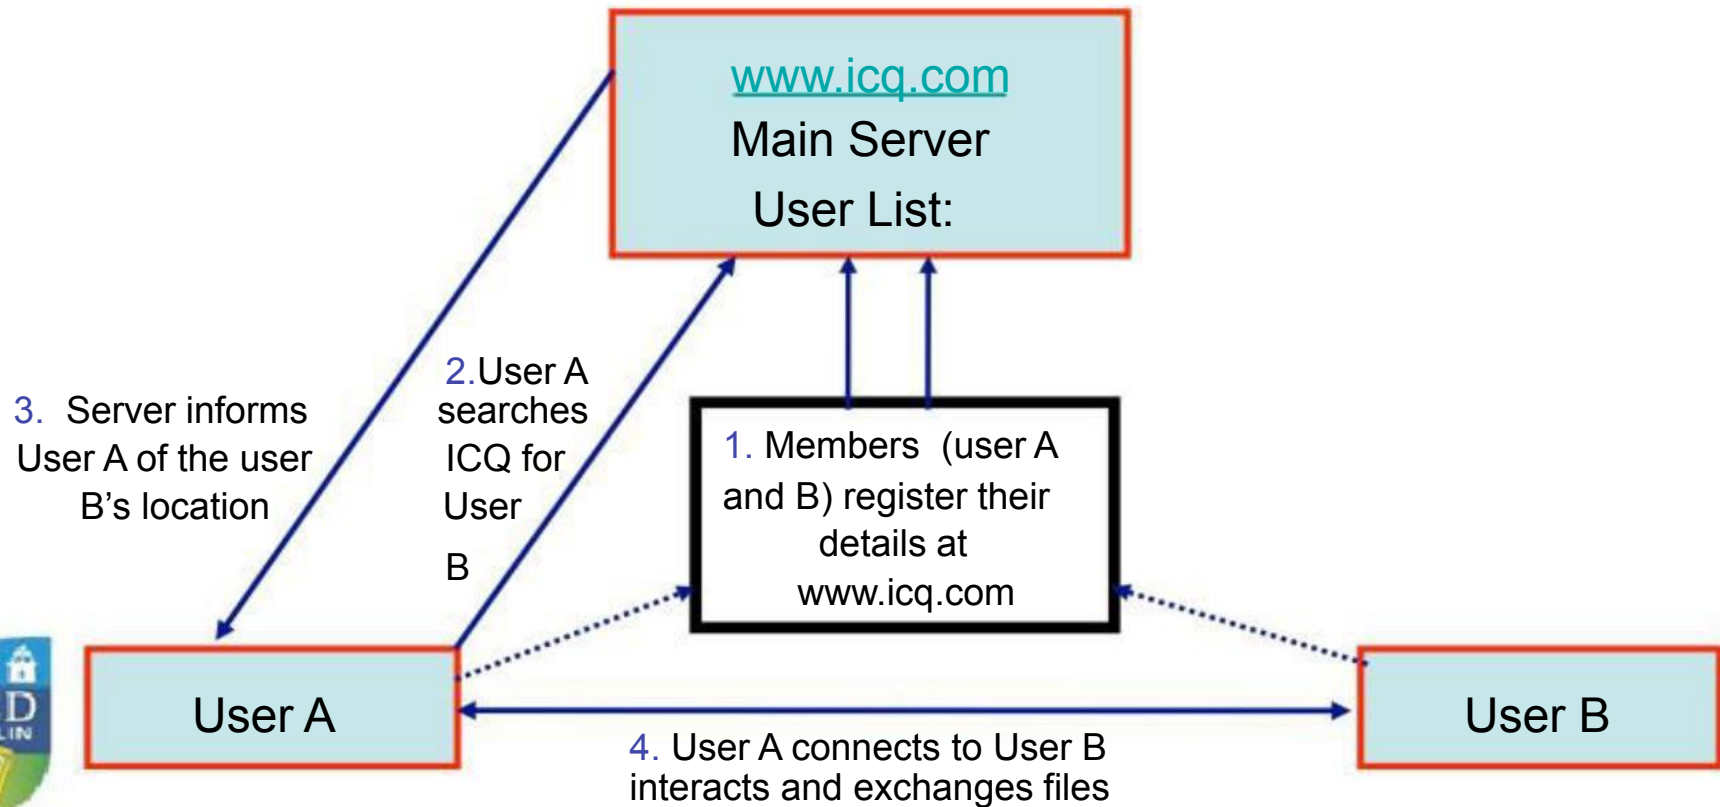
Instance Messaging Applications

- ICQ: pioneer of Instant Messaging (IM) - Released in November 1996
- Type of Hybrid P2P system
 - centralized discovery, decentralized communication
- uses central server to monitor which users are online - notifies users when their friends come online
- all other communication is conducted between peers directly
- Users can send messages to their friends (instant messaging)
- Also allows users to exchange files
- Other Examples: MSN, QQ, gmail chat, skype chat, Yahoo! Messenger



Instance Messaging Applications


ICQ General Architecture



Second Generation P2P Systems

Example 1: Gnutella - First Decentralised P2P System

- Two types of Node:
 - Peers - Normal nodes
 - Ultrapeers - Ad-hoc servers.
- Peers get promoted to ultrapeers dynamically, based on their bandwidth.
- Discovery based on broadcasts (also known as 'query flooding')
 - Nodes send a search message to its ultrapeer
 - Ultrapeers forward message to all other connected peers (in Gnutella 0.4, typically 5)
 - Repeat recursively until resource is found OR all nodes are searched OR message reaches maximum hops (typically 7)
- Searches generate a lot of traffic: $O(n)$ complexity
 - Not scalable!



The Gnutella Protocol

- Gnutella 0.4 is the original protocol
- Operated a purely query flooding-based search protocol
- Consists of five packet types:
 - ping: discover hosts on network
 - pong: reply to ping
 - query: search for a file
 - query hit: reply to query
 - push: download request for firewalled servers
- These packet types deal with health checks between peers and searching the network for files.
- File downloads are handled by HTTP.
- Many extensions now exist - deal with issues such as intelligent query routing



Summary so far...

- P2P is still an emerging paradigm.
 - It promotes a view of distributed systems as a open systems that are comprised of a heterogeneous collection of nodes that have equal responsibilities and capabilities.
 - From the users perspective, a P2P system is a shared resource space that they have global access to
- The key problem to be addressed when building a P2P system is how to locate the resources that you need
 - First generation systems addressed this problem through a centralised server that contained a global resource list (not scalable, legal issues)
 - Second generation systems decentralised the resource list to improve robustness (and to sidestep the legal issues ;-)
 - Third generation systems have focused on making file sharing an anonymous activity and on techniques for improving resource search speed



Peer-to-Peer Middleware Systems

P2P Middleware Systems

- With the growing popularity of P2P systems, P2P researchers worked to develop generalized versions of the first P2P system
 - that utilise existing routing, naming, data replication and security techniques in new ways to provide a reliable resource sharing layer over an unreliable and untrusted collection of computers and networks
- Aims:
 - To be application independent
 - scalability: to share resources, storage and data present in computers at the edges of the internet on a global scale
- They provide reference architecture that allowed researchers to focus on specific P2P issues.
- These systems are commonly known as P2P Middleware Systems.



P2P Middleware Systems

- P2P Middleware systems are based on 2nd generation architectures:
 - Every node in the system is a peer.
 - Some special infrastructure nodes included to improve network performance.
 - Peer nodes connect to infrastructure nodes (which may also be peers), and route searches via these infrastructure nodes.
- Perhaps the most central issue in P2P research is the issue of resource discovery (searching).
 - algorithms used to carry out the search within P2P middleware systems are known as routing overlay algorithms



Routing Overlays - basic idea

- P2P routing overlays are used to locate objects.
 - The middleware is a layer that is responsible for routing requests from any node to the node that holds the target resource.
 - However, P2P systems allow for the migration of resources between nodes.
 - As a result, the term overlay is used to clarify that it is an application layer routing mechanism (it is not a network layer mechanism like IP routing).
 - The first system to employ routing overlays was the Plaxton Mesh. Others to follow include Pastry and Tapestry.
- Routing overlays ensure that any node can access any resource by routing each request through a sequence of nodes.
 - It exploits knowledge at each of them in order to locate the resource.
 - Because P2P systems can store replicas, the routing algorithm must maintain the location of all copies.
 - It then delivers requests to the nearest “live” node



Routing Overlay - basic idea

- The main task of a routing overlay is:
 - Clients wishing to invoke an operation on a resource submit a request, which includes a GUID (unique identifier) for the resource, to the routing overlay algorithm.
 - Globally Unique IDentifiers (GUID)
 - It directs the request to a node on which a replica (or the original) of the resource resides.
- In addition the routing overlay is responsible for:
 - Creating GUID's for all new resources
 - Announcing the existence of the new resources to ensure that the resource is reachable by all clients.
 - Handling the removal of references to shared resources that are no longer in the system
 - managing the addition/removal of nodes to/from the system.



Routing Overlays - Resource GUIDS

- In a P2P System, resources are identified by a globally unique identifier (GUID) that is location independent.
 - For object-based resources the GUID is calculated using a hash function, such as SHA-1.
 - This hash is derived from some/all of the resources state. This means that clients receiving the resource can check the validity of the hash (self certifying aspect)
 - Uniqueness is verified by searching for another object with the same GUID.
- GUID's are used to determine the placement of objects, and to retrieve them
- Because of this, routing overlay routing systems are sometimes constructed using Distributed Hash Tables (DHT -this is reflected in the simple API or operations that is used to access and add resources)
- The DHT layer takes responsibility for:
 - choosing a location for each resource
 - storing it (with replicas to ensure availability)
 - providing access to it via a get() operation.



Routing Overlays - Distributed Hash Tables (DHT)

- DHTs form an infrastructure that can be used to build peer-to-peer networks.
- Question: GUID's are used to determine the placement of objects, and to retrieve them - Within the DHT model, how does this happen.
- Answer: A data item with GUID X is stored at the node whose GUID is numerically closest to X , and at the r hosts with GUIDs numerically closest to X .
 - where r is a replication factor chosen to ensure a very high probability of availability



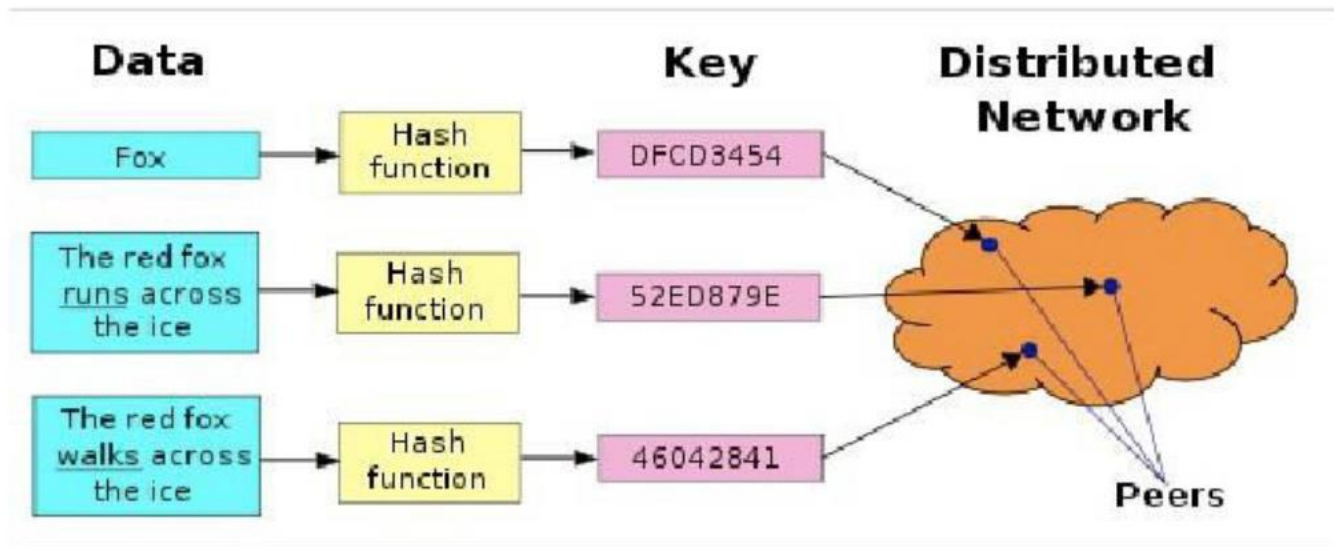
Distributed Hash Table Operations

- put(GUID, data)
 - The data is stored in replicas at all nodes responsible for the object identified by GUID
- remove(GUID)
 - Deletes all references to GUID and the associated data
- data = get(GUID)
 - The data associated with GUID is retrieved from one of the nodes responsible for it.



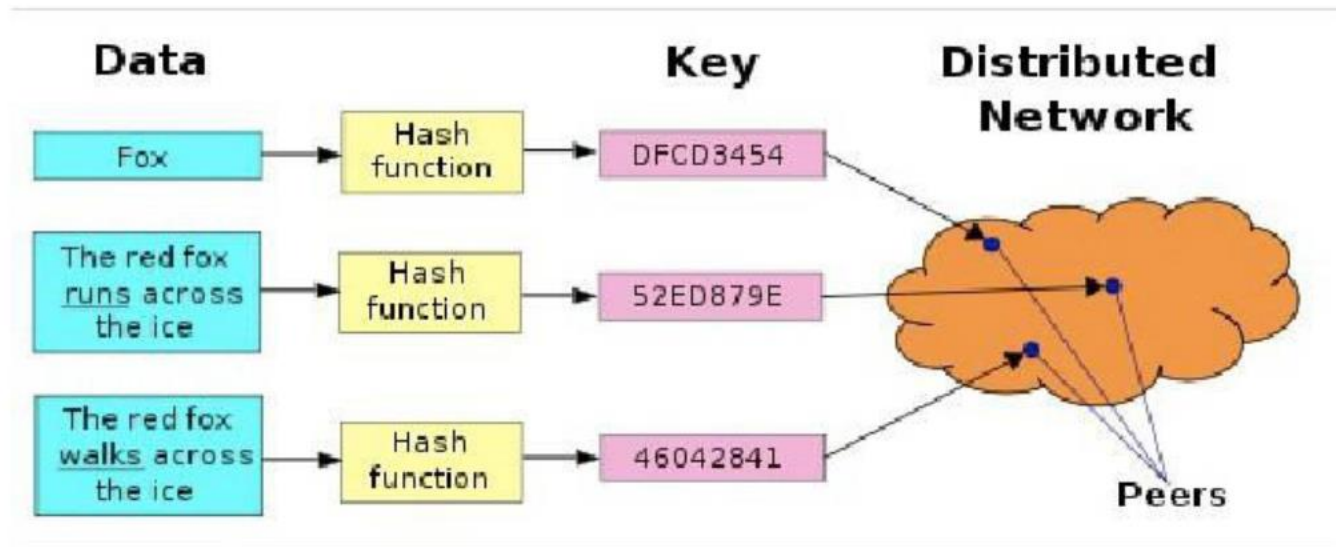
Distributed Hash Tables (DHT's)

- Networks that use DHTs include:
 - BitTorrent
- provide a lookup service similar to a hash table
 - (key, value) pairs are stored in the DHT
 - any participating node can efficiently retrieve the value (data) associated with a given key



Distributed Hash Tables (DHT's)

- Responsibility for maintaining the mapping from keys to values (data) is distributed among the nodes
 - can scale to extremely large numbers of nodes
 - handle continual node arrivals, departures, and failures.



DHT's - Properties

- Decentralization: the nodes collectively form the system without any central coordination.
- Scalability: the system should function efficiently even with thousands or millions of nodes.
- Fault tolerance: the system should remain reliable even with nodes continuously joining, leaving, and failing
 - mainly through the replication of objects





Thank you

For general enquiries, contact:

Please contact the Head Teaching Assistant: Xingyu Pan (Star), Xingyu.Pan@ucdconnect.ie