

Dr SABER Takfarinas
takfarinas.saber@ucd.ie

COMP2009J
Computer Networks

Dynamic Routing



IP Routing

- There are two approaches for calculating the routing tables

1. Static Routing

- by hand


2. Dynamic Routing

- automatically calculated by a routing protocol



Autonomous Systems

- An **autonomous system** is a region of the Internet that is administered by a single entity
- E.g.,
 - Campus network
 - Backbone network
 - National Internet Service Provider
- Routing is done differently:
 - **intradomain routing**: within an autonomous system
 - **interdomain routing**: between autonomous system



Interdomain vs Intradomain Routing

• Intradomain Routing

- Routing within an AS
- Ignores the Internet outside the AS

- Protocols for Intradomain routing are also called **Interior Gateway Protocols** or **IGP's**

- Popular protocols: RIP (simple, old), OSPF (better)

• Interdomain Routing


- Routing between AS's
- Assumes that the Internet consists of a collection of interconnected AS's
- Normally, there is one dedicated router in each AS that handles interdomain traffic

- Protocols for interdomain routing are also called **Exterior Gateway Protocols** or **EGP's**

- Popular protocols: EGP, BGP (more recent)

Why Do We Need a Routing Algorithm?

- Need to make sure that every network is reachable!
- In a network, there might be different paths to go from a node to another
- In addition, links do not have equal costs
 - E.g., latency
 - We need to find the best path
- Static routing not ideal. Why?
 - Cost on links keeps changing
 - We need a dynamic routing



Requirements of a dynamic Routing?

1. Send and Receive reachability information about network to other routers
2. Calculate optimal routes using a shortest path algorithm
3. Advertise and react to topology changes



Routing Algorithms

- There are two main types of routing algorithms:

Distance Vector Routing

- Every node knows the distance (i.e., cost) to its directly-connected neighbours
- A node sends periodically a list of routing updates to its directly-connected neighbours
- If all nodes update their distances, the routing tables eventually converge

Exterior Gateway Protocol (EGP)

Link State Routing

- Each node knows the distance (i.e., cost) to its directly connected neighbours
- The distance information of every link is broadcasted to all nodes in the network
- Each node calculates the routing tables independently

Open Shortest Path First (OSPF)



Distance Vector Routing

- A.k.a., Bellman-Ford forwarding
- Idea:
 - Each router holds a Distance Vector (DV) for all available destinations
 - Each router shares information it has about the network with its neighbours
 - Repeated until all routers have all the information necessary to route to all routers (converge)



When to Update Tables in Distance Vector Routing


- 1. Periodic Updates:** every time period (e.g., every 90s)
- 2. Triggered Updates:** If a metric changes on a link, a router immediately sends out an update without waiting for the end of the update period
- 3. Full Routing Table Update:** send the entire routing table to the neighbours (not only entries which change).
- 4. Route invalidation timers:** Routing table entries are invalid if they are not refreshed for a certain number of updates (e.g., no updates received after 6 update periods)



Example

Initial State: every node knows the distance to his direct neighbours


	A	B	C	D	E
A	0	7	∞	∞	2
B	7	0	1	∞	3
C	∞	1	0	2	∞
D	∞	∞	2	0	4
E	2	3	∞	4	0



Example

B sends its distance vector to its neighbours: A, E, C

	A	B	C	D	E
A	0	7	8	∞	2
B	7	0	1	∞	3
C	∞	1	0	2	∞
D	∞	∞	2	0	4
E	2	3	∞	4	0



A receives B's distance vector

7	0	1	∞	3
---	---	---	----------	---

Adds the distance between A and B (i.e., 7)

14	7	7	∞	10
----	---	---	----------	----

Compares the results to its current distance vector

0	7	∞	∞	2
---	---	----------	----------	---

Update distances with shorter ones if they exist

0	7	8	∞	2
---	---	---	----------	---


A changed his distance vector \rightarrow A has to inform his neighbours



Example

B sends its distance vector to its neighbours: A, E, C

	A	B	C	D	E
A	0	7	8	∞	2
B	7	0	1	∞	3
C	∞	1	0	2	∞
D	∞	∞	2	0	4
E	2	3	4	4	0



E receives B's distance vector

7	0	1	∞	3
---	---	---	----------	---

Adds the distance between E and B (i.e., 3)

10	3	4	∞	6
----	---	---	----------	---

Compares the results to its current distance vector

2	3	∞	4	0
---	---	----------	---	---

Update distances with shorter ones if they exist

2	3	4	4	0
---	---	---	---	---

E changed his distance vector \rightarrow E has to inform his neighbours



Example

B sends its distance vector to its neighbours: A, E, C

	A	B	C	D	E
A	0	7	8	∞	2
B	7	0	1	∞	3
C	8	1	0	2	4
D	∞	∞	2	0	4
E	2	3	4	4	0

C receives B's distance vector

7	0	1	∞	3
---	---	---	----------	---

Adds the distance between C and B (i.e., 1)

8	1	2	∞	4
---	---	---	----------	---


Compares the results to its current distance vector

∞	1	0	2	∞
----------	---	---	---	----------

Update distances with shorter ones if they exist

8	1	0	2	4
---	---	---	---	---


C changed his distance vector \rightarrow C has to inform his neighbours



Example

A sends its distance vector to its neighbours: B, E

	A	B	C	D	E
A	0	7	8	∞	2
B	7	0	1	∞	3
C	8	1	0	2	4
D	∞	∞	2	0	4
E	2	3	4	4	0




We keep repeating this until convergence

	A	B	C	D	E
A	0	5	6	6	2
B	5	0	1	3	3
C	6	1	0	2	4
D	6	3	2	0	4
E	2	3	4	4	0

We repeat this processes until no change in the distance vector of all nodes

DONE



Link State Routing

- Each router shares information about its neighbours with the rest of the network
- Each router stores the complete topology of the network
- All routers on the network store the same information
- Each router uses an algorithm to solve the best path using:
 - Dijkstra's algorithm
- "Flood" link information throughout the network using Link State Packets (LSPs)
- Convergence occurs when all routers have received a LSP from each other router in the network

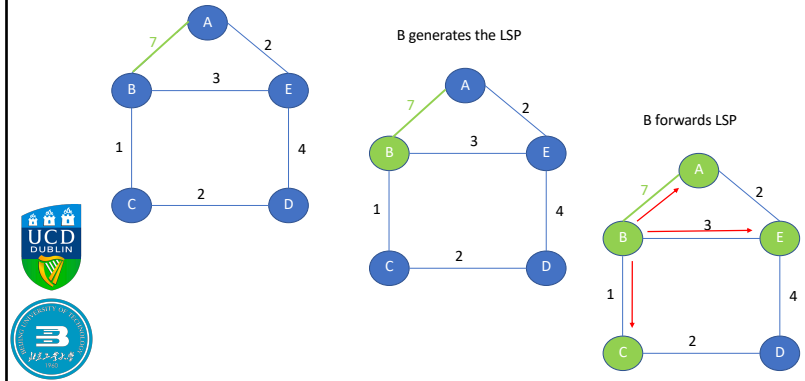


- Reliable flooding
 - Each router transmits a **Link State Packet (LSP)** on all links
 - A neighboring router forwards out all links except incoming
 - » Keep a copy locally; don't forward previously-seen LSPs
- Challenges
 - Packet loss
 - Out-of-order arrival
- Solutions
 - Acknowledgments and retransmissions
 - Sequence numbers
 - Time-to-live for each packet

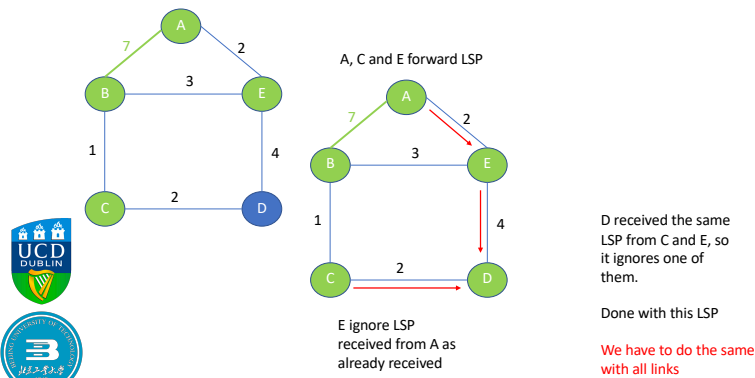
nks



Example of LSP Broadcasting



Example of LSP Broadcasting



When to Flood ?

- Periodically:
 - For example: every 30 minutes
- After a topology change
 - Link or node failure
 - Modification in link metric




After Sending all LSPs

- All nodes should have the same link-state database
- Each node calculates the shortest paths
 - Using Dijkstra Algorithm
- Forwards its packets on the shortest path



Dijkstra Algorithm



Link State Database
at node C

	B	E
A	7	2

	A	C	E
B	7	1	3

	B	D
C	1	2

	C	E
D	2	4


	A	B	D
E	2	3	4

Shortest Paths From C

Add distance label to each node

Take the node with the smallest distance, from reachable nodes

Explore the outgoing link



	B	E
A=∞	7	2

	A	C	E
B=∞	7	1	3

	B	D
C=0	1	2

	C	E
D=∞	2	4

	A	B	D
E=∞	2	3	4




Shortest Paths From C

Add distance label to each node

Take the node with the smallest distance, from reachable nodes

Explore the outgoing link



	B	E
A=∞	7	2

	A	C	E
B=1	7	1	3

	B	D
C=0	1	2

	C	E
D=2	2	4

	A	B	D
E=∞	2	3	4




Shortest Paths From C

Add distance label to each node

Take the node with the smallest distance, from reachable nodes

Explore the outgoing link



	B	E
A=8	7	2

	A	C	E
B=1	7	1	3

	B	D
C=0	1	2

	C	E
D=2	2	4

	A	B	D
E=4	2	3	4




Shortest Paths From C

Add distance label to each node

Take the node with the smallest distance, from reachable nodes

Explore the outgoing link



	B	E
A=8	7	2

	A	C	E
B=1	7	1	3

	B	D
C=0	1	2

	C	E
D=2	2	4

	A	B	D
E=4	2	3	4




Shortest Paths From C

Add distance label to each node

Take the node with the smallest distance, from reachable nodes

Explore the outgoing link



	B	E
A=6	7	2

	A	C	E
B=1	7	1	3

	B	D
C=0	1	2

	C	E
D=2	2	4

	A	B	D
E=4	2	3	4



No more nodes to visit.
DONE!