

Databases and Info Systems

Relational Databases

Dr. Seán Russell
`sean.russell@ucd.ie`,

School of Computer Science,
University College Dublin

February 19, 2020

Table of Contents

- 1 Relational Model
- 2 Structure
- 3 Integrity
- 4 Manipulating Relational Databases
- 5 Closure

Relational Model

- There are three aspects of the relational model that we must understand
 - How it is **structured**
 - How it maintains **integrity**
 - How it can be **manipulated**

Table of Contents

- 1 Relational Model
- 2 Structure
- 3 Integrity
- 4 Manipulating Relational Databases
- 5 Closure

Database Structure

- Data in the relational models is stored in a collection of **relations**
 - We will use the term **table** when discussing relations
- The relationships between tables is not explicitly stored

Database Structure Example

students

| <u>student_num</u> | name | major | year_of_entry |
|--------------------|--------------|-------|---------------|
| 17206777 | Sean Russell | SE | 2017 |
| 18205333 | David Lillis | IOT | 2016 |
| 16205777 | Brett Becker | EIE | 2016 |

modules

| <u>module_code</u> | title | teacher |
|--------------------|-----------------|--------------|
| COMP2013J | Database | Sean Russell |
| COMP2011J | OOP | Sean Russell |
| COMP2003J | Data Structures | David Lillis |

Structure Example

- The structure of a database is described in terms of **relations** and **attributes**
- We use this format to describe them
 - `relation_name (attribute_name1, a_name2, a_name3)`
- The relations from our example are
 - `students(student_num, name, major, year_of_entry)`
 - `modules(module_code, title, teacher)`
 - `results(student_num, grade, module_code)`

Connections

- We can see that data in some of these tables are related
- But this is not recorded in the structure of the database
- We have to create these connections when we are searching for data

Table of Contents

- 1 Relational Model
- 2 Structure
- 3 Integrity
 - Keys
- 4 Manipulating Relational Databases
- 5 Closure

Database Integrity

- Database integrity is the correctness of the data
- Integrity is ensured using a number of constraints (rules)
 - Domain Integrity Constraint
 - Entity Integrity Constraint
 - Referential Integrity Constraint

Domain Integrity Constraint

- The domain integrity constraint states that all columns in a relational database are in a defined **domain**
- This is similar to data types in programming
 - E.g. `student_num` can only allow integer values
 - E.g. `module_code` can only allow text values

Entity Integrity Constraint

- The entity integrity is concerned with the concept of primary keys
- The rule states that every table must have its own **primary key**
- The value of this attribute must be **unique** and **not null** for every row

Referential Integrity Constraint

- The referential integrity constraint is the concept of foreign keys
- The rule states that the foreign key value can be in two states.
 - The first state is that the **foreign key** value would refer to a **primary key** value of another table
 - The second state is that it can be **null**
- Being null could simply mean that there are no relationships, or that the relationship is unknown

Keys

- The entity and referential integrity constraints both referred to **keys**
- A key is an attribute or set of attributes that **uniquely identifies** rows in a table
 - Remember no duplicate rows are allowed

Primary Key

- Any key that uniquely identifies a relation is called a **candidate key**
- From the candidate keys a single **primary key** is chosen to identify the relation
- In many modern databases, this attribute is added to make sure each record is **unique**
 - This is why you have a student number
 - Many students may have the same name

Keys

- Primary keys are shown by underlining the attributes
- Multiple attributes can be used together as the primary key
- Our example redone would be:

Example

```
students( student_num, name, major, year_of_entry )  
modules( module_code, title, teacher )  
results( student_num, grade, module_code )
```


Combined Keys

- When using a combined primary key we can have duplicates of **part** of the key but not all of it
 - `results(student_num, grade, module_code)`
- In this example the same `student_num` can be repeated and so can the same `module_code`
- But you cannot have the same combination twice

Foreign Keys

- A foreign key is an attribute or set of attributes in a relation that is a key in another relation
- Foreign keys are used to link data together in different relations
- When using foreign keys, we combine data together only where the foreign key in our table matches the primary key in another table

Foreign Keys

Example

```
students( student_num, name, major, year_of_entry )  
modules( module_code, title, teacher )  
results( student_num, grade, module_code )
```

- Here student_num and module_code in the results relation are both **foreign** keys
 - student_num matches to the primary key of students
 - module_code matches to the primary key of modules
- Foreign keys do not need the same name as the primary key
- But it is common to have the same or similar names to help identify foreign keys

Foreign Keys Example

| <u>module_code</u> | title | teacher |
|--------------------|-------------------|--------------|
| COMP2013J | Databases | Sean Russell |
| COMP2014J | Data Structures 2 | David Lillis |
| COMP2011J | OOP | Sean Russell |

| <u>student_num</u> | grade | <u>module_code</u> |
|--------------------|---------------|----------------------|
| 1312345 | A- | COMP2013J |
| 1218985 | B+ | COMP2011J |
| 1312345 | A+ | COMP2011J |
| 1412345 | C+ | COMP2015J |

The last row contains a module code that does not exist
(violates the referential integrity constraint)

Table of Contents

- 1 Relational Model
- 2 Structure
- 3 Integrity
- 4 **Manipulating Relational Databases**
 - Example Data
 - Project Operator
 - Restrict Operator
 - Join Operator
- 5 Closure

Operators

- There are three main categories of operators we can use to manipulate relational databases

Project - choose which columns/attributes we want to see

Restrict - choose which rows/tuples we want to see (sometimes called select)

Join - combine two or more tables

students Relation

| <u>student_num</u> | name | major | year_of_entry |
|--------------------|----------------|-------|---------------|
| 19206555 | Jordan Larmour | IOT | 2019 |
| 18206123 | Jordan Murphy | SE | 2017 |
| 15205999 | Jamie Heaslip | FIN | 2015 |
| 16206456 | Gary Ringrose | EIE | 2017 |
| 18205555 | Gavin Henshaw | SE | 2018 |

courses Relation

| <u>module_code</u> | title | teacher |
|--------------------|-----------------------------------|--------------|
| COMP1002J | Intro to Programming 2 | John Dunnion |
| COMP1004J | Intro to Program Construction 1 | Sean Russell |
| COMP2013J | Databases and Information Systems | Sean Russell |
| COMP2014J | Data Structures and Algorithms 2 | David Lillis |

results Relation

| <u>student_num</u> | grade | <u>module_code</u> |
|--------------------|-------|--------------------|
| 18206123 | A+ | COMP2014J |
| 18205555 | C | COMP1004J |
| 19206555 | FM | COMP1002J |
| 16206456 | B | COMP2014J |

Projection

- The **project** operator allows you to choose which attributes/columns of a relation you want to see
- **All** rows in the relation will be returned
- This is performing a vertical cut of your relation/table

SQL Projection

- Getting data from an SQL database, requires the use of the SELECT command
- The syntax of SELECT is:
`SELECT` projection `FROM` relation;
 - E.g. `SELECT` name, major `FROM` students;
- The attributes that we list are the only ones returned in the result
- If we want all of the attributes we use `*` instead
 - E.g. `SELECT` `*` `FROM` courses;

Example of Projection

students relation

| <u>student_num</u> | name | major | year_of_entry |
|--------------------|----------------|-------|---------------|
| 19206555 | Jordan Larmour | IOT | 2019 |
| 18206123 | Jordan Murphy | SE | 2017 |
| 15205999 | Jamie Heaslip | FIN | 2015 |
| 16206456 | Gary Ringrose | EIE | 2017 |
| 18205555 | Gavin Henshaw | SE | 2018 |

```
SELECT name, major FROM students;
```

| name | major |
|----------------|-------|
| Jordan Larmour | IOT |
| Jordan Murphy | SE |
| Jamie Heaslip | FIN |
| Gary Ringrose | EIE |
| Gavin Henshaw | SE |

Example of Projection

students relation

| <u>student_num</u> | name | major | year_of_entry |
|--------------------|----------------|-------|---------------|
| 19206555 | Jordan Larmour | IOT | 2019 |
| 18206123 | Jordan Murphy | SE | 2017 |
| 15205999 | Jamie Heaslip | FIN | 2015 |
| 16206456 | Gary Ringrose | EIE | 2017 |
| 18205555 | Gavin Henshaw | SE | 2018 |

```
SELECT * FROM students;
```

| <u>student_num</u> | name | major | year_of_entry |
|--------------------|----------------|-------|---------------|
| 19206555 | Jordan Larmour | IOT | 2019 |
| 18206123 | Jordan Murphy | SE | 2017 |
| 15205999 | Jamie Heaslip | FIN | 2015 |
| 16206456 | Gary Ringrose | EIE | 2017 |
| 18205555 | Gavin Henshaw | SE | 2018 |

Restriction

- The **restrict** operator allows you to choose which **tuples/rows** of a relation you want to see
- A **subset** of the rows in the relation will be returned
- This is performing a **horizontal** cut of your relation/table

SQL Restriction

- The restrict operator adds to the SELECT command
- The updated syntax is: `SELECT` projection `FROM` relation `WHERE` restriction;
 - E.g. `SELECT * FROM students WHERE major="SE";`
 - E.g. `SELECT * FROM results WHERE grade="A+";`

Example of Restriction

students relation

| <u>student_num</u> | name | major | year_of_entry |
|--------------------|----------------|-------|---------------|
| 19206555 | Jordan Larmour | IOT | 2019 |
| 18206123 | Jordan Murphy | SE | 2017 |
| 15205999 | Jamie Heaslip | FIN | 2015 |
| 16206456 | Gary Ringrose | EIE | 2017 |
| 18205555 | Gavin Henshaw | SE | 2018 |

```
SELECT * FROM students WHERE major="SE";
```

| <u>student_num</u> | name | major | year_of_entry |
|--------------------|---------------|-------|---------------|
| 18206123 | Jordan Murphy | SE | 2017 |
| 18205555 | Gavin Henshaw | SE | 2018 |

Example of Restriction

results relation

| <u>student_num</u> | grade | <u>module_code</u> |
|--------------------|-------|--------------------|
| 18206123 | A+ | COMP2014J |
| 18205555 | C | COMP1004J |
| 19206555 | FM | COMP1002J |
| 16206456 | B | COMP2014J |

```
SELECT * FROM results WHERE grade="A+";
```

| <u>student_num</u> | grade | <u>module_code</u> |
|--------------------|-------|--------------------|
| 18206123 | A+ | COMP2014J |

Where Clause

- The WHERE clause of a SELECT statement indicates the rows that we are interested in
- E.g. `SELECT * FROM students WHERE major="SE";`
- **Every** row in the students relation is individually checked to see if the value for major is "SE"
 - Rows where the value of major is "SE" are returned
 - Other rows are not

Combining Restriction and Projection

- The restrict and project operations can be combined into one SELECT command
- `SELECT name FROM students WHERE major="SE";`
 - First, we **restrict** to only the students studying Software Engineering (SE)
 - Then, we **project** the result to only see the names of the students

Example of Restriction and Projection

students relation

| <u>student_num</u> | name | major | year_of_entry |
|--------------------|----------------|-------|---------------|
| 19206555 | Jordan Larmour | IOT | 2019 |
| 18206123 | Jordan Murphy | SE | 2017 |
| 15205999 | Jamie Heaslip | FIN | 2015 |
| 16206456 | Gary Ringrose | EIE | 2017 |
| 18205555 | Gavin Henshaw | SE | 2018 |

```
SELECT name FROM students WHERE major="SE";
```

| name |
|---------------|
| Jordan Murphy |
| Gavin Henshaw |

Selecting from Multiple Tables

- Selecting from **two** tables gets the Cartesian product of the tables
 - It combines every row from the first table with every row from the second table
- Often, this Cartesian product does not make logical sense and more is required to get a meaningful result

Cartesian Product Example 1

students relation

| <u>student_num</u> | name | major | year_of_entry |
|--------------------|----------------|-------|---------------|
| 19206555 | Jordan Larmour | IOT | 2019 |
| 18206123 | Jordan Murphy | SE | 2017 |
| 15205999 | Jamie Heaslip | FIN | 2015 |

results relation

| <u>student_num</u> | grade | <u>module_code</u> |
|--------------------|-------|--------------------|
| 18206123 | A+ | COMP2014J |
| 18205555 | C | COMP1004J |
| 19206555 | FM | COMP1002J |
| 16206456 | B | COMP2014J |

Cartesian Product Example 2

Cartesian Product of students and results

| <u>student_num</u> | name | major | year_of_entry | <u>student_num</u> | grade | <u>module_code</u> |
|--------------------|----------------|-------|---------------|--------------------|-------|--------------------|
| 19206555 | Jordan Larmour | IOT | 2019 | 18206123 | A+ | COMP2014J |
| 19206555 | Jordan Larmour | IOT | 2019 | 18205555 | C | COMP1004J |
| 19206555 | Jordan Larmour | IOT | 2019 | 19206555 | FM | COMP1002J |
| 19206555 | Jordan Larmour | IOT | 2019 | 16206456 | B | COMP2014J |
| 18206123 | Jordan Murphy | SE | 2017 | 18206123 | A+ | COMP2014J |
| 18206123 | Jordan Murphy | SE | 2017 | 18205555 | C | COMP1004J |
| 18206123 | Jordan Murphy | SE | 2017 | 19206555 | FM | COMP1002J |
| 18206123 | Jordan Murphy | SE | 2017 | 16206456 | B | COMP2014J |
| 15205999 | Jamie Heaslip | FIN | 2015 | 18206123 | A+ | COMP2014J |
| 15205999 | Jamie Heaslip | FIN | 2015 | 18205555 | C | COMP1004J |
| 15205999 | Jamie Heaslip | FIN | 2015 | 19206555 | FM | COMP1002J |
| 15205999 | Jamie Heaslip | FIN | 2015 | 16206456 | B | COMP2014J |

Join Operator

- The JOIN operator takes records from two relations based on some **join condition**

```
1 SELECT name, grade, module_code FROM students, results  
   WHERE students.student_num = results.student_num;
```

- The join condition is the clause
`students.student_num = results.student_num`

Dot Membership Operator

- The `.` in the join condition is called the dot membership operator
- `students.student_num = results.student_num`
 - `students.student_num` refers to the `student_num` attribute in `students`
 - `results.student_num` refers to the `student_num` attribute in `results`

Join Condition

- The join condition only selects the rows in the cartesian product, where the two student number values are the same

Cartesian Product of students and results

| <u>s.student_num</u> | name | major | year.of.entry | <u>r.student_num</u> | grade | <u>module_code</u> |
|----------------------|----------------|-------|---------------|----------------------|-------|--------------------|
| 19206555 | Jordan Larmour | IOT | 2019 | 18206123 | A+ | COMP2014J |
| 19206555 | Jordan Larmour | IOT | 2019 | 18205555 | C | COMP1004J |
| 19206555 | Jordan Larmour | IOT | 2019 | 19206555 | FM | COMP1002J |
| 19206555 | Jordan Larmour | IOT | 2019 | 16206456 | B | COMP2014J |
| 18206123 | Jordan Murphy | SE | 2017 | 18206123 | A+ | COMP2014J |
| 18206123 | Jordan Murphy | SE | 2017 | 18205555 | C | COMP1004J |
| 18206123 | Jordan Murphy | SE | 2017 | 19206555 | FM | COMP1002J |
| 18206123 | Jordan Murphy | SE | 2017 | 16206456 | B | COMP2014J |
| 15205999 | Jamie Heaslip | FIN | 2015 | 18206123 | A+ | COMP2014J |
| 15205999 | Jamie Heaslip | FIN | 2015 | 18205555 | C | COMP1004J |
| 15205999 | Jamie Heaslip | FIN | 2015 | 19206555 | FM | COMP1002J |
| 15205999 | Jamie Heaslip | FIN | 2015 | 16206456 | B | COMP2014J |

Join Example

```
1 SELECT name, grade, module_code FROM students, results WHERE  
   students.student_num = results.student_num;
```

Result

| name | grade | <u>module_code</u> |
|----------------|-------|--------------------|
| Jordan Larmour | FM | COMP1002J |
| Jordan Murphy | A+ | COMP2014J |

Table of Contents

- 1 Relational Model
- 2 Structure
- 3 Integrity
- 4 Manipulating Relational Databases
- 5 Closure

Closure

- The fact that the result of any operation is another relation is known as the **closure** property.
- It means that the output from one operation can be the input to another operation.
- `SELECT` name, major `FROM` students;
- The records returned by this query obey all the rules of the relational model
- This means that we can perform **another query** on the result

Closure

- The Closure property means it is possible to write nested expressions
 - This means one query inside another
- When we say the output of an operation is a relation, we mean that logically it is a relation and so is available for the next operation
- How it is actually stored, or not, is a matter for the DBMS and not something we need to worry about

Set Operations

- All of these operations are manipulating and producing relations containing data.
 - That is sets of data
- They do not work at the single record level.
- A relation of one row, or of no rows, is a valid relation and may arise because one row, or no rows, meet the criteria set by the operation
- Dealing with sets of data is a different approach to procedural data manipulation as with a procedural programming language