

Databases and Info Systems

Structured Query Language (SQL)

Dr. Seán Russell
`sean.russell@ucd.ie`,

School of Computer Science,
University College Dublin

February 26, 2020

Table of Contents

- 1 Basic SELECT Queries
 - Example Tables
 - Importing a Database
 - Basic Query Structure
- 2 More Complex SELECT Queries
- 3 Aliases and Joins
- 4 NULL Values and Duplicates

Example Tables

employees

emp_id	name	title	salary	dept_id	join_date
1234	Sean Russell	Trainer	50000	10	2018-03-01
4567	Jamie Heaslip	Manager	47000	10	2004-10-21
6542	Leo Cullen	Trainer	45000	10	2012-12-01
1238	Brendan Macken	Technician	25000	20	2001-09-10
1555	Sean O'Brien	Designer	50000	20	1999-06-24
1899	Brian O'Driscoll	Manager	45000	20	1998-02-27
2525	Peter Stringer	Designer	25000	30	2017-01-16
1585	Denis Hickey	Architect	20000	30	2009-08-07
1345	Ronan O'Gara	Manager	29000	30	2019-12-25

departments

dept_id	dept_name	office	division	manager_id
10	Training	Lansdowne	D1	4567
20	Design	Belfield	D2	1899
30	Implementation	Donnybrook	D1	1345
40	Strategy	Terenure	D2	NULL

Example Tables

appointments

title	date	start_time	end_time
Head of School	2020-02-25	15:20:00	15:40:00
Football	2020-02-25	17:30:00	19:00:00
Sleeping	2020-02-25	22:30:00	

- These tables are available in the file `week4.db` on moodle

Importing a Database

- To import a database from a file, we must complete the following steps
 - 1 Log in to mysql using your username and password (make sure the file is in the same folder as your command prompt)
 - 2 Create the database e.g. `CREATE DATABASE week4;`
 - 3 Select the database e.g. `USE week4`
 - 4 Load the file e.g. `source week4.db`
- Now all of the tables and the data stored in them will be loaded and ready to query

SQL Select Query Syntax

- SQL queries are expressed by the SELECT statement.
- The basic syntax is:

```
1 SELECT attr_expr {, attr_expr } FROM table_name {,  
   table_name } [ WHERE condition ] ;
```

- [] means an optional expression
- {} means an optional list of expressions

Select Query Parts

```
1 SELECT attr_expr {, attr_expr } FROM table_name {,  
   table_name } [ WHERE condition ] ;
```

- The three parts of the query are usually called:
 - Target list (the attributes you want to retrieve, and/or expressions based on these attributes).
 - From clause (the table(s) to select from)
 - Where clause (the condition on which to select rows)

Order

- Queries are completed in a particular order
 - 1 The Cartesian product of the tables in the **FROM** clause is generated
 - 2 The result is restricted to the rows that satisfy the condition in there **WHERE** clause
 - 3 The attribute expressions in the target list are evaluated

Simple Query Examples

```
1 SELECT salary FROM employees WHERE title =  
   "Technician";
```

```
1 +-----+  
2 | salary |  
3 +-----+  
4 | 25000  |  
5 +-----+  
6 1 row in set (0.00 sec)
```

Simple Query Examples

```
1 SELECT * FROM employees WHERE title =  
   "Manager";
```

```
1 |-----|  
2 | emp_id | name          | title   | salary | dept_id | join_date |  
3 |-----|  
4 | 1345   | Ronan O' Gara | Manager | 29000  | 30      | 2019-12-25 |  
5 | 1899   | Brian O' Driscoll | Manager | 45000  | 20      | 1998-02-27 |  
6 | 4567   | Jamie Heaslip | Manager | 47000  | 10      | 2004-10-21 |  
7 |-----|  
8 3 rows in set (0.00 sec)
```

Table of Contents

1 Basic SELECT Queries

2 More Complex SELECT Queries

- Comparison Operators
- Complex Clauses
- Comparison Functions
- String Functions
- Date and Time Functions
- Functions in Target List

3 Aliases and Joins

4 NULL Values and Duplicates

Comparison Operators

- SQL allows the use of the typical comparison operators
 - = Equality
 - < Less than
 - > Greater than
 - <= Less than or equal
 - >= Greater than or equal
 - <> Not equal
 - != Not equal

Complex WHERE Clauses

- The **WHERE** clause is a boolean expression
 - it is evaluated for every row in the table
- We can use boolean operators **AND**, **OR**, **NOT** to combine simple expressions into more complex ones
- Find all managers who earn at least 30000

```
1 SELECT * FROM employees WHERE  
   title="Manager" AND salary >= 30000;
```

```
1 |-----|  
2 | emp_id | name           | title      | salary | dept_id | join_date |  
3 |-----|  
4 | 1899   | Brian O'Driscoll | Manager    | 45000  | 20      | 1998-02-27 |  
5 | 4567   | Jamie Heaslip   | Manager    | 47000  | 10      | 2004-10-21 |  
6 |-----|  
7 2 rows in set (0.00 sec)
```

Brackets in Where Clause

- You can use parentheses (round brackets) to group boolean statements correctly
- Find all managers who work in either department 10 or department 20

```
1 SELECT * FROM employees WHERE  
   title="Manager" AND (dept_id="10" OR  
   dept_id="20");
```

```
1 |-----|  
2 | emp_id | name           | title    | salary | dept_id | join_date |  
3 |-----|  
4 | 1899   | Brian O'Driscoll | Manager  | 45000  | 20      | 1998-02-27 |  
5 | 4567   | Jamie Heaslip   | Manager  | 47000  | 10      | 2004-10-21 |  
6 |-----|  
7 2 rows in set (0.00 sec)
```

Comparison Functions

- Full list of comparison functions and operators in MySQL:
 - <https://dev.mysql.com/doc/refman/8.0/en/comparison-operators.html>
- Interesting operators:
 - BETWEEN ... AND
 - IN

BETWEEN ... AND

- Search for values in a particular range
- Find all employees paid between 25000 and 40000

```
1 SELECT * FROM employees WHERE salary BETWEEN 25000  
   AND 40000;
```

- Is the same as

```
1 SELECT * FROM employees WHERE salary >= 25000 AND  
   salary <= 40000;
```

```
1 |-----|  
2 | emp_id | name          | title      | salary | dept_id | join_date |  
3 |-----|  
4 | 1238   | Brendan Macken | Technician | 25000  | 20      | 2001-09-10 |  
5 | 1345   | Ronan O'Gara   | Manager   | 29000  | 30      | 2019-12-25 |  
6 | 2525   | Peter Stringer | Designer   | 25000  | 30      | 2017-01-16 |  
7 |-----|  
8 3 rows in set (0.00 sec)
```


IN

- Check if a value is contained in a set of values
- Find all employees with the title Trainer, Designer, or Architect

```
1 SELECT * FROM employees WHERE title IN  
   ("Trainer", "Designer", "Architect");
```

- Same as the clause `WHERE title="Trainer" OR title="Designer" OR title="Architect"`

```
1 |-----|  
2 | emp_id | name           | title      | salary | dept_id | join_date |  
3 |-----|  
4 | 1234   | Sean Russell   | Trainer    | 50000  | 10      | 2018-03-01 |  
5 | 1555   | Sean O'Brien  | Designer   | 50000  | 20      | 1999-06-24 |  
6 | 1585   | Denis Hickey   | Architect  | 20000  | 30      | 2009-08-07 |  
7 | 2525   | Peter Stringer | Designer   | 25000  | 30      | 2017-01-16 |  
8 | 6542   | Leo Cullen     | Trainer    | 45000  | 10      | 2012-12-01 |  
9 |-----|  
10 5 rows in set (0.00 sec)
```

String Functions

- There are many functions that we can apply to attributes in the WHERE clause
- A full list of the MYSQL String functions can be found here:
 - <https://dev.mysql.com/doc/refman/8.0/en/string-functions.html>
- Interesting functions
 - LIKE
 - LENGTH
 - RIGHT

LIKE

- Search for patterns in string data
- Find all employees where the second letter of their name is e

```
1 SELECT * FROM employees WHERE name LIKE "_e%";
```

- "_" represents a single character
- "%" represents any number of characters

```
1 |-----|
2 | emp_id | name      | title   | salary | dept_id | join_date |
3 |-----|
4 | 1234   | Sean Russell | Trainer | 50000  | 10      | 2018-03-01 |
5 | 1555   | Sean O'Brien | Designer | 50000  | 20      | 1999-06-24 |
6 | 1585   | Denis Hickey | Architect | 20000  | 30      | 2009-08-07 |
7 | 2525   | Peter Stringer | Designer | 25000  | 30      | 2017-01-16 |
8 | 6542   | Leo Cullen   | Trainer | 45000  | 10      | 2012-12-01 |
9 |-----|
10 5 rows in set (0.00 sec)
```

LIKE

- Find all employees whose job title contains exactly seven letters

```
1 SELECT * FROM employees WHERE title LIKE " _ _ _ _ _ _ _ ";
```

```
1 |-----|
2 | emp_id | name          | title    | salary | dept_id | join_date |
3 |-----|
4 | 1234    | Sean Russell  | Trainer  | 50000  | 10      | 2018-03-01 |
5 | 1345    | Ronan O'Gara  | Manager  | 29000  | 30      | 2019-12-25 |
6 | 1899    | Brian O'Driscoll | Manager  | 45000  | 20      | 1998-02-27 |
7 | 4567    | Jamie Heaslip | Manager  | 47000  | 10      | 2004-10-21 |
8 | 6542    | Leo Cullen    | Trainer  | 45000  | 10      | 2012-12-01 |
9 |-----|
10 5 rows in set (0.00 sec)
```

LENGTH

- Get the length of a string in an attribute
- Find all employees who's job title contains exactly seven letters

```
1 SELECT * FROM employees WHERE LENGTH(title)=7;
```

```
1 |-----|
2 | emp_id | name      | title   | salary | dept_id | join_date |
3 |-----|
4 | 1234   | Sean Russell | Trainer | 50000  | 10      | 2018-03-01 |
5 | 1345   | Ronan O'Gara | Manager | 29000  | 30      | 2019-12-25 |
6 | 1899   | Brian O'Driscoll | Manager | 45000  | 20      | 1998-02-27 |
7 | 4567   | Jamie Heaslip | Manager | 47000  | 10      | 2004-10-21 |
8 | 6542   | Leo Cullen  | Trainer | 45000  | 10      | 2012-12-01 |
9 |-----|
10 5 rows in set (0.00 sec)
```

RIGHT (and LEFT)

- The `RIGHT(str,num)` function returns the last `num` characters of the string `str`
- Find all employees whose names end in "ll"

```
1 SELECT * FROM employees WHERE RIGHT(name,2)="ll";
```

```
1 |-----|
2 | emp_id | name           | title    | salary | dept_id | join_date |
3 |-----|
4 | 1234   | Sean Russell   | Trainer  | 50000  | 10      | 2018-03-01 |
5 | 1899   | Brian O'Driscoll | Manager  | 45000  | 20      | 1998-02-27 |
6 |-----|
7 2 rows in set (0.00 sec)
```

Date and Time Functions

- Other data types also have available functions
- A full list of the MySQL date and time functions can be found here:
 - <https://dev.mysql.com/doc/refman/8.0/en/date-and-time-functions.html>
- Interesting functions
 - `CURDATE()`
 - `CURTIME()`
 - `NOW()`
 - `DATEDIFF`

CURDATE

- Gets the current date (also works in target list)

```
1 SELECT CURDATE();
```

```
1 +-----+
2 | CURDATE() |
3 +-----+
4 | 2020-02-25 |
5 +-----+
6 1 row in set (0.00 sec)
```

- Find employees who were hired today

```
1 SELECT * FROM employees WHERE join_date =
    CURDATE();
```


CURTIME

- Gets the current time (also works in target list)

```
1 SELECT CURTIME();
```

```
1 +-----+
2 | CURTIME() |
3 +-----+
4 | 17:01:07 |
5 +-----+
6 1 row in set (0.00 sec)
```

- Find all appointments today that have finished:

```
1 SELECT * FROM appointments WHERE date = CURDATE()
   AND end_time < CURTIME();
```

```
1 +-----+-----+-----+-----+
2 | title          | date       | start_time | end_time   |
3 +-----+-----+-----+-----+
4 | Head of School | 2020-02-25 | 15:20:00   | 15:40:00   |
5 | Football       | 2020-02-25 | 17:30:00   | 19:00:00   |
6 +-----+-----+-----+-----+
7 2 rows in set (0.00 sec)
```

NOW

- Work very similarly to CURTIME, but returns a DATETIME object instead of a TIME object
- Works in both target list and WHERE clause

DATEDIFF

- This function returns the number of days between two DATE or DATETIME values
- Syntax: DATEDIFF(expr1,expr2)
- Find all employees who have joined in the last 90 days

```
1 SELECT * FROM employees WHERE DATEDIFF(CURDATE() ,  
    join_date) < 90;
```

```
1 +-----+-----+-----+-----+-----+-----+  
2 | emp_id | name      | title  | salary | dept_id | join_date |  
3 +-----+-----+-----+-----+-----+-----+  
4 | 1345   | Ronan O'Gara | Manager | 29000  | 30      | 2019-12-25 |  
5 +-----+-----+-----+-----+-----+-----+  
6 1 row in set (0.00 sec)
```

Expressions in Target List

- Expressions can be used in the target list
- Assuming the database contains annual salary, find the monthly salary of every employee

```
1 SELECT name, salary/12 FROM employees;
```

```
1 |-----|
2 | name          | salary/12 |
3 |-----|
4 | Sean Russell  | 4166.6667 |
5 | Brendan Macken| 2083.3333 |
6 | Ronan O'Gara  | 2416.6667 |
7 | Sean O'Brien  | 4166.6667 |
8 | Denis Hickey  | 1666.6667 |
9 | Brian O'Driscoll| 3750.0000 |
10 | Peter Stringer| 2083.3333 |
11 | Jamie Heaslip | 3916.6667 |
12 | Leo Cullen    | 3750.0000 |
13 |-----|
14 9 rows in set (0.00 sec)
```

Expressions in Target List

- The previous example gave ugly output
- We can use a function to round the answer
 - There are other functions available here:
 - <https://dev.mysql.com/doc/refman/8.0/en/numeric-functions.html>

```
1 SELECT name, ROUND(salary/12) FROM employees;
```

```
1 +-----+-----+
2 | name          | ROUND(salary/12) |
3 +-----+-----+
4 | Sean Russell  |          4167    |
5 | ...           |                  |
6 +-----+-----+
7 9 rows in set (0.00 sec)
```

Table of Contents

- 1 Basic SELECT Queries
- 2 More Complex SELECT Queries
- 3 Aliases and Joins
 - Aliases
 - More Complex Joins
- 4 NULL Values and Duplicates

Aliases

- Using expressions and functions in the target list gives us complex headings for our results
- We can use an alias to temporarily give a descriptive name

```
1 SELECT name, ROUND(salary/12) AS month_salary FROM  
   employees;
```

```
1 |-----|-----|  
2 | name          | month_salary |  
3 |-----|-----|  
4 | Sean Russell  |          4167 |  
5 | ...          |              |  
6 |-----|-----|  
7 9 rows in set (0.00 sec)
```

Aliases on Tables

- We can also use aliases for tables to shorten our queries
- This becomes very useful for complex joins

```
1 SELECT name, dept_name FROM employees, departments  
   WHERE employees.dept_id=departments.dept_id;
```

- Is the same as:

```
1 SELECT name, dept_name FROM employees AS e,  
   departments AS d WHERE e.dept_id=d.dept_id;
```

- This is more useful the more joins you have

More Complex Joins

- One of the big challenges of using databases is translating a requirement in human language and writing a suitable SQL query to answer it
- Find the names of the employees who work in the Lansdowne office of division D1

More Complex Joins

- Find the names of the employees who work in the Lansdowne office of division D1
- Query:

```
1 SELECT name FROM employees AS e, departments AS d
   WHERE e.dept_id = d.dept_id AND d.division =
      'D1' AND d.office = 'Lansdowne';
```

- Result:

```
1 |-----|
2 | name
3 |-----|
4 | Sean Russell
5 | Jamie Heaslip
6 | Leo Cullen
7 |-----|
8 3 rows in set (0.00 sec)
```

More Complex Joins

- Find the names of the employees who work in either the Lansdowne office or the Belfield office
- Query 1:

```
1 SELECT name FROM employees AS e, departments AS d
   WHERE e.dept_id = d.dept_id AND
      (d.office="Belfield" OR d.office="Lansdowne");
```

- Query 2:

```
1 SELECT name FROM employees AS e, departments AS d
   WHERE e.dept_id = d.dept_id AND d.office
      IN("Belfield", "Lansdowne");
```

Table of Contents

- 1 Basic SELECT Queries
- 2 More Complex SELECT Queries
- 3 Aliases and Joins
- 4 NULL Values and Duplicates
 - NULL Values
 - Duplicates

NULL Values

- NULL values may mean that:
 - a value is not applicable
 - a value is applicable but unknown
 - it is unknown if a value is applicable or not

Comparing NULL Values

- Previous standards of SQL used two-valued logic
 - Only the values TRUE and FALSE
 - Comparing against a NULL value returns FALSE
- SQL-2 (and later) use a three-valued logic
 - Uses the values TRUE, FALSE and UNKNOWN
 - Comparing against a NULL value returns UNKNOWN

Testing for NULL Values

- In a SELECT query, we may want to test if an attribute contains a NULL value.
 - attribute IS NULL
 - attribute IS NOT NULL
- Find the names of all departments that do not have a manager

```
1 SELECT dept_name FROM departments WHERE manager_id  
   IS NULL;
```

```
1 +-----+  
2 | dept_name |  
3 +-----+  
4 | Strategy |  
5 +-----+  
6 1 row in set (0.00 sec)
```

Testing for NOT NULL Values

- In a SELECT query, we may want to test if an attribute contains a NULL value.
 - attribute IS NULL
 - attribute IS NOT NULL
- Find the names of all departments that have a manager

```
1 SELECT dept_name FROM departments WHERE manager_id  
   IS NOT NULL;
```

```
1 +-----+  
2 | dept_name |  
3 +-----+  
4 | Training  |  
5 | Design    |  
6 | Implementation |  
7 +-----+  
8 3 rows in set (0.00 sec)
```


Duplicates

- In SQL, results may have identical rows
- Duplicates can be removed using the keyword **DISTINCT**

```
1 SELECT title FROM employees;
```

```
1 +-----+
2 | title |
3 +-----+
4 | Trainer
5 | Technician
6 | Manager
7 | Designer
8 | Architect
9 | Manager
10 | Designer
11 | Manager
12 | Trainer
13 +-----+
14 9 rows in set (0.00 sec)
```

Duplicates

- In SQL, results may have identical rows
- Duplicates can be removed using the keyword **DISTINCT**

```
1 SELECT DISTINCT title FROM employees;
```

```
1 +-----+
2 | title
3 +-----+
4 | Trainer
5 | Technician
6 | Manager
7 | Designer
8 | Architect
9 +-----+
10 5 rows in set (0.00 sec)
```