

Databases and Info Systems

Structured Query Language (SQL)

Dr. Seán Russell
`sean.russell@ucd.ie`,

School of Computer Science,
University College Dublin

March 4, 2020

Table of Contents

1 Example Data

- Example Tables → week5.db

2 Joining Tables

3 Ordering and Limiting Data

4 Aggregate Queries

5 Grouping Data

lecturers

lec_id	name
1	Sean Russell
2	David Lillis
3	Catherine Mooney
4	Shen Wang
5	Ruhai Dong
6	Brett Becker
7	John Dunnion

modules

code	module_name	lec_id
COMP1005J	Programming 2	7
COMP2001J	Computer Networks	4
COMP1002J	Program Construction 2	6
COMP2004J	Databases	5
COMP2007J	Computer Organisation	4
COMP2013J	Databases	1
COMP2003J	Data Structures 2	8
MATH4023J	Really Hard Maths	15

employees

emp_id	name	title	salary	dept_id	join_date
1234	Sean Russell	Trainer	50000	10	2018-03-01
4567	Jamie Heaslip	Manager	47000	10	2004-10-21
6542	Leo Cullen	Trainer	45000	10	2012-12-01
1238	Brendan Macken	Technician	25000	20	2001-09-10
1555	Sean O'Brien	Designer	50000	20	1999-06-24
1899	Brian O'Driscoll	Manager	45000	20	1998-02-27
2525	Peter Stringer	Designer	25000	30	2017-01-16
1585	Denis Hickey	Architect	20000	30	2009-08-07
1345	Ronan O'Gara	Manager	29000	30	2019-12-25

departments

dept_id	dept_name	office	division	manager_id
10	Training	Lansdowne	D1	4567
20	Design	Belfield	D2	1899
30	Implementation	Donnybrook	D1	1345
40	Strategy	Terenure	D2	NULL

Table of Contents

- 1 Example Data
- 2 Joining Tables
 - INNER JOIN
 - LEFT JOIN
 - RIGHT JOIN
 - FULL JOIN
 - USING and NATURAL JOIN

3 Ordering and Limiting Data

4 Aggregate Queries

5 Grouping Data

Joining Tables

- SQL-2 introduced a new syntax for joins
- Joins are represented explicitly in the FROM clause

```
SELECT attr_expr [AS alias] {, attr_expr [AS  
    alias] } FROM table [AS alias] { [JoinType]  
    JOIN table [AS alias] ON JoinConditions } [  
    WHERE OtherCondition ]
```

- [] means an optional expression
- {} means an optional list of expressions

Join Types

- It is useful to learn the Cartesian product type of join (that we have seen before) when you start learning databases
- But generally they are not widely used.
- Instead, we normally use one of the following types of join:
 - INNER
 - RIGHT [OUTER]
 - LEFT [OUTER]
 - FULL [OUTER]

Inner Join

- Inner Join joins two tables together based on some condition
- Results are only returned for rows where there is a matching result in both tables
- If there is no match for a row in either table, then the row is not shown

Inner Join Example

- Query: Find all of the lecturers and the modules that they teach
- SQL:

```
SELECT * FROM lecturers AS l INNER JOIN modules AS m ON  
l.lec_id = m.lec_id;
```

- Result:

lec_id	name	code	module_name	lec_id
6	Brett Becker	COMP1002J	Program Construction 2	6
7	John Dunnion	COMP1005J	Programming 2	7
4	Shen Wang	COMP2001J	Computer Networks	4
9	Henry Mcloughlin	COMP2002J	Data Structures 1	9
5	Ruhai Dong	COMP2004J	Databases	5
4	Shen Wang	COMP2007J	Computer Organisation	4
1	Sean Russell	COMP2013J	Databases	1
7 rows in set (0.00 sec)				

Left Join

- Left Join joins two tables together based on some condition
- Results are only returned for **every row** in the table on the **left** of the join
- If there is no match for a row in the right table, then the columns all show the value **NULL**

Left Join Example

- Query: Find the lecturers and their modules, including the lecturers that don't teach modules
- SQL:

```
SELECT * FROM lecturers AS l LEFT JOIN modules AS m ON l.lec_id = m.lec_id;
```

- Result:

lec_id	name	code	module_name	lec_id
6	Brett Becker	COMP1002J	Program Construction 2	6
7	John Dunnion	COMP1005J	Programming 2	7
4	Shen Wang	COMP2001J	Computer Networks	4
9	Henry Mccloughlin	COMP2002J	Data Structures 1	9
5	Ruhai Dong	COMP2004J	Databases	5
4	Shen Wang	COMP2007J	Computer Organisation	4
1	Sean Russell	COMP2013J	Databases	1
2	David Lillis	NULL	NULL	NULL
3	Catherine Mooney	NULL	NULL	NULL
8	Takfarinas Saber	NULL	NULL	NULL

10 rows in set (0.00 sec)

Right Join

- Join two tables together based on some condition
- Results are returned for **every row** in the table on the **right** of the join
- If there is no match for a row in the left table, then the columns all show the value **NULL**

Right Join Example

- Query: Find the lecturers and modules, including the modules that do not have a lecturer in the database
- SQL:

```
SELECT * FROM lecturers AS l RIGHT JOIN modules AS m ON
l.lec_id = m.lec_id;
```

- Result:

lec_id	name	code	module_name	lec_id
6	Brett Becker	COMP1002J	Program Construction 2	6
7	John Dunnion	COMP1005J	Programming 2	7
4	Shen Wang	COMP2001J	Computer Networks	4
9	Henry Mccloughlin	COMP2002J	Data Structures 1	9
5	Ruhai Dong	COMP2004J	Databases	5
4	Shen Wang	COMP2007J	Computer Organisation	4
1	Sean Russell	COMP2013J	Databases	1
NULL	NULL	MATH4023J	Really Hard Maths	15

8 rows in set (0.00 sec)

Full Join

- Join two tables together based on some condition
- Results are returned for **every row** in the **both tables**
- If there is no match for a row in the left table, then the columns all show the value **NULL**
- If there is no match for a row in the right table, then the columns all show the value **NULL**

Right Join Example

- Query: Find the lecturers and modules, including the modules that do not have a lecturer in the database
- SQL:

```
SELECT * FROM lecturers AS l FULL JOIN modules AS m ON l.lec_id  
= m.lec_id;
```

- Result:

```
ERROR 1064 (42000): You have an error in your SQL syntax; check the manual that  
corresponds to your MySQL server version for the right syntax to use near  
'FULL JOIN modules AS m ON l.lec_id = m.lec_id' at line 1
```

- Full join is not supported in MySQL...

Alternative Full Join

- We can implement a full join, by performing 2 queries and putting the results together
- A left join and a right join put together using a UNION
- SQL:

```
SELECT l.lec_id, name, code, module_name FROM lecturers AS l LEFT  
JOIN modules AS m ON l.lec_id = m.lec_id UNION SELECT  
m.lec_id, name, code, module_name FROM lecturers AS l RIGHT  
JOIN modules as m ON l.lec_id = m.lec_id;
```

Alternative Full Join

- We can implement a full join, by performing 2 queries and putting the results together
- A left join and a right join put together using a UNION
- Result:

lec_id	name	code	module_name
6	Brett Becker	COMP1002J	Program Construction 2
7	John Dunnion	COMP1005J	Programming 2
4	Shen Wang	COMP2001J	Computer Networks
9	Henry Mcloughlin	COMP2002J	Data Structures 1
5	Ruhai Dong	COMP2004J	Databases
4	Shen Wang	COMP2007J	Computer Organisation
1	Sean Russell	COMP2013J	Databases
2	David Lillis	NULL	NULL
3	Catherine Mooney	NULL	NULL
8	Takfarinas Saber	NULL	NULL
15	NULL	MATH4023J	Really Hard Maths

11 rows in set (0.00 sec)

LEFT JOIN to find missing relationships

- LEFT JOIN will return every row in the left table
- If it does not match anything in the right table, these attributes will contain NULL values
- Find the names of all lecturers who do not teach any module
- SQL:

```
SELECT name FROM lecturers AS l LEFT JOIN modules as m ON  
l.lec_id = m.lec_id WHERE code IS NULL;
```

- Result:

name
David Lillis
Catherine Mooney
Takfarinas Saber

3 rows in set (0.00 sec)

USING

- Joining tables is a very common operation in SQL
- When computer scientists do something a lot, we find ways to simplify it
- Many joins between two tables are that contain an attribute with the same name that we are comparing to find where they are the same
- In this situation we can use the keyword USING
- Syntax:

```
table [JoinType] JOIN table USING(attribute)
```

Using Examples

- Inner Join:

```
SELECT * FROM lecturers INNER JOIN modules USING(lec_id);
```

- Left Join:

```
SELECT * FROM lecturers LEFT JOIN modules USING(lec_id);
```

- Right Join:

```
SELECT * FROM lecturers RIGHT JOIN modules USING(lec_id);
```


Natural Join

- Natural Join takes this logic of saving time a little further
- Using a natural join, we perform an inner join using any attributes with the same name in both tables
- In our examples, this would be on the column `lec_id`, but it could be others
- Syntax:

```
table NATURAL JOIN table
```

Natural Example

- Inner Join:

```
SELECT * FROM lecturers NATURAL JOIN modules;
```

- Result:

lec_id	name	code	module_name
6	Brett Becker	COMP1002J	Program Construction 2
7	John Dunnion	COMP1005J	Programming 2
4	Shen Wang	COMP2001J	Computer Networks
9	Henry Mcloughlin	COMP2002J	Data Structures 1
5	Ruhai Dong	COMP2004J	Databases
4	Shen Wang	COMP2007J	Computer Organisation
1	Sean Russell	COMP2013J	Databases

7 rows in set (0.00 sec)

Table of Contents

- 1 Example Data
- 2 Joining Tables
- 3 Ordering and Limiting Data
 - Ordering Data
 - Limiting the Number of Rows
- 4 Aggregate Queries
- 5 Grouping Data

Ordering Data

- To order our data in a particular way, we use the `ORDER BY` clause
- Example

```
SELECT name, salary FROM employees ORDER BY salary;
```

- Employees are ordered by salary
- The default behaviour is to sort in ascending order (i.e. from lowest to highest).
- You can reverse the sort by using the `DESC` keyword (short for **descending**)

Ordering Example

- Example

```
SELECT name, salary FROM employees ORDER BY salary;
```

- Result

name	salary
Denis Hickey	20000
Brendan Macken	25000
Peter Stringer	25000
Ronan O'Gara	29000
Brian O'Driscoll	45000
Leo Cullen	45000
Jamie Heaslip	47000
Sean Russell	50000
Sean O'Brien	50000

9 rows in set (0.00 sec)

Ordering Example

- Example

```
SELECT name, salary FROM employees ORDER BY salary DESC;
```

- Result

name	salary
Sean Russell	50000
Sean O'Brien	50000
Jamie Heaslip	47000
Brian O'Driscoll	45000
Leo Cullen	45000
Ronan O'Gara	29000
Brendan Macken	25000
Peter Stringer	25000
Denis Hickey	20000

9 rows in set (0.00 sec)

More Ordering

- Some employees in the database have equal salaries. You can ask these to be sorted by something else

```
SELECT name, salary FROM employees ORDER BY salary DESC,  
name;
```

- Employees are ordered by salary (descending)
- If two employees have the same salary, these will be ordered by their name.
- Salary is ordered in descending order, but name in ascending order

Ordering Example

- Example

```
SELECT name, salary FROM employees ORDER BY salary DESC,  
name;
```

- Result

name	salary
Sean O'Brien	50000
Sean Russell	50000
Jamie Heaslip	47000
Brian O'Driscoll	45000
Leo Cullen	45000
Ronan O'Gara	29000
Brendan Macken	25000
Peter Stringer	25000
Denis Hickey	20000

9 rows in set (0.00 sec)

- We can use the LIMIT keyword to reduce the number of rows returned by a SELECT query
- This is useful for:
 - Exploring a large new database (you may not need to see all of the data)
 - Programs that display part of the data to users (there is no need to get all the data in a query)
- Syntax:

```
SELECT * FROM employees LIMIT 5;
```

Multiple Arguments

- We can give two arguments to LIMIT
 - The first is the starting index
 - The second is the maximum number of rows
- First 5 rows:

```
SELECT name, salary FROM employees ORDER BY salary, name  
LIMIT 0,5;
```

- Next 5 rows:

```
SELECT name, salary FROM employees ORDER BY salary, name  
LIMIT 5,5;
```

```
SELECT name, salary FROM employees ORDER BY salary, name LIMIT 0,5;
```

name	salary
Denis Hickey	20000
Brendan Macken	25000
Peter Stringer	25000
Ronan O'Gara	29000
Brian O'Driscoll	45000

5 rows in set (0.00 sec)

```
SELECT name, salary FROM employees ORDER BY salary, name LIMIT 5,5;
```

name	salary
Leo Cullen	45000
Jamie Heaslip	47000
Sean O'Brien	50000
Sean Russell	50000

4 rows in set (0.00 sec)

Table of Contents

- 1 Example Data
- 2 Joining Tables
- 3 Ordering and Limiting Data
- 4 Aggregate Queries
 - Count
 - Sum
 - Max
 - Min
 - Average
 - Complex Aggregate Queries

Aggregate Queries

- An aggregate query performs some calculation on multiple rows
- SQL-2 has 5 aggregate operators
 - COUNT - count the number of rows
 - SUM - calculate the sum of an attribute for a set of rows
 - MAX - find the maximum value of an attribute for a set of rows
 - MIN - find the minimum value of an attribute for a set of rows
 - AVG - calculate the average of an attribute for a set of rows

Count

- COUNT returns the number of rows or distinct values;
- It can count the number of rows

```
SELECT COUNT(*) FROM departments;
```

```
+-----+  
| COUNT(*) |  
+-----+  
|         4 |  
+-----+  
1 row in set (0.00 sec)
```

Count

- COUNT returns the number of rows or distinct values;
- Find the number of values in a particular column (NULL values are not counted)

```
SELECT COUNT(manager_id) FROM departments;
```

```
+-----+  
| COUNT(manager_id) |  
+-----+  
|                   3 |  
+-----+  
1 row in set (0.00 sec)
```

Count

- COUNT returns the number of rows or distinct values;
- Find the number of different values in a particular column (duplicates values are not counted)

```
SELECT COUNT(DISTINCT division) FROM departments;
```

```
+-----+  
| COUNT(DISTINCT division) |  
+-----+  
|                           2 |  
+-----+  
1 row in set (0.00 sec)
```


Sum

- The SUM function returns the sum of a set of value (ignores NULL values)
- Example: Find the total cost of all employee salaries
- SQL:

```
SELECT SUM(salary) FROM employees;
```

- Result

```
+-----+  
| SUM( salary ) |  
+-----+  
|          336000 |  
+-----+  
1 row in set (0.00 sec)
```

Max

- The MAX function returns the maximum value in a set of value
- Example: Find the highest salary of all employee
- SQL:

```
SELECT MAX(salary) FROM employees;
```

- Result

```
+-----+  
| MAX(salary) |  
+-----+  
|          50000 |  
+-----+  
1 row in set (0.00 sec)
```

Min

- The MIN function returns the minimum value in a set of value
- Example: Find the lowest salary of all employee
- SQL:

```
SELECT MIN(salary) FROM employees;
```

- Result

```
+-----+  
| MIN(salary) |  
+-----+  
|          20000 |  
+-----+  
1 row in set (0.00 sec)
```

Average

- The AVG function returns the average value of a set of values (ignores NULL values).
- Example: Find the average salary for all employee
- SQL:

```
SELECT AVG(salary) FROM employees;
```

- Result

```
+-----+  
| AVG(salary) |  
+-----+  
| 37333.3333 |  
+-----+  
1 row in set (0.00 sec)
```

More Complex Aggregate Queries

- Aggregate query with restriction
- Find the highest paid employee who is a manager

```
SELECT MAX(salary) FROM employees WHERE title="Manager";
```

- Result:

MAX(salary)
47000

1 row in set (0.00 sec)

More Complex Aggregate Queries

- Aggregate query with join
- Find the average salary of all employees in the design department

```
SELECT AVG(salary) FROM employees INNER JOIN departments  
      USING(dept_id) WHERE dept_name="Design";
```

- Result:

```
+-----+  
| AVG( salary ) |  
+-----+  
| 40000.0000    |  
+-----+  
1 row in set (0.00 sec)
```

Aggregate Queries and Target List

- Attributes in a query must have same number of results
- You can't mix aggregated and non-aggregated attributes
 - unless you use GROUP BY, which we will discuss later
- Example:

```
SELECT name, MAX(salary) FROM employees;
```

- Result:

```
ERROR 1140 (42000): In aggregated query without GROUP BY, expression #1 of  
SELECT list contains nonaggregated column 'week5.employees.name'; this is  
incompatible with sql_mode=only_full_group_by
```

Aggregate Queries and Target List

- We can have queries using multiple aggregate functions
- They will all have the same number of results (1)
- Example:

```
SELECT MAX(salary), MIN(salary), AVG(salary) FROM employees;
```

- Result:

MAX(salary)	MIN(salary)	AVG(salary)
50000	20000	37333.3333
1 row in set (0.00 sec)		

Table of Contents

- 1 Example Data
- 2 Joining Tables
- 3 Ordering and Limiting Data
- 4 Aggregate Queries
- 5 Grouping Data**
 - **GROUP BY**

- Queries may apply aggregate functions to only a **subset of rows**
- This allows us to use aggregate and non-aggregate attributes together
 - non-aggregate attributes must be in the GROUP BY clause
- For example, we might want to find the sum of all salaries for each department
- SQL:

```
SELECT dept_name, SUM(salary) FROM employees NATURAL JOIN
departments GROUP BY dept_name;
```

- Result:

dept_name	SUM(salary)
Training	142000
Design	120000
Implementation	74000

3 rows in set (0.00 sec)

Group predicates

- When conditions are on the result of an aggregate operator, it is necessary to use the HAVING clause
- For example, we might want to Find the departments where the average salary is over 40000
- SQL:

```
SELECT dept_id, AVG(salary) FROM employees GROUP BY dept_id  
HAVING AVG(salary) > 40000;
```

- Result:

dept_id	AVG(salary)
10	47333.3333

1 row in set (0.00 sec)

How It Works

- 1 The query is executed without GROUP BY and without aggregate operators (the WHERE clause is applied at this stage)
- 2 The query result is divided in subsets with the same values for the attributes appearing after the group by clause
- 3 The aggregate operator is applied separately to each subset (the HAVING clause is applied at this stage)

WHERE or HAVING

- Only expressions containing **aggregate operators** should appear in the argument of the HAVING clause, because the HAVING clause is applied after aggregation takes place
- The WHERE clause is applied before the aggregation takes place
- For example, we might want to find the name of the departments where the average salary is over 40000
- SQL:

```
SELECT dept_name FROM employees INNER JOIN departments  
    USING(dept_id) GROUP BY dept_name HAVING AVG(salary) >  
    40000;
```

WHERE and HAVING

- WHERE and HAVING can be used together
- Find the names of all departments that have an average salary less than 40000 for its employees that are not managers
- SQL:

```
SELECT dept_name, AVG(salary)
FROM departments INNER JOIN employees USING(dept_id)
WHERE title!="Manager"
GROUP BY dept_id
HAVING AVG(salary) < 40000;
```

- Note: WHERE must be before GROUP BY

Example

- Average salary of all departments

```
SELECT dept_name, AVG(salary) FROM departments INNER JOIN  
employees USING(dept_id) GROUP BY dept_id;
```

dept_name	AVG(salary)
Training	47333.3333
Design	40000.0000
Implementation	24666.6667

3 rows in set (0.00 sec)

Example

- Average salary less than 40000 (for all employees)

```
SELECT dept_name, AVG(salary) FROM departments INNER JOIN  
employees USING(dept_id) GROUP BY dept_id HAVING  
AVG(salary) < 40000;
```

dept_name	AVG(salary)
Implementation	24666.6667

1 row in set (0.00 sec)

Example

- Average salary less than 40000 (non-managers only)

```
SELECT dept_name, AVG(salary) FROM departments INNER JOIN  
employees USING(dept_id) WHERE title != 'Manager' GROUP BY  
dept_id HAVING AVG(salary) < 40000;
```

dept_name	AVG(salary)
Design	37500.0000
Implementation	22500.0000
2 rows in set (0.00 sec)	