

Databases and Info Systems

Normalisation

Dr. Seán Russell
`sean.russell@ucd.ie`,

School of Computer Science,
University College Dublin

March 27, 2020



Table of Contents

1 What is Normalisation? and Why Normalise Databases?

- Problem of Redundancy
- Integrity Constraints
- Why Normalise Databases?
- Saving Space
- Prevent Anomalies
- Avoiding NULL Values

2 Functional Dependencies

3 Normal Forms

- **Normalisation** is the process of transforming data from a problem into relations, ensuring **data integrity** and eliminating **data redundancy**.
 - Data Integrity: Database is consistent and satisfies all constraint rules.
 - Data Redundancy: If data can be found in two places in a single database (direct redundancy) or calculated using data from different parts of the database (indirect redundancy) then redundancy exists.
- Normalisation should remove redundancy, but not at the expense of data integrity.

- If redundancy exists then this can cause problems during normal database operations:
 - When data is inserted into the database, the data must be duplicated wherever redundant versions of that data exists.
 - When data is updated, all redundant data must be updated at the same time.

- An **integrity constraint** is a rule that restricts the values that may be present in the database.
- **Entity integrity**: the rows (or tuples) in a relation represent entities, and each one must be uniquely identified.
 - We must have a **primary key** that must have a unique non-null value for every row
- **Referential integrity**: Involves the foreign keys
 - Foreign keys tie the relations together, so it is important that the links are correct.
 - Every foreign key must either be null, or its value must be the actual value of a key in another relation

- There are many benefits to normalised database
 - Saves space
 - Prevents anomalies
 - Avoid NULL values

- While many modern computers have a lot of disk space, this can still be of benefit
- If we can make our database small enough to fit in memory, queries will be much quicker (no disk reads)
- This is also important for apps on devices with less memory, like phones

Example

employees

<u>Eid</u>	name	Bdate	salary	<i>Dnum</i>
1234	Sean Russell	1985-07-21	50000	20
4567	Jamie Heaslip	1982-05-04	47000	10
6542	Leo Cullen	1978-01-07	45000	10
1238	Brendan Macken	1990-12-03	25000	20
1555	Sean O'Brien	1986-04-09	50000	30

departments

<u>Dnum</u>	Dname	<i>Dmgr_id</i>
10	Training	4567
20	Design	1238
30	Implementation	1555

Example - Alternate

employees_departments

<u>Eid</u>	name	Bdate	salary	Dnum	Dname	Dmgr_id
1234	Sean Russell	1985-07-21	50000	20	Design	1238
4567	Jamie Heaslip	1982-05-04	47000	10	Training	4567
6542	Leo Cullen	1978-01-07	45000	10	Training	4567
1238	Brendan Macken	1990-12-03	25000	20	Design	1238
1555	Sean O'Brien	1986-04-09	50000	30	Implementation	1555

- This configuration of the table contains the same information as the previous
- But we are recording the same pieces of information in several places



- One of the biggest concerns with a database is integrity
- Anomalies in that database are where our data is not correct
- This can happen in a number of ways
 - When inserting data
 - When deleting data
 - When modifying data

Insertion Anomalies

- Every time we insert new data, we must make sure it is consistent with the data already in the table
- For example, when inserting a new employee into department 10, we must make sure that the department name and manager are correct

employees_departments

Eid	name	Bdate	salary	Dnum	Dname	Dmgr_id
1234	Sean Russell	1985-07-21	50000	10	Design	1238
4444	Paul Heaslip	1983-05-04	44000	10	Training	4567

- With this data it is impossible to know what the name of department 10 is and who is the manager



Deletion Anomalies

- Because we have represented the employees and departments together deleting data can cause problems
- If we delete employee 1555, we will lose all records of department number 30

employees_departments

<u>Eid</u>	name	Bdate	salary	Dnum	Dname	Dmgr_id
1234	Sean Russell	1985-07-21	50000	20	Design	1238
4567	Jamie Heaslip	1982-05-04	47000	10	Training	4567
6542	Leo Cullen	1978-01-07	45000	10	Training	4567
1238	Brendan Macken	1990-12-03	25000	20	Design	1238
1555	Sean O'Brien	1986-04-09	50000	30	Implementation	1555

Modification Anomalies

- Because we have represented the employees and departments together changing data can be more complicated
- If we want to rename one of the departments, we must rename it in every row where that department is used
- If we miss some rows, our data will be inconsistent

- In some database designs, we may create tables that have many attributes that do not apply to all rows
- This means that we end up with many rows with NULL values in those attributes
- This can cause problems with wasted space and may make some join operations more difficult to understand

employees

Eid	name	Bdate	salary	Dnum	typing_speed	eng_type
1234	Sean Russell	1985-07-21	50000	20	140	NULL
4567	Jamie Heaslip	1982-05-04	47000	10	NULL	Civil
6542	Leo Cullen	1978-01-07	45000	10	NULL	NULL
1238	Brendan Macken	1990-12-03	25000	20	NULL	NULL
1555	Sean O'Brien	1986-04-09	50000	30	NULL	Demolition

Table of Contents

1 What is Normalisation? and Why Normalise Databases?

2 Functional Dependencies

- Formal Definition
- Identifying Functional Dependencies
- Partial Dependencies

3 Normal Forms

- If $R = \{A_1, A_2, \dots, A_n\}$ is a relation that describes all of the tables in our database
- Then a functional dependency, denoted by $X \rightarrow Y$, between two sets of attributes X and Y that are subsets of R specifies a constraint on the possible tuples that can form a relation state r of R .
- The constraint is that, for any two tuples t_1 and t_2 in r that have $t_1[X] = t_2[X]$, they must also have $t_1[Y] = t_2[Y]$

Explanation

- A functional dependency means that the values for the attributes on the right depend on the attributes on the left
- So if we have the following functional dependency:
 $Eid \rightarrow name$
 - Then whenever we see particular id, we will know what name is on that row
 - The opposite is not true (two employees may have the same name)
- Functional dependencies always involve keys on the left (candidate and primary keys)



Example

employees

<u>Eid</u>	name	Bdate	salary	Dnum
1234	Sean Russell	1985-07-21	50000	20
4567	Jamie Heaslip	1982-05-04	47000	10
6542	Leo Cullen	1978-01-07	45000	10
1238	Brendan Macken	1990-12-03	25000	20
1555	Sean O'Brien	1986-04-09	50000	30

departments

<u>Dnum</u>	Dname	Dmgr_id
10	Training	4567
20	Design	1238
30	Implementation	1555

- $Eid \rightarrow \{name, Bdate, salary, Dnum\}$
- $Dnum \rightarrow \{Dname, Dmgr_id\}$



Understanding the Data

- You cannot find the functional dependencies in a database by looking at the data
- You need to understand the meaning of the data
- You can prove that something is not a functional dependency by finding a single example that contradicts it

Keys

- Identifying functional dependencies comes down to keys
- We describe keys using several terms
 - superkey
 - candidate key
 - primary key
 - surrogate key
- Lets look at an example using the following relation and assumptions:
 - Table: `students(UCD_id, name, major , Bdate)`
 - Assumptions: `UCD_id` is unique and the combination of `name` and `Bdate` is also unique



Superkeys

- A **superkey** is a **set** of attributes that uniquely identify every row in a database
- From the example, we can identify the following superkeys:
 - $\{UCD_id, name, major, Bdate\}$
 - $\{UCD_id, name, major\}$
 - $\{UCD_id, name, Bdate\}$
 - $\{UCD_id, major, Bdate\}$
 - $\{UCD_id, name\}$
 - $\{UCD_id, major\}$
 - $\{UCD_id, Bdate\}$
 - $\{UCD_id\}$
 - $\{name, major, Bdate\}$
 - $\{name, Bdate\}$



Keys

- A **key**, is a superkey with one additional property, removing any attribute from the set will make it no longer a superkey
- This means if we remove any of the attributes, the set is no longer unique and cannot be used to identify each row
- The main difference between keys and super keys is that keys are minimal

Keys and Superkeys

- From our previous example:

1	$\{UCD_id, name, major, Bdate\}$	superkey
2	$\{UCD_id, name, major\}$	superkey
3	$\{UCD_id, name, Bdate\}$	superkey
4	$\{UCD_id, major, Bdate\}$	superkey
5	$\{UCD_id, name\}$	superkey
6	$\{UCD_id, major\}$	superkey
7	$\{UCD_id, Bdate\}$	superkey
8	$\{UCD_id\}$	key
9	$\{name, major, Bdate\}$	superkey
10	$\{name, Bdate\}$	key

- If we remove an attribute from 8 it is empty
- If we remove any attribute from 10, it can no longer identify the rows

Candidate and Primary Keys

- If a table in our database has more than one key, each key is called a **candidate key**
- One of the keys is chosen to be the **primary key**
 - How you choose does not matter
 - In the previous example, I would choose UCD_id, because it is only made up of a single attribute
 - This could make join operations easier
- Every table **must** have a primary key



Surrogate Key

- Every table **must** have a primary key
 - If you can't identify any candidate keys, we can use all of the attributes as the primary key
- Alternatively, we can make up a primary key by adding a new attribute
 - This is called a **surrogate key**
 - Usually, automatically generated integers are used
 - This is how you get your student numbers

Functional Dependencies and Keys

- There is always a functional dependency from each candidate key to all other information in the table
- `students(UCD_id, name, major , Bdate)`
 - $\{UCD_id\} \rightarrow \{name, major, Bdate\}$
 - $\{name, Bdate\} \rightarrow \{UCD_id, major\}$

Functional Dependency Types

- For the some normal forms, we must have an understanding of the different types of functional dependencies
- There are two types of functional dependency
 - Full functional dependency
 - Partial functional dependency
- Partial functional dependency is something we need to check for when our primary key contains multiple attributes
- A primary key with only a single attribute is automatically a full functional dependency



Full Functional Dependency

- A functional dependency $X \rightarrow Y$ is a **full** functional dependency if we **cannot** remove any attributes from X and have it stay as a functional dependency
- Consider the relation:
project_hours(E_id , $Pnum$, $hours$)
- We have the functional dependency $\{E_id, Pnum\} \rightarrow \{hours\}$
- This is a full functional dependency because neither $\{E_id\} \rightarrow \{hours\}$ or $\{Pnum\} \rightarrow \{hours\}$ hold as functional dependencies

Partial Functional Dependency

- A functional dependency $X \rightarrow Y$ is a **partial** functional dependency if we **can** remove an attribute from X and have it stay as a functional dependency
- Consider the relation:
 $emp_proj(\underline{E_id}, \underline{Pnum}, Ename, Pname, Plocation)$
- We have the functional dependencies
 $\{E_id\} \rightarrow \{Ename\}$, $\{E_id, Pnum\} \rightarrow \{hours\}$ and
 $\{Pnum\} \rightarrow \{Pname, Plocation\}$
- Lets look at the dependencies on the key we have
 $\{E_id, Pnum\} \rightarrow$
 $\{hours, Ename, Pname, Plocation\}$



Partial Functional Dependency

- So we are checking if $\{E_id, Pnum\} \rightarrow \{hours, Ename, Pname, Plocation\}$ is a partial functional dependency
 - $\{E_id, Pnum\} \rightarrow \{hours\}$ - Cannot remove part of key
 - $\{E_id, Pnum\} \rightarrow \{Ename\}$ - Can remove Pnum
 - $\{E_id, Pnum\} \rightarrow \{Pname\}$ - Can remove E_id
 - $\{E_id, Pnum\} \rightarrow \{Plocation\}$ - Can remove E_id
- This is a **partial** functional dependency because $\{E_id\} \rightarrow \{Ename\}$ and $\{Pnum\} \rightarrow \{Pname, Plocation\}$ both hold as functional dependencies

Table of Contents

1 What is Normalisation? and Why Normalise Databases?

2 Functional Dependencies

3 Normal Forms

- First Normal Form (1NF)
- Second Normal Form (2NF)
- Third Normal Form (3NF)
- Boyce-Codd Normal Form (BCNF)
- Why BCNF?

- The data in a database can be considered to be in one of a number of **normal forms**.
- Basically, the normal form of the data indicates how much redundancy is in the data.
- The normal forms have a strict ordering
 - 1st Normal Form
 - 2nd Normal Form
 - 3rd Normal Form
 - Boyce-Codd Normal Form (BCNF)
 - 4th Normal Form
 - 5th Normal Form



- First Normal Form (1NF) deals with the shape of the record type.
- A relation is in 1NF if, and only if, it contains no repeating attributes or groups of attributes
- $\{UCD_id, name, Bdate, (subject, grade)\}$
 - For this example, we have a single row for each student, but can have multiple pairs of subject and grade
 - This is not usually a feature of modern database systems

students

<u>UCD_id</u>	name	Bdate	subject	grade
12345	Smith.J	1985-07-21	Java	B
			Soft Eng	C
			Databases	A
23456	White.A	1990-12-03	Java	D
			Soft Eng	B
34567	Moore.T	1986-04-09	Databases	A
			Soft Eng	B
			Networks	C
45678	Smith.J	1998-11-02	Soft Eng	C

- This table with the repeating group is not in 1NF.
- To remove the repeating group, one of two things can be done:
 - "Flatten" the relation (fill in the empty attribute spaces) and extend the key, or
 - "Decompose" the relation (divide into multiple relations)



Flatten

students

<u>UCD_id</u>	name	Bdate	<u>subject</u>	grade
12345	Smith.J	1985-07-21	Java	B
12345	Smith.J	1985-07-21	Soft Eng	C
12345	Smith.J	1985-07-21	Databases	A
23456	White.A	1990-12-03	Java	D
23456	White.A	1990-12-03	Soft Eng	B
34567	Moore.T	1986-04-09	Databases	A
34567	Moore.T	1986-04-09	Soft Eng	B
34567	Moore.T	1986-04-09	Networks	C
45678	Smith.J	1998-11-02	Soft Eng	C

- The empty spaces are filled in with the data from the row
 - But this would mean we have multiple rows with the same primary key
- To correct this, the key is changed to $\{UCD_id, subject\}$
- This is now in **1NF**



Problems with Flattened Tables

- Redundant Data

- We are wasting storage space by storing the name and date of birth of each student many times

- Anomalies

Insertion : Because the key is now UCD_id and subject, we cannot insert a new student until they have completed a course

Update : To change the name of a student, we must update it in every row for that student in the table

Deletion : If we delete all rows containing data about a subject, we will also delete any student who has only completed that subject



Decompose

- The alternative approach is to split the table into two relations: one for the repeating groups and one for the non-repeating groups.
- The primary key for the original relation is included in both of the new relations.
- We can return to the original table by using a JOIN operation on these relations

Decomposed

students

<u>UCD_id</u>	name	Bdate
12345	Smith.J	1985-07-21
23456	White.A	1990-12-03
34567	Moore.T	1986-04-09
45678	Smith.J	1998-11-02

grades

<u>UCD_id</u>	subject	grade
12345	Java	B
12345	Soft Eng	C
12345	Databases	A
23456	Java	D
23456	Soft Eng	B
34567	Databases	A
34567	Soft Eng	B
34567	Networks	C
45678	Soft Eng	C

Results with Decomposed Tables

- There is no redundancy

- Anomalies

Insertion : We can insert a new student before they have completed a course

Update : We only need to update a single row to change the name of a student

Deletion : We can delete all of the rows for a subjects grade, and we will not delete any students



- A relation is in 2NF if, and only if, it is in 1NF and every non-key attribute is **fully** functionally dependent on the **whole** key
- The relation must be in 1NF and all non-key attributes must depend on the whole key, not just part of it.
 - In other words: there must be no partial key dependencies.
- The problem arises when there is a compound key, e.g. in the Grades relation: UCD_id, subject
- In this case it is possible for non-key attributes to depend on only part of the key (i.e. on only one of the key attributes).

2NF Example 1

- Consider the flattened student relation:
students(UCD_id, *name*, *Bdate*, *subject*, *grade*)
- There are no repeating groups: already in 1NF.
- However, there is a compound primary key, so we must check that the non-key attributes depend on the whole key.
- There are three non-key attributes: *name*, *Bdate* and *grade*

2NF Example 1

- So we are checking if $\{UCD_id, subject\} \rightarrow \{name, Bdate, grade\}$ is a partial functional dependency
 - $\{UCD_id, subject\} \rightarrow \{name\}$ - Can remove subject
 - $\{UCD_id, subject\} \rightarrow \{Bdate\}$ - Can remove subject
 - $\{UCD_id, subject\} \rightarrow \{grade\}$ - Cannot remove part of key
- This relation is not in 2NF. It appears to be two tables squashed into one.
- Solution: decompose the relation

Heath's Theorem

- Heath's Theorem allows us to perform non-loss decomposition.
- If R is a relation made up of sets of attributes $\{A, B, C\}$.
- If $A \rightarrow B$, then R can be non-loss decomposed into:
 - $\{A, B\}$
 - $\{A, C\}$
- R can be created again using $\{A, B\} \text{ JOIN } \{A, C\}$



Heath's Theorem Steps

- 1 Create a new relation that contains all of the attributes that are solely dependent on UCD_id (UCD_id is the primary key of the new relation).
- 2 Create a new relation that contains all the attributes that are solely dependent on subject (subject is the primary key of the new relation).
 - In this example, there are no attributes that depend only on subject.
- 3 Create a new relation with all the attributes that depend on both UCD_id and subject.
 - The primary key is UCD_id, Subject



Result

students(*UCD_id*, *name*, *Bdate*)

grades(*UCD_id*, *subject*, *grade*)

- All attributes in each relation are fully functionally dependent on the primary key.
 - $\{UCD_id\} \rightarrow \{name, Bdate\}$
 - $\{UCD_id, subject\} \rightarrow \{grade\}$
- Both relations are now in 2NF

- 3NF is an even stricter normal form that removes almost all redundant data.
- A relation is in 3NF if, and only if, it is in 2NF and there are no transitive functional dependencies.
- Transitive functional dependencies are when one non-key attribute is functionally dependent on another non-key attribute.
- By definition, a transitive functional dependency can only happen if there is more than one non-key attribute

3NF Example

- Consider the following table
projects(*Pnum*, *manager*, *address*)

projects

<u>Pnum</u>	manager	address
P1	Black. B	32 High St
P2	Smith.J	11 New St
P3	Black. B	32 High St
P4	Black. B	32 High St

- Projects has more than one non-key field, so we must check for transitive dependency

3NF Example

- In this example, we are told that address depends on the value of manager.
 - If we know a project's manager, we can find the address
- This gives us a functional dependency of $\{manager\} \rightarrow \{address\}$
- In this case, address is transitively dependent on manager.
- The primary key is Pnum, but the functional dependency makes no reference to this key

3NF Example

- Data redundancy can come from this:
 - We duplicate the address if a manager is in charge of more than one project.
 - Causes problems if we have to change the address, because it must be changed in several places.
- Solution: decompose:
 - Create two relations: one with the transitive dependency in it, and another for all of the remaining attributes.
 - Split projects into projects and managers
 - In the projects relation, we keep the same primary key: Pnum
 - In the managers relation we use the left side of the functional dependency as the primary key: manager



projects

<u>Pnum</u>	<i>manager</i>
P1	Black. B
P2	Smith.J
P3	Black. B
P4	Black. B

managers

<u>manager</u>	address
Black. B	32 High St
Smith.J	11 New St

- Now we store the address of each manager only once
- If we need to know a manager's address, we can look it up in the managers table
- The manager attribute in the Projects relation is now a foreign key.
- These relations are now in 3NF

- Boyce-Codd Normal Form (BCNF) is named after Raymond Boyce and Edgar Codd who developed it in 1974 to fix some anomalies not addressed in 3NF
 - It is sometimes called 3.5NF
- A relation in 3NF is usually also in BCNF
- Only if a relation in 3NF has **overlapping candidate keys** can it not be in BCNF

Overlapping Candidate Keys

- Overlapping candidate keys means that we have composite candidate keys with at least one attribute in common
 - E.G If we have relation containing prices for products:
products(Pid, start_date, end_date, price)
 - Then we would have candidate keys of $\{Pid, start_date\}$ and $\{Pid, end_date\}$
- Here there is an overlap, because *Pid* is in both candidate keys

BCNF

- BCNF is based on the concept of a **determinant**.
 - A determinant is any attribute (or group of attributes) that some other attribute is fully functionally dependent on
 - I.e. the left side of a functional dependency
- A relation is in BCNF if, and only if, every determinant is a candidate key.

BCNF: Theory

- Consider the following relation and dependencies:
- $R(a, b, c, d)$
 - $\{a, c\} \rightarrow \{b, d\}$
 - $\{a, d\} \rightarrow \{b\}$
- To be in BCNF, all determinants must be candidate keys. Let's look at each in turn
- $\{a, c\} \rightarrow \{b, d\}$
 - a and c together can determine both b and d
 - Therefore $\{a, c\}$ is a candidate key for this relation
- $\{a, d\} \rightarrow \{b\}$
 - $\{a, d\}$ does not determine c, so it can't be a candidate key.
 - Therefore R is not in BCNF.

BCNF: Hospital Example

PatNum	PatName	AppSlot	Time	Doctor
1	Eamonn	0	09:00	Octopus
2	Eoin	0	09:00	Evil
3	Arnold	1	10:00	Octopus
4	Stephen	0	13:00	Evil
5	Patricia	1	14:00	Octopus

- Extra information:

- Every patient has a unique patient number.
- Patients with names beginning with a letter before 'P' get morning appointments.
- The Appointment slots start at 0 for the first appointment of the morning or afternoon, 1 for the second and so on.



BCNF: Hospital Example

- *Appointments*(*PatNum*, *PatName*, *AppSlot*, *Time*, *Doctor*)
- We are given some functional dependencies (mostly based on the extra information):
 - $\{PatNum\} \rightarrow \{PatName\}$
 - $\{PatNum, AppSlot\} \rightarrow \{Time, Doctor\}$
 - $\{Time\} \rightarrow \{AppSlot\}$
- We have two options for selecting a primary key:
 - *Appointments*(*PatNum*, *PatName*, *AppSlot*, *Time*, *Doctor*): example A
 - *Appointments*(*PatNum*, *PatName*, *AppSlot*, *Time*, *Doctor*): example B

BCNF: Example A

- *Appointments*(*PatNum*, *PatName*, *AppSlot*, *Time*, *Doctor*)
- No repeating groups, so in 1NF.
- 2NF – eliminate partial key dependencies:
 - *Appointments*(*PatNum*, *AppSlot*, *time*, *Doctor*)
 - *Patients*(*PatNum*, *PatName*)
- 3NF – no transient dependencies so it's already in 3NF
- Now try BCNF

BCNF: Example A

- BCNF: Every determinant must be a candidate key
 - *Appointments*(*PatNum*, *AppSlot*, *Time*, *Doctor*)
 - *Patients*(*PatNum*, *PatName*)
- For each functional dependency, we will look at each table and ask the following questions:
 - 1 Are the attributes present in the table?
 - 2 Is the determinant the key of the table?

BCNF: Example A

- BCNF: Every determinant must be a candidate key
 - *Appointments*(*PatNum*, *AppSlot*, *Time*, *Doctor*)
 - *Patients*(*PatNum*, *PatName*)
- We have 3 functional dependencies:
 - 1 $\{PatNum\} \rightarrow \{PatName\}$
 - 2 $\{PatNum, AppSlot\} \rightarrow \{Time, Doctor\}$
 - 3 $\{Time\} \rightarrow \{AppSlot\}$

1. $\{PatNum\} \rightarrow \{PatName\}$

- *Patients*(*PatNum*, *PatName*)

- ① All attributes are present, so this is relevant

- ② $\{PatNum\}$ is the key of the table (and also a candidate key), so this is OK

- *Appointments*(*PatNum*, *AppSlot*, *Time*, *Doctor*)

- ① *PatNum* is present in *Appointments*, but not *PatName*, so this is not relevant

2. $\{PatNum, AppSlot\} \rightarrow \{Time, Doctor\}$

- *Patients*(*PatNum*, *PatName*)

- ① Not all attributes are present, so this is not relevant

- *Appointments*(*PatNum*, *AppSlot*, *Time*, *Doctor*)

- ① All attributes are present, so this is relevant

- ② $\{PatNum, AppSlot\}$ is the key of the table (and also a candidate key), so this is OK

3. $\{Time\} \rightarrow \{AppSlot\}$

- $Patients(\underline{PatNum}, PatName)$
 - ① No attributes are present, so this is not relevant
- $Appointments(\underline{PatNum}, \underline{AppSlot}, Time, Doctor)$
 - ① All attributes are present, so this is relevant
 - ② $\{Time\}$ is not a candidate key of the table,
- This is not in BCNF

Rewrite to BCNF

- Original 3NF version:
 - *Appointments*(*PatNum*, *AppSlot*, *Time*, *Doctor*)
 - *Patients*(*PatNum*, *PatName*)
- Rewritten to BCNF
 - *Appointments*(*PatNum*, *Time*, *Doctor*)
 - *Patients*(*PatNum*, *PatName*)
 - *Slots*(*Time*, *AppSlot*)
- “Time” is enough to work out the appointment slot of a patient
- Now BCNF is satisfied, and the final relations shown are in BCNF

BCNF: Example B

- *Appointments*(*PatNum*, *PatName*, *AppSlot*, *Time*, *Doctor*)
- No repeating groups, so in 1NF.
- 2NF – eliminate partial key dependencies:
 - *Appointments*(*PatNum*, *Time*, *Doctor*)
 - *Patients*(*PatNum*, *PatName*)
 - *Slots*(*Time*, *AppSlot*)
- 3NF – no transient dependencies so it's already in 3NF
- Now try BCNF

BCNF: Example B

- BCNF: Every determinant must be a candidate key
 - *Appointments*(*PatNum*, *Time*, *Doctor*)
 - *Patients*(*PatNum*, *PatName*)
 - *Slots*(*Time*, *AppSlot*)
- We have 3 functional dependencies:
 - 1 $\{PatNum\} \rightarrow \{PatName\}$
 - 2 $\{PatNo, AppSlot\} \rightarrow \{Time, Doctor\}$
 - 3 $\{Time\} \rightarrow \{AppSlot\}$

1. $\{PatNum\} \rightarrow \{PatName\}$

- *Appointments*(*PatNum*, *Time*, *Doctor*)
 - ① *PatNum* is present in *Appointments*, but not *PatName*, so this is not relevant
- *Patients*(*PatNum*, *PatName*)
 - ① All attributes are present, so this is relevant
 - ② $\{PatNum\}$ is the key of the table (and also a candidate key), so this is OK
- *Slots*(*Time*, *AppSlot*)
 - ① No attributes are present, so this is not relevant

2. $\{PatNum, AppSlot\} \rightarrow \{Time, Doctor\}$

- *Appointments*(*PatNum*, *Time*, *Doctor*)
 - ① *PatNum*, *Time* and *Doctor* are present, but not *AppSlot*, so this is not relevant
- *Patients*(*PatNum*, *PatName*)
 - ① *PatNum* is present in, but none of the others, so this is not relevant
- *Slots*(*Time*, *AppSlot*)
 - ① *Time* and *AppSlot* are present, but not the others, so this is not relevant

3. $\{Time\} \rightarrow \{AppSlot\}$

- *Appointments*(PatNum, Time, Doctor)
 - ① Time is present, but not AppSlot, so this is not relevant
- *Patients*(PatNum, PatName)
 - ① No attributes are present, so this is not relevant
- *Slots*(Time, AppSlot)
 - ① Time and AppSlot are present, so this is relevant
 - ② $\{Time\}$ is the key of the table (and also a candidate key), so this is OK
- Relations are in BCNF

- It can take a bit of time and effort to get a database design into BCNF
- So what is the benefit of having done so?
- We will look at an example that shows some of the problems that BCNF overcomes

students

<u>Snum</u>	<u>major</u>	supervisor
123	Physics	Einstein
123	Music	Mozart
456	Biology	Darwin
789	Physics	Bohr
999	Physics	Einstein

- No repeating groups, so it's in 1NF
- No partial key dependencies, so it's in 2NF
- There's only one non-key attribute (Supervisor) so it must be in 3NF

students

<u>Snum</u>	<u>major</u>	supervisor
123	Physics	Einstein
123	Music	Mozart
456	Biology	Darwin
789	Physics	Bohr
999	Physics	Einstein

- We have the following functional dependencies:
 - $\{Snum, major\} \rightarrow \{supervisor\}$
 - $\{supervisor\} \rightarrow \{Major\}$

students

<u>Snum</u>	<u>major</u>	supervisor
123	Physics	Einstein
123	Music	Mozart
456	Biology	Darwin
789	Physics	Bohr
999	Physics	Einstein

- If the record for student 456 is deleted we lose not only information about that student, but also the fact that Darwin advises in Biology.
- We cannot record the fact that Watson can advise on Computing until we have a student doing a project on Computing that has Watson as an advisor.



- In BCNF we have two tables, and these problems are eliminated:

supervisors

<u>Snum</u>	<u>supervisor</u>
123	Einstein
123	Mozart
456	Darwin
789	Bohr
999	Einstein

majors

<u>supervisor</u>	major
Einstein	Physics
Mozart	Music
Darwin	Biology
Bohr	Physics

Summary

- A relation is in 1NF if, and only if, it contains no repeating groups.
- A relation is in 2NF if, and only if, it is in 1NF and every non-key attribute is fully functionally dependent on the whole key.
- A relation is in 3NF if, and only if, it is in 2NF and has no transitive functional dependencies.
- A relation is in BCNF if, and only if, it is in 3NF and every determinant is a candidate key.

