

# Operating Systems

Dr. Vivek Nallur ([vivek.nallur@ucd.ie](mailto:vivek.nallur@ucd.ie))

# First things, first – enrol on Moodle

---

- ▶ Register yourself [ucdconnect account] on Moodle:

<https://csmoodle.ucd.ie/moodle/login/index.php>

- ▶ Search for: COMP 2006J
- ▶ Find a module entitled: COMP2006J Operating Systems 2017-2018

Group Number	Enrollment Key
153711	<a href="#">COMP2006JVN153711</a>
153712	<a href="#">COMP2006JVN153712</a>
153713	<a href="#">COMP2006JVN153713</a>

Group Number	Enrollment Key
163721	<a href="#">COMP2006JVN163721</a>
163722	<a href="#">COMP2006JVN163722</a>
163723	<a href="#">COMP2006JVN163723</a>





# Plagiarism & UCD Computer Science

---

- **Plagiarism is a serious academic offence**
  - [Student Code, sections 6.2 & 6.3] or [UCD Registry Plagiarism Policy] or [CS Plagiarism policy and procedures]
- Our staff and demonstrators are **proactive** in looking for possible plagiarism in all submitted work
- Suspected plagiarism is reported to the CS Plagiarism subcommittee for investigation
  - Usually includes an interview with student(s) involved
  - 1st offence: **usually** 0 or NG in the affected components
  - 2nd offence: may be referred to the **University disciplinary committee**
- Student who enables plagiarism is equally responsible

[http://www.ucd.ie/registry/academicsecretariat/docs/plagiarism\\_po.pdf](http://www.ucd.ie/registry/academicsecretariat/docs/plagiarism_po.pdf)

[http://www.ucd.ie/registry/academicsecretariat/docs/student\\_code.pdf](http://www.ucd.ie/registry/academicsecretariat/docs/student_code.pdf)

<http://libguides.ucd.ie/academicintegrity>

# Operating Systems Structure

## Operating System Architecture

# Operating System Structure

---

- ▶ In order to understand any operating system structure, we must consider its **components** and their organisation
- ▶ What are the essential components?
- ▶ How do they relate to each other?



# Operating System Structure

---

- ▶ It is important to remember that
- ▶ The concepts we study will exist in some form in every operating system
  - ▶ But they will be implemented in different ways
- ▶ Divisions between components are not always clearly defined



# Operating System Components

---

- ▶ **Kernel**: software containing the core OS components; it may typically include:
  - ▶ Memory Manager
    - ▶ Provides efficient memory allocation and deallocation of memory
  - ▶ I/O manager
    - ▶ Handles input and output requests from and to hardware devices (through device drivers)



# Operating System Components

---

- ▶ **Inter-process communication (IPC) manager**
  - ▶ Provides communication between different processes (programs in execution)
- ▶ **Process Manager (scheduler)**
  - ▶ Handles what is executed when and where (if more than one CPU)





# Operating System Components

---

- ▶ A OS kernel may consist of many more components:
  - ▶ System service routines
  - ▶ File System (FS) manager
  - ▶ Error handling systems
  - ▶ Accounting systems
  - ▶ System programs
  - ▶ And many more



# Operating System Interface

---

- ▶ Original OS interfaces were very simple and called **Command Line Interface** (CLI) or **Command Interpreter** (CI)
- ▶ This is the like the command prompt on windows and terminal or shell on Mac/Linux
- ▶ The user types a command and the CI executes it



# Operating System Interface

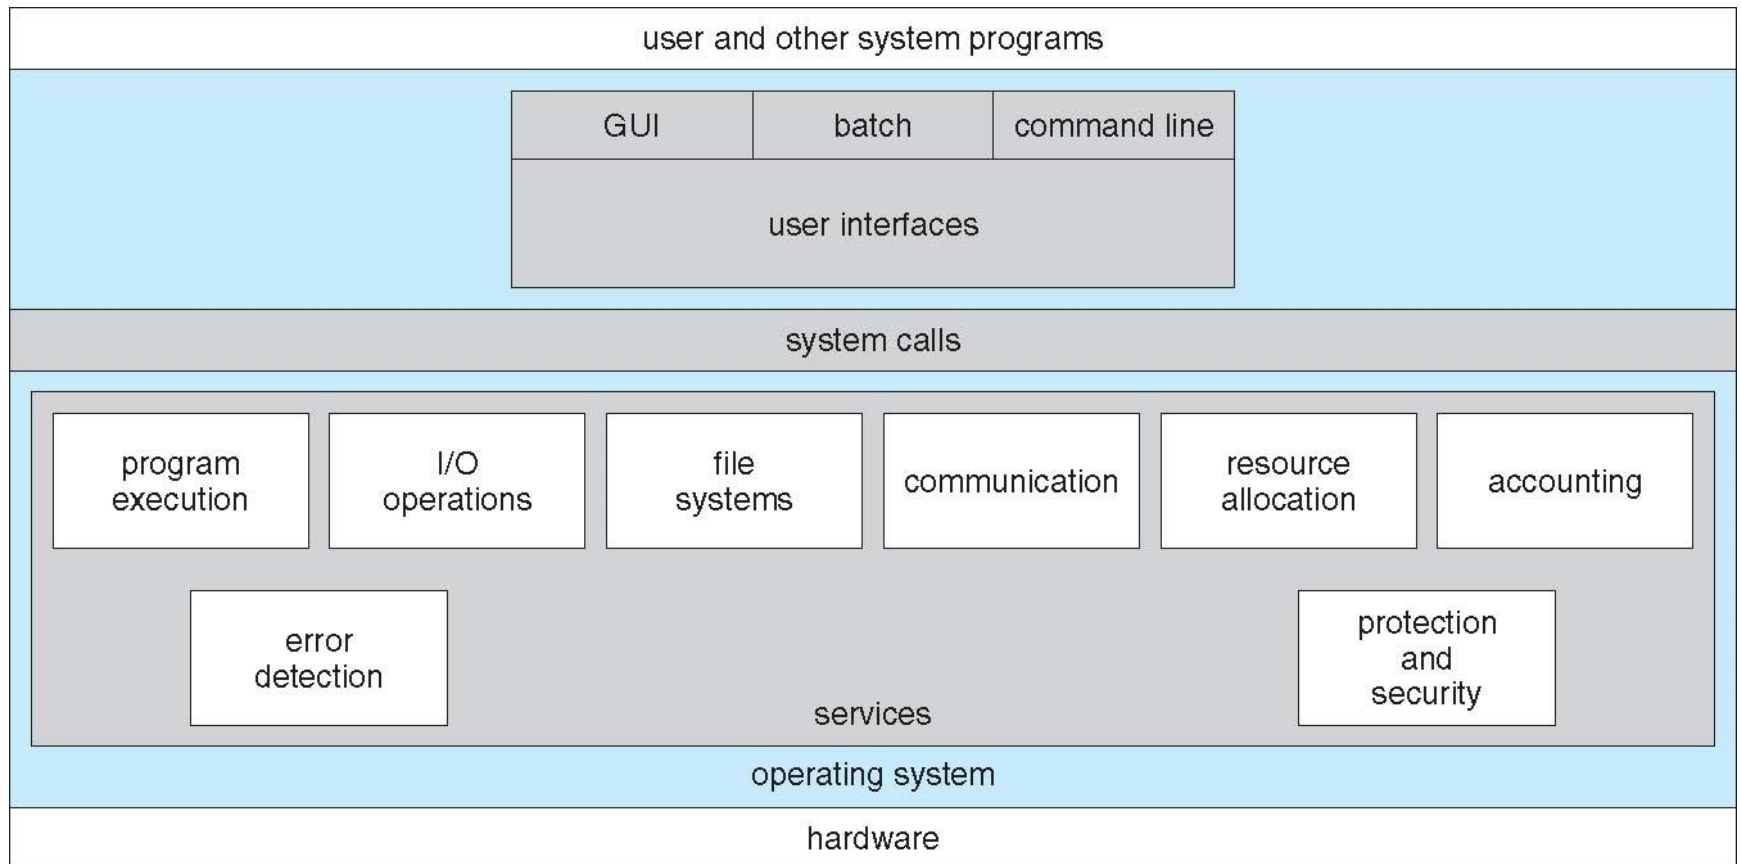
---

- ▶ Later systems used a more user friendly Graphical User Interface (GUI)
- ▶ Based on Desktop idea
  - ▶ Usually mouse, keyboard, and monitor
  - ▶ WIMP, (windows, icons, menus, pointing)
  - ▶ Icons represent files, programs, actions, etc
- ▶ Used on almost all systems today



# Example system

---



# OS Structure Issues:

---

- ▶ How are all of these components organised?
- ▶ What are the entities involved and where do they exist?
- ▶ how do these entities cooperate?



# Operating System Goals

---

- ▶ When we design an operating system we want it to be:
  - ▶ Efficient – (High **throughput**)
  - ▶ Interactive
  - ▶ **Robust** – (Fault tolerant & reliable)
  - ▶ Secure
  - ▶ Scalable
  - ▶ Extensible
  - ▶ Portable



# Monolithic Architecture

---

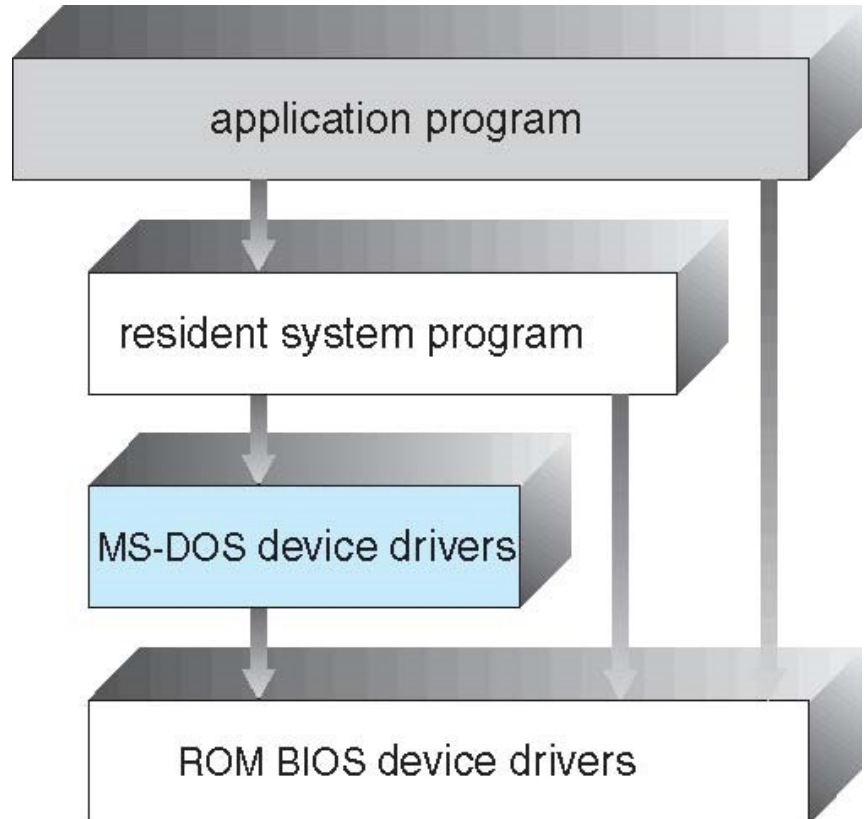
- ▶ Traditionally, systems were built around **monolithic** kernels
  - ▶ every OS component is contained in the kernel
  - ▶ any component can directly communicate with any other (by means of function calls)
  - ▶ due to this they tend to be highly efficient (performance)



# Example : MS DOS

---

- ▶ Written to provide the most functionality in the least space
- ▶ Not divided into modules
- ▶ Interfaces and levels of functionality are not well separated





# Problems with Monolithic Structure

---

- ▶ Because they are unstructured they are hard to understand modify and maintain
- ▶ Susceptible to damage from errant or malicious code, as all kernel code runs with unrestricted access to the system



# Layered Structure

---

- ▶ To try and solve the problems with the monolithic structure, the **layered structure** was developed
- ▶ Components are grouped into layers that perform similar functions



# Advantages of Layered Structure

---


- ▶ Designing the system as a number of modules gives the system structure and consistency
- ▶ This allows easier debugging, modification and reuse



# Layered Architecture

---

- ▶ Each layer communicates only with layers immediately above and below it
  - ▶ each layer is a virtual machine to the layer above
  - ▶ a higher layer provides a higher-level virtual machine



# First Layers-based OS

---

- ▶ Layering was first used in Dijkstra's THE OS (1968)
- ▶ Each layer “sees” a logical machine provided by lower layers



# The Layers

---

layer 4	User Programs
layer 3	I/ O Management
layer 2	Console Device (commands), IPC
layer 1	Memory Management
layer 0	CPU Scheduling (multiprogramming)
	Hardware



# First Layers-based OS

---

- ▶ Each layer “sees” a logical machine provided by lower layers
  - ▶ layer 4 (user space) sees virtual I/O drivers
  - ▶ layer 3 sees virtual console
  - ▶ layer 2 sees virtual memory
  - ▶ layer 1 sees virtual processors
- ▶ Based on a static set of cooperating processes
- ▶ Each process can be tested and verified independently



# Problems with layering

---

- ▶ **Appropriate definition of layers is difficult**
  - ▶ A layer is implemented using only those operations provided by lower-level layers
  - ▶ A real system structure is often **more complex** than the strict hierarchy required by layering





# Problems with layering

---

- ▶ The secondary memory (disk) driver would normally placed **above** the CPU scheduler (because an I/O wait may trigger a CPU rescheduling operation)
- ▶ However, in a large system the CPU scheduler may need more memory than can fit in memory: parts of the memory can be swapped to disk (virtual memory), and then the secondary memory driver should be **below** the CPU scheduler
- ▶ Both things cannot be achieved at the same time; therefore the layered approach is not flexible



# Problems with Layering

---

- ▶ **Performance issues**

- ▶ Processes' requests might pass through many layers before completion (layer crossing)
- ▶ System throughput can be lower than in monolithic kernels



# Problems with Layering

---


- ▶ Still susceptible to malicious/errant code if all layers can have unrestricted access to the system
  - ▶ As we will see later, this can only be avoided through hardware
- ▶ As a consequence, (imperfect) layers are often used for convenience in system design, but any OS implementation cannot be purely layered



# Example: MS DOS

---

- ▶ Structure is somewhat layered
- ▶ Interfaces and levels of functionality are not well separated



# Microkernel Architecture

---

- ▶ The **microkernel** ( $\mu$ -kernel) architecture was designed to minimise the services offered by the kernel
- ▶ This was an attempt to keep it small and scalable



# Microkernel Architecture

---

- ▶ There is no agreement about minimal set of services inside the microkernel
- ▶ At least: minimal process and memory management capabilities, plus inter-process communications
- ▶ Services such as networking and file system tend to run nonprivileged at the user process level



# Benefits of Microkernel

---

- ▶ **Modularity**
- ▶ It promotes uniform interfaces
- ▶ **Distributed systems support:**
  - ▶ modules communicate through the microkernel, even through a network



# Benefits of Microkernel

---

- ▶ Reliability
- ▶ Scalability
- ▶ Portability
- ▶ Easy to extend and customise





# Disadvantages of Microkernel

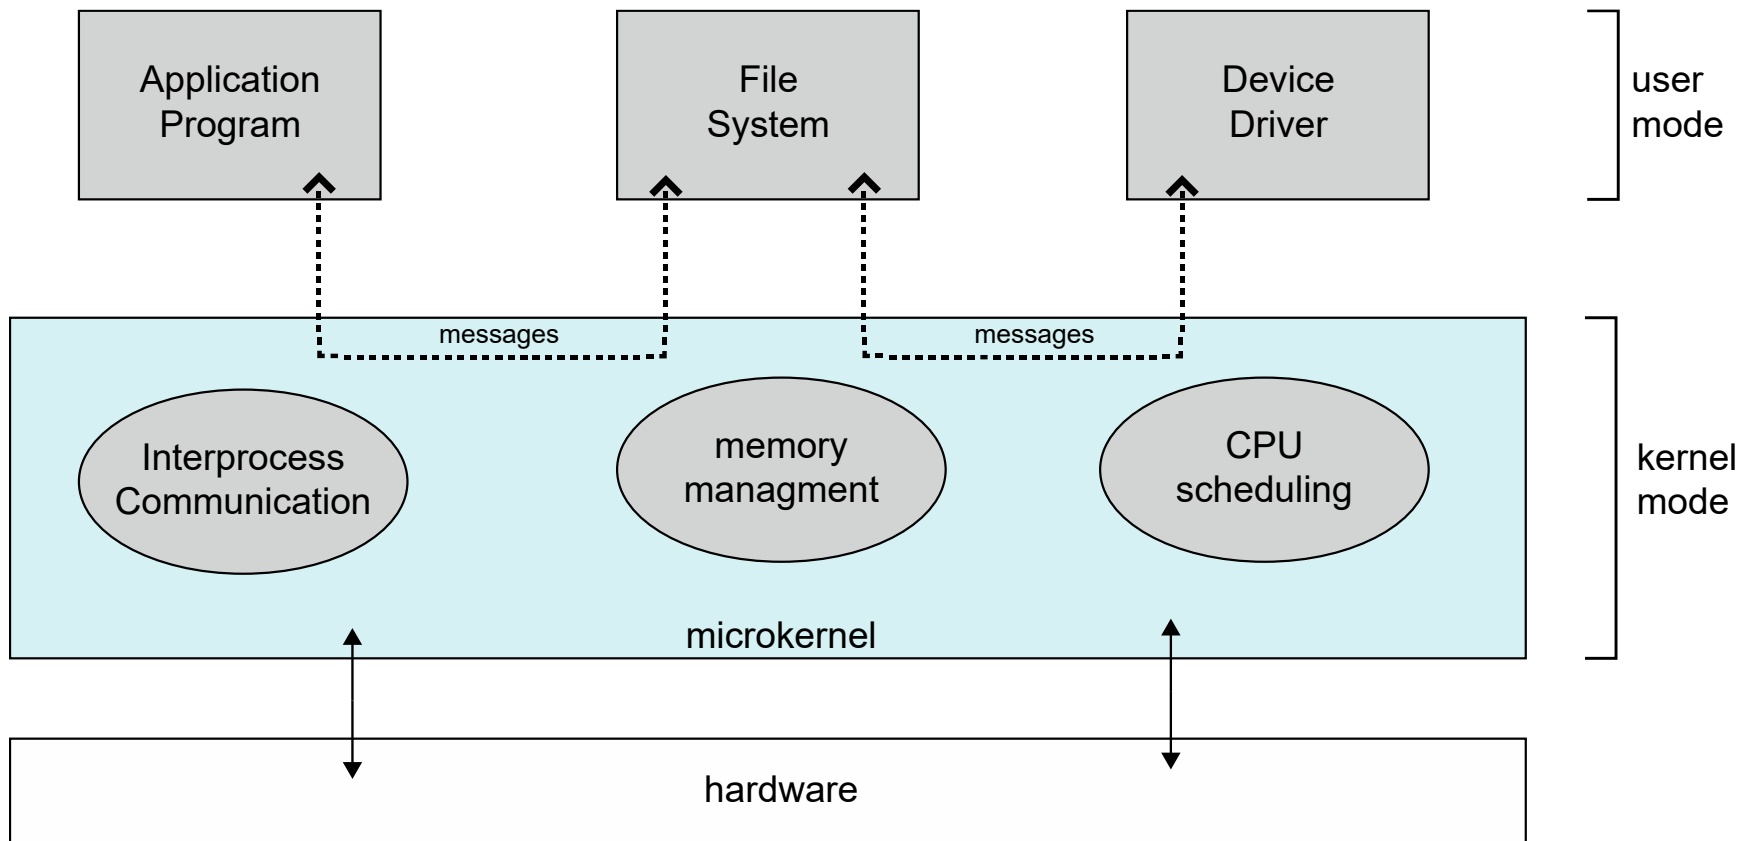
---

- ▶ System performance can be worse than in monolithic kernels, especially if kernel minimisation is taken too far



# Microkernel System Structure

---



# Modules

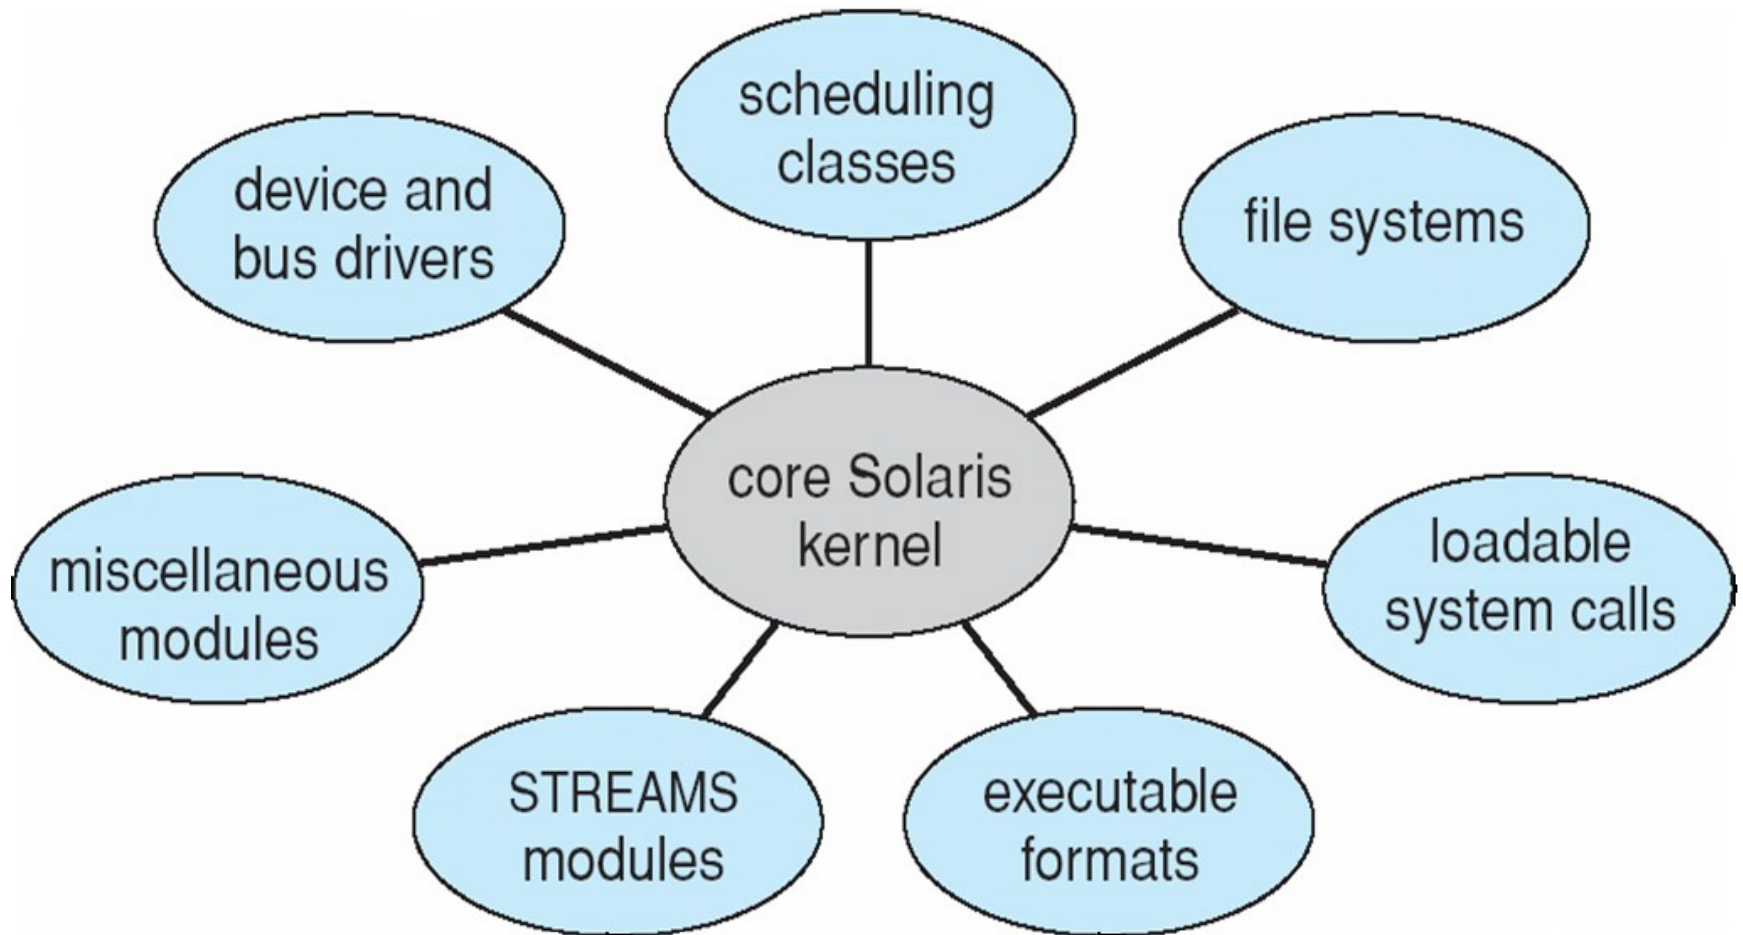
---

- ▶ Many modern operating systems implement loadable kernel modules
  - ▶ Uses object-oriented approach
  - ▶ Each core component is separate
  - ▶ Each talks to the others over known interfaces
  - ▶ Each is loadable as needed within the kernel
- ▶ Overall, similar to layers but with more flexibility
  - ▶ Linux, Solaris, etc



# Solaris Modular Approach

---



# Hybrid Systems

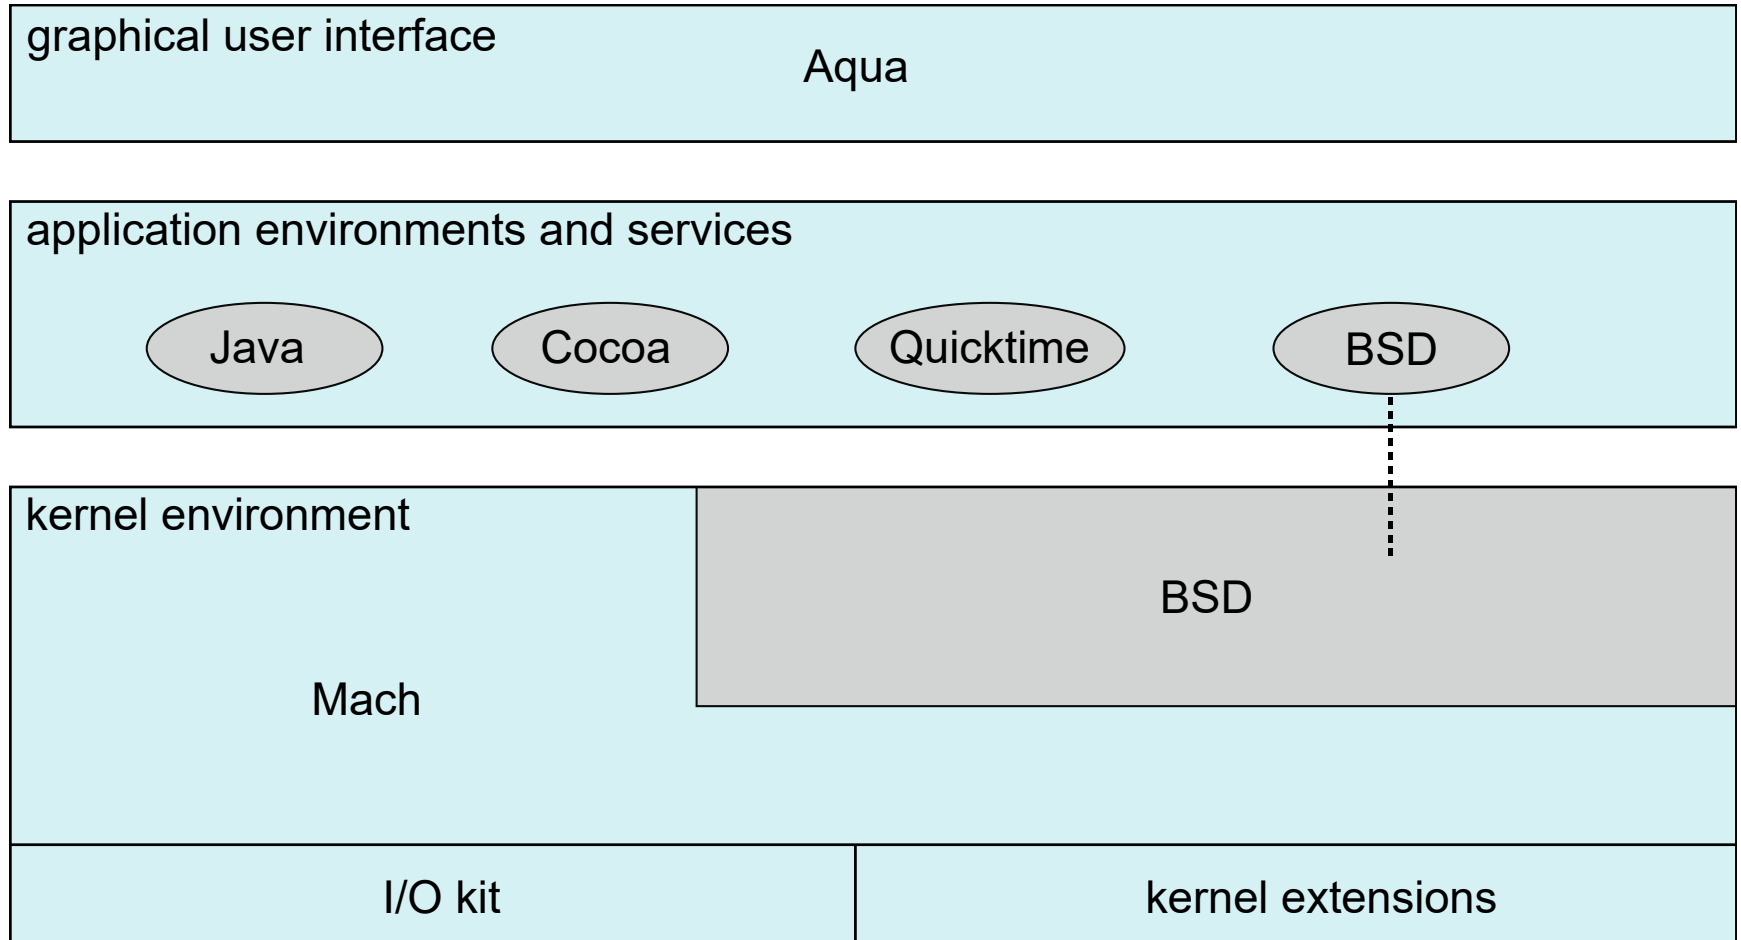
---

- ▶ Most modern operating systems are actually not one pure model
  - ▶ Hybrid combines multiple approaches to address performance, security, usability needs
  - ▶ Linux and Solaris kernels are in kernel address space, so they are monolithic, but also modular for dynamic loading of functionality
  - ▶ Apple Mac OS X combines a layered approach with a kernel consisting of Mach microkernel and BSD Unix parts, plus I/O kit and dynamically loadable modules



# Mac OS X Structure

---



# iOS

---

- ▶ **Based on Mac OS X**

- ▶ Cocoa Touch Objective-C API for developing apps
- ▶ Media services layer for graphics, audio, video
- ▶ Core services provides cloud computing, databases
- ▶ Core operating system, based on Mac OS X kernel

Cocoa Touch

Media Services

Core Services

Core OS



# Andriod

---

- ▶ Open Source OS Developed by Open Handset Alliance (mostly Google)
- ▶ Similar stack to IOS
- ▶ Based on Linux kernel but modified
  - ▶ Provides process, memory, device-driver management
  - ▶ Adds power management





# Android

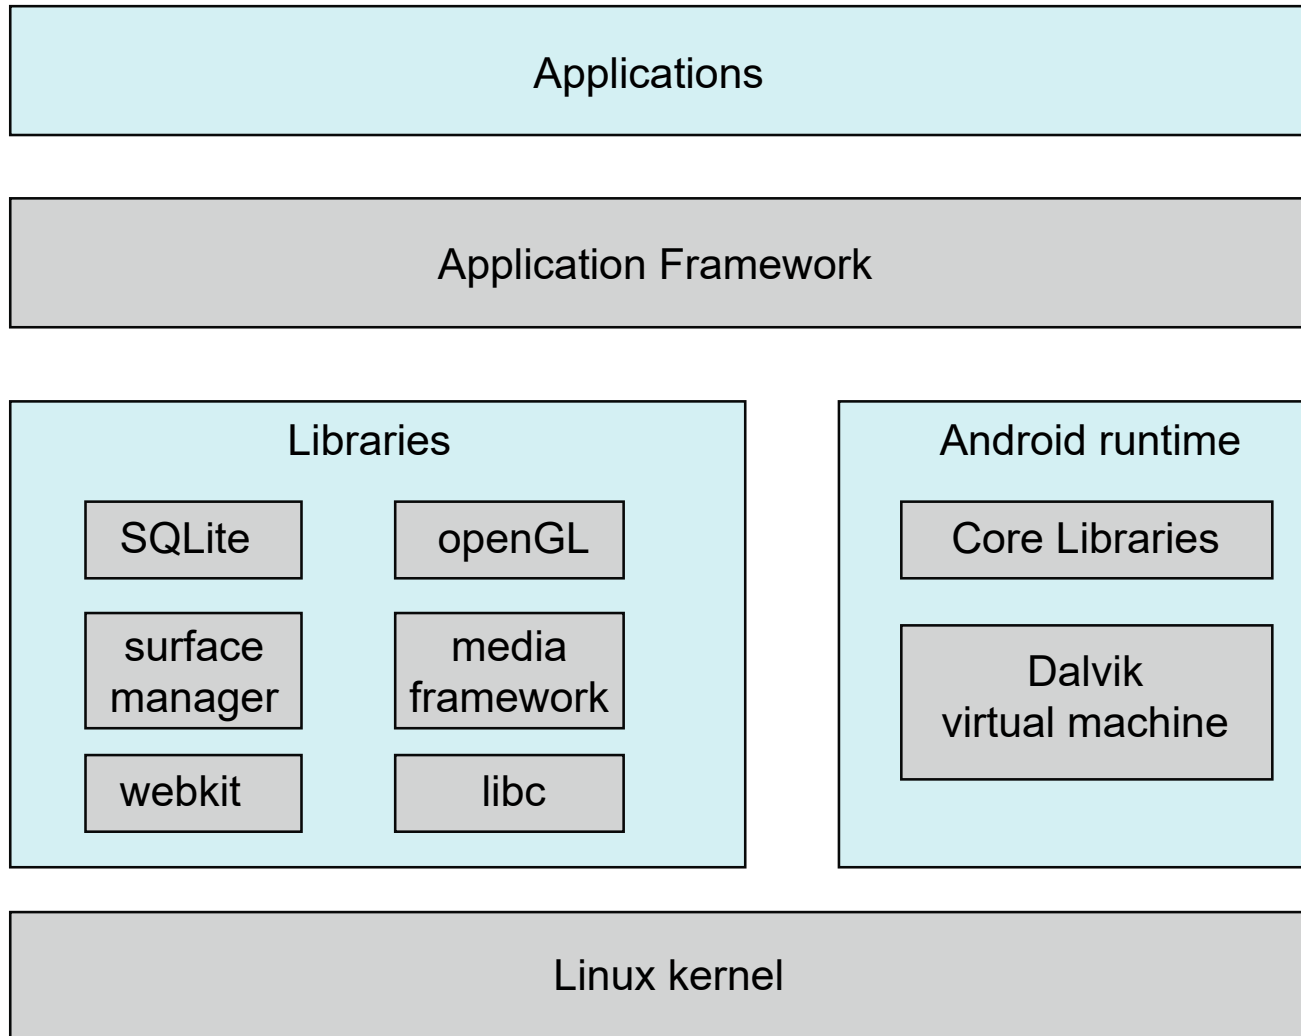
---

- ▶ Runtime environment includes core set of libraries and Dalvik/ART(2015) virtual machine
- ▶ Apps developed in Java plus Android API
- ▶ Java class files compiled to Java bytecode then translated to executable then runs in Dalvik VM



# Android Architecture

---



# Operating Systems Structure

Operating System and Hardware  
Features

# Operating System and Hardware

---

- ▶ An OS is partly dedicated to a specific hardware architecture
- ▶ Good hardware support can greatly simplify its design
- ▶ Hardware features are many times motivated by OS needs



# Protected Instructions

---

- ▶ In modern systems some instructions are typically restricted to the OS (**protected or privileged instructions**)
- ▶ This is to prevent a faulty or malicious user program from affecting the whole system



# Protected Instructions

---

- ▶ Typically, users are not allowed to
  - ▶ Directly access I/O (disk, printer,...)
  - ▶ Directly manage memory
  - ▶ Execute CPU halt instructions
- ▶ These operations are always handled through privileged instructions or memory mapping



# Dual Mode Operation

---

- ▶ The implementation of protected instructions requires some type of hardware mechanism
- ▶ The HW must support - at least - two operation modes
  - ▶ kernel mode: access to all the CPU instruction set
    - ▶ also called monitor/system/privileged mode
  - ▶ user mode: access restricted to a subset of the instruction set



# Dual Mode Operation

---

- ▶ The mode is indicated by a **status bit** (mode bit) in a protected processor register
  - ▶ OS programs & protected instructions executed in kernel mode
  - ▶ user programs executed in user mode
- ▶ Examples of protection in older and newer systems:
  - ▶ MS-DOS (based on Intel 8088): no protection modes
  - ▶ Windows 2000/XP, OS/2, Linux (based on Intel x86 systems): protection modes





# Crossing Protection Boundries

---

- ▶ User-mode programs cannot execute privileged instructions, but they still need kernel-mode services (I/O operations, memory management, etc)
- ▶ To execute a privileged instruction, a user must call an OS procedure: **system call**



# Crossing Protection Boundries

---

- ▶ A system call causes a **trap**, which jumps to the trap handler in the kernel
- ▶ When called the trap handler
  - ▶ uses call parameters to determine which system routine to run
  - ▶ saves caller's state: program counter (PC), mode bit, ...



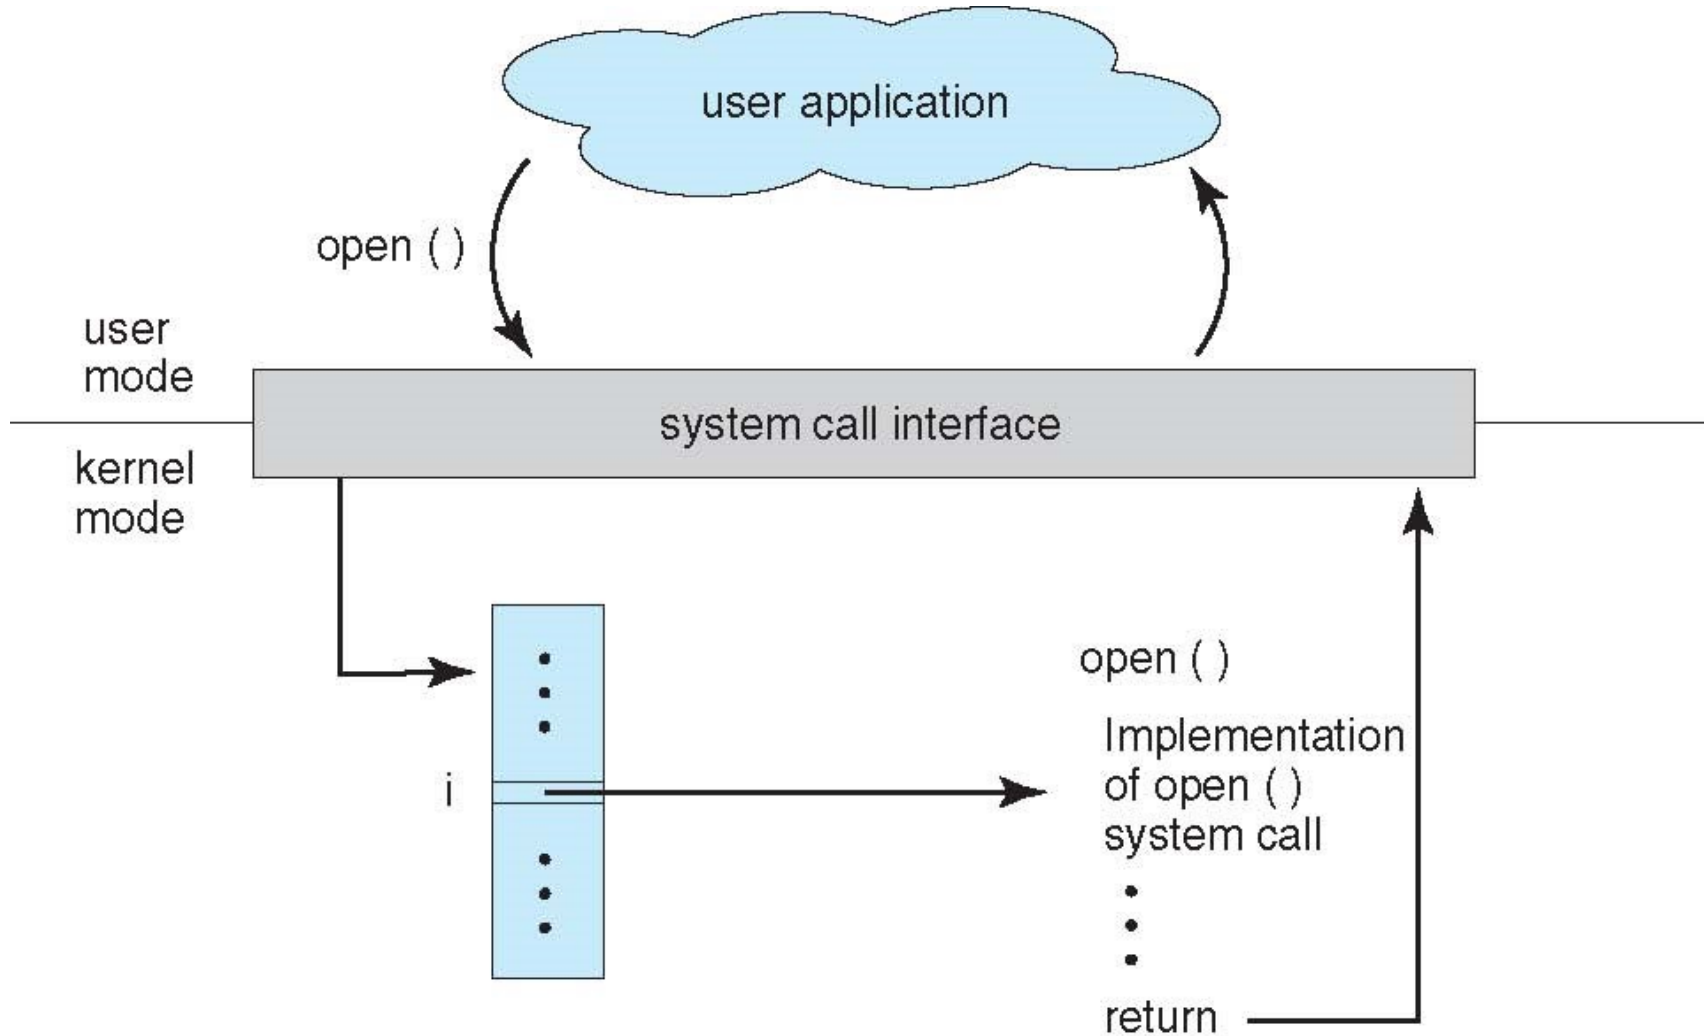
# Crossing Protection Boundries

---

- ▶ After this the hardware must
  - ▶ Implement caller's parameters verification (e.g. memory pointers should only be allowed within user's section)
  - ▶ Return to user-mode when trap system call finished
- ▶ The trap is treated by the hardware as a **software-initiated interrupt**



# The Open System Call



# Exceptions

---

- ▶ **Exception** (hardware-initiated interrupt):  
basically the same as a trap
- ▶ Exceptions are automatically triggered by an error or a particular situation rather than on purpose (like in a system call)



# Exceptions

---

- ▶ **Exceptions also transfer control to a handler within the OS**
  - ▶ system status can be saved on exceptions (memory dump), so that faulty processes can be later debugged
- ▶ **Exceptions decrease performance**
  - ▶ exception conditions could be detected by inserting extra instructions in the code, but at a high performance cost



# Typical Exceptions

---

- ▶ memory access out of user space
- ▶ overflow, underflow
- ▶ trace traps (debugging)
- ▶ illegal use of privileged instructions
- ▶ virtual memory (paging): page faults, write to read-only page



# Memory Protection

---

- ▶ A memory protection mechanism must protect
  - ▶ user programs from each other
  - ▶ OS (kernel) from user programs
- ▶ Simplest scheme is to use **base** and **limit** registers

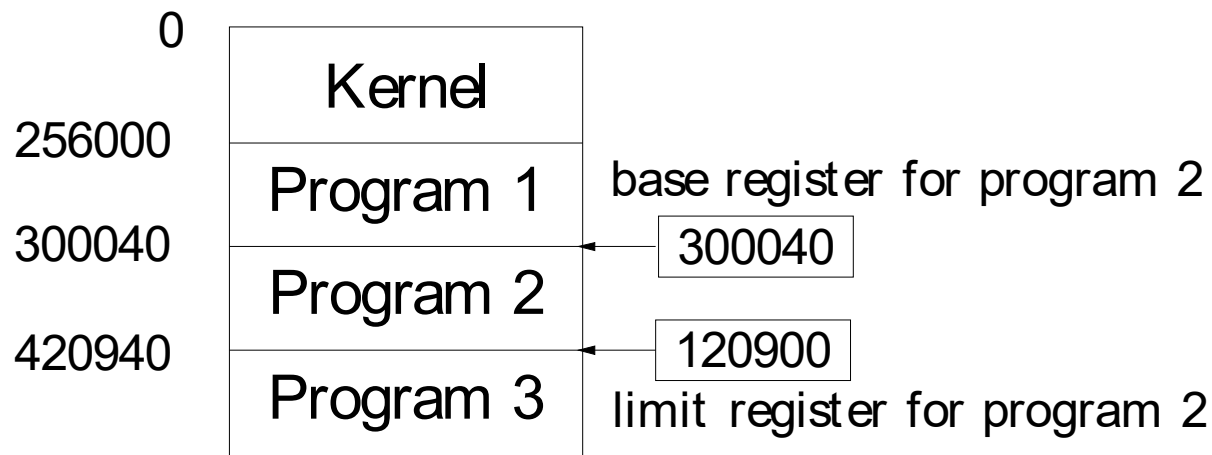




# Base and Limit Registers

---

- ▶ Base and limit registers are loaded by the OS before starting the execution of any program in user mode



- ▶  $\text{base} \leq \text{address} < \text{base} + \text{limit}$ ; otherwise exception raised
-

# Memory Protection

---

- ▶ Currently memory protection is more complex than this
- ▶ However, the simple base and limit scheme is the basis of modern virtual memory when a strategy called “segmentation” is used



# I/O Control

---

- ▶ **All I/O instructions are privileged**
  - ▶ This is because a program could disrupt the whole system by issuing illegal I/O instructions
- ▶ **Two situations must be considered:**
  - ▶ I/O start: handled by system calls
  - ▶ I/O completion and I/O events: handled by interrupts



# I/O Control

---

- ▶ Interrupts are the basis for **asynchronous** I/O
  - ▶ I/O devices have small processors that allow them to run autonomously (i.e. asynchronously with respect to the CPU)
  - ▶ I/O devices send interrupt signals when done with an operation; CPU switches to address corresponding to interrupt
    - ▶ an interrupt vector table contains the list of kernel routine addresses that handle different events



# CPU Protection

---

- ▶ Apart from protecting memory and I/O, we must ensure that the OS always maintains control
- ▶ A user program might get stuck into an infinite loop and never return control to the OS
- ▶ **Timer**: it generates an interrupt after a fixed or variable amount of execution time



# CPU Protection

---

- ▶ When an interrupt is generated by the Timer the OS may choose to treat the interrupt as a fatal error (and stop program execution) or allocate more execution time
- ▶ note: in time-sharing systems a timer interrupt is periodically generated after a fixed period of time (for scheduling a new program)



# Summary

---

- ▶ An OS provides a number of services
- ▶ These relate to managing
  - ▶ The hardware
  - ▶ The processes
  - ▶ The users
  - ▶ Communication between all of them
- ▶ Multiple ways of structuring the kernel, the system programs and user programs
- ▶ Each has advantages and disadvantages



## Summary - II

OS Requirement	Hardware Feature
Dual kernel/user modes	Protected instructions
System calls	Trap instructions and vectors
Exceptions, signals	Interrupt vectors
Memory protection	Base and limit registers
I/O control	Interrupts
CPU protection, scheduling	Timer (clock)





# That's all, folks!

---

► Questions?

