



University College Dublin  
An Coláiste Ollscoile, Baile Átha Cliath

# Operating Systems Process Synchronisation

Dr. Vivek Nallur ([vivek.nallur@ucd.ie](mailto:vivek.nallur@ucd.ie))

# Concurrent Execution

# Concurrent Execution

- Concurrent Execution is when two processes are executing at the **same** time
- Each cooperative processes may affect the execution of the other process

# Concurrent Execution Example

int A;

Process 1

```
A = 1;  
if(A == 1)  
    printf("Process 1 wins");
```

Process 2

```
A = 2;  
if(A == 2)  
    printf("Process 2 wins");
```

# Concurrent Execution Example

- The variable **A** is shared by both processes
- When both processes are run concurrently on one processor, which one "wins"?
  - the outcome of the concurrent execution depends on which assignment takes place first (**race condition**)

# What about threads?

- If we replaced the processes with threads the outcome will be still the same

# Atomic Operations

- Is it possible to get a different value in the variable A if we try and execute the two assignments statements at the same time?
- E.g.  $A = 1$  and  $A = 2$  happen at the same time

# Atomic Operations

- It is **not** possible.
- References and assignments are all **atomic** in the CPU
  - This means all read and write operations happen as a single step



# Atomic Operations

- An atomic operation cannot be interrupted
- This is to prevent illogical things happening

# Atomic Instructions

- This atomicity is provided by hardware
- Higher-level constructs are not atomic in general
  - A higher-level construct is any sequence of two or more instructions

# Higher-level Example

int B;

Process 1

```
B = 0;  
while(B < 10) B++;  
printf("Process 1 finished");
```

Process 2

```
B = 0;  
while(B > -10) B--;  
printf("Process 2 finished");
```

+ Variable B is shared

+ But increments and decrements are atomic

+ Will the process finish?

# Higher-level Example

int B;

Process 1

```
B = 0;  
while(B < 10) B++;  
printf("Process 1 finished");
```

Process 2

```
B = 0;  
while(B > -10) B--;  
printf("Process 2 finished");
```

- + The **while** sections above are not atomic
- + The process theoretically might never reach 10 or -10

# Higher-level Example

int B;

Process 1

```
B = 0;  
while(B < 10) B++;  
printf("Process 1 finished");
```

Process 2

```
B = 0;  
while(B > -10) B--;  
printf("Process 2 finished");
```

+ Process synchronisation is all about making **high-level** constructs **behave atomically**

# Mutual Exclusion & Critical Section

# The Milk Problem

- Two flat mates sharing a fridge
- They want to have **at most one** bottle of milk in the fridge at any time

# Possible Events

| Time  | Flatmate A                      | Flatmate B                      |
|-------|---------------------------------|---------------------------------|
| 10:00 | Check the fridge: NO MILK       |                                 |
| 10:05 | Walk to Shop                    | Check the fridge: NO MILK       |
| 10:10 | Arrive at the shop; buy milk    | Walk to Shop                    |
| 10:15 | Arrive home; put milk in fridge | Arrive at the shop; buy milk    |
| 10:20 |                                 | Arrive home; put milk in fridge |
|       |                                 | TOO MUCH MILK                   |



# The Milk Problem

- The reason for this problem is that the behaviour of A and B is not atomic
- To solve this problem A and B must **synchronise** (cooperate)

# Important Definitions

- **Synchronisation**: ensuring proper cooperation among processes (or threads), by relying on atomic operations
- **Mutual exclusion** (ME): ensuring that only one process at a time holds or modifies a shared resource
  - ME ensures atomicity

# Important Definitions

- **Critical section** (CS): a section of code in a program in which shared resources are manipulated
- Mutual exclusion in a critical section requires serialisation of process access to the critical section

# Milk Problem

- The critical section is
  - Check fridge, go shopping, put milk in fridge
- Mutual exclusion:
  - Only let one flatmate do this at a time

# Locking and Mutual Exclusion

- Achieving mutual exclusion in a critical section always involve some sort of **locking mechanism**
- Locking is where we prevent someone else for doing something with the shared resource

# Locking Rules

- Locking involves three rules:
  1. You must **lock** before you enter a critical section
  2. You must **unlock** when leaving a critical section
  3. You must **wait** when trying to enter critical section if it is locked

# Milk Problem Lock

1. Leave note before going shopping (“gone shopping”)
2. Remove note after shopping
3. Don’t go shopping if note has been left

# Milk Problem (1st Solution Attempt)

- Let both A and B execute (in a closed loop)

```
if(no_milk) {  
    if(no_note) {  
        put note;  
        buy milk;  
        remove note;  
    }  
}
```



# Milk Problem (1st Solution Attempt)

- Mutual exclusion is not guaranteed
  - A and B could execute "if(no\_note)" before the note is put on the fridge
- This is really hard to debug
  - It might work most of the time but occasionally fail

## Milk Problem (2<sup>nd</sup> Attempt)

- As a workaround, let us change the meaning of the note:
  - B buys if there is a note
  - A buys if there is no note

# Milk Problem (2<sup>nd</sup> Attempt)

## Process A

```
if (no_note) {  
    if (no_milk) {  
        buy milk;  
        add note;  
    }  
}
```

## Process B

```
if (note) {  
    if (no_milk) {  
        buy milk;  
        remove note;  
    }  
}
```

## Milk Problem (2<sup>nd</sup> Attempt)

- Is the critical section mutually exclusive?
- Yes, but what happens if B goes on holidays?
  - A cannot buy any milk!

## Milk Problem (2<sup>nd</sup> Attempt)

- We have the same problem if B is very slow
  - The **relative speed** of processes can be an issue
- Processes must take turns to be in the critical section
  - This means that the same flatmate **cannot** buy milk twice in a row

## Milk Problem (3<sup>rd</sup> Attempt)

- In order to try to avoid the last issue, let us use **two** notes (**note\_A** and **note\_B**) and some basic courtesy protocol
- Each process can now examine each other's status, but not modify it.

## Milk Problem (3<sup>rd</sup> Attempt)

Process A

put note\_A;

if(**no\_note\_B**) {

    if(no\_milk) {

        buy milk;

    }

}

remove note\_A;

Process B

put note\_B;

if(**no\_note\_A**) {

    if(no\_milk) {

        buy milk;

    }

}

remove note\_B;

# Milk Problem (3<sup>rd</sup> Attempt)

- Does this work?
- Better than before.
  - Relative speed is still an issue
  - What happens if they both leave notes at the exact same time?



# Milk Problem (4<sup>th</sup> Attempt)

- Two processes use slightly different code
  - One process stays the same as before
  - The other **waits** until the note is removed

## Process A

```
put note_A; if(no_note_B)
{
    if(no_milk)
        buy milk;
} else {
    while(note_B){ }
    if(no_milk)
        buy milk;
}
remove note_A;
```

## Process B

```
put note_B; if(no_note_A) {
    if(no_milk) {
        buy milk;
    }
}
remove note_B;
```

# Milk Problem (4<sup>th</sup> Attempt)

- Is this a solution?
  - It was actually the first solution ever given to the mutual exclusion problem
- But it is not satisfactory

# Issues with the Solution

- The solution is complicated and just for this problem
  - We need to pay careful attention to make sure that it really works
  - A similar solution might not be good for more complex problems

# Issues with the Solution

- The solution is **asymmetric**
  - Not everybody is performing the same steps
- The solution is **not scalable**
  - How would we make it work for more people sharing milk

# Issues with the Solution

- The solution is **inefficient**
  - While A is waiting it is consuming processor time
    - Called **busy-waiting**
- The solution is not scalable
  - How would we make it work for more people sharing milk

# Consequence

- Because of the problems with the 4<sup>th</sup> solution it is necessary to have standard synchronisation mechanisms in an operating system
- These can automatically fulfill some minimum requirements

# Requirements for True Solution for CS

## 1. Mutual exclusion

- One process **at most** inside the CS at any time

## 2. Bounded waiting (no starvation)

- A process attempting to enter its CS will eventually do so



# Requirements for True Solution for CS

## 3. Progress

- A process executing outside a CS cannot prevent another process from entering it
- If several processes are attempting to enter a CS at the same time the decision on which one goes in cannot be indefinitely postponed
- A process can't stay inside its CS forever (or exit in there)

# Requirements for True Solution for CS

- These conditions are **necessary and sufficient** for process synchronisation
  - Provided that basic operations are **atomic**
- No assumptions are made about: number of processes, relative speed of processes, or underlying hardware

# Desirable Properties of a ME Mechanism

- Simple
- Systematic and easy to use
  - E.g. just bracket the critical section
- Easy to maintain

# Desirable Properties of a ME Mechanism

- Efficient
  - Does not use a lot of resources while waiting
    - No busy-waiting
  - Overhead due to entering and leaving critical sections has to be small
    - At least smaller than the work inside it
- Scalable:
  - It should work when many threads share the critical section

# Implementations

1. There are three basic mechanisms for implementing mutual exclusion in critical sections
  - Semaphores
    - Simple, but hard to program with
    - Very low level

# Implementations

## 2. Monitors

- Higher level mechanism
- Requires higher-level programming languages

## 3. Messages

- Synchronisation without shared memory
- Uses IPC messages instead

# Semaphores

# Semaphore

- A semaphore is a protected integer variable **S** with an associated queue of waiting processes
- Only two atomic operations **P()** and **V()** can be performed on S



# Semaphore Operations

P(S)

```
if (S > 0) {  
    S--;  
} else {  
    addToQueue();  
}
```

V(S)

```
if (queue_not_empty) {  
    nextInQueue();  
} else {  
    S++  
}
```

# Semaphores

- Initial value of S is **how many** threads can enter critical section at the same time
- Processes in the queue are **blocked**
- The operations are **atomic**, so only one process at a time can execute them

# Mutual Exclusion using Semaphores

- To implement mutual exclusion using semaphores, we do the following:
  - Initialise S at 1
  - To enter critical section we execute **P** on its semaphore
  - When leaving the critical section we execute **V** on its semaphore

# Counting Semaphores

- If we initialise  $S > 1$ , more than one process at a time can get into the CS
  - This means that there is no mutual exclusion
  - This is used in a particular type of semaphores called **counting semaphores**
- A semaphore with the initial value of 1, is also called a **binary semaphore** or **mutex**

# Why P and V?

- Dijkstra, who first proposed semaphores, was Dutch
- In Dutch
  - **P**roberen means to probe
  - **V**erhogen means to increment

# Creating Semaphores

- Creating semaphores is different for each language
- All will require an initial value
- Other information like who can access may also be required

# Semaphores Example

- Atomically increment a shared variable X by n concurrent processes
- Setup:

```
int x;  
semaphore S (1, NULL) ;
```

- Code in threads:

```
P (S) ;  
x = x + 1 ;  
V (S) ;
```

# Semaphores for Milk problem

- Declare semaphore:

```
semaphore S (1, NULL);
```

- Processes code:

```
P (S) ;  
if (no_milk)  
    buy_milk  
V (S) ;
```



# The Produce/Consumer Problem

- This is a common synchronisation problem
- Consider a **producer process** supplying a resource to a **consumer process**
  - Producer: creates instances of a resource
  - Consumer: uses up instances of a resource

# The Produce/Consumer Problem

- Producer and consumer share a **buffer**
  - The producer put resources into the buffer
  - The consumer takes resources from the buffer
  - The buffer has a **finite** size

# Problem Constraints

- The consumer must wait for producer if the buffer is empty
- Producer must wait for consumer if the buffer is full

# Synchronisation

- We need to keep the producer and consumer synchronised
  - The buffer is the critical section in this problem
  - We usually need mutually exclusive access to the buffer

# P/C Solution with Semaphores

## Producer

```
while(true) {  
    msg = produce();  
    P(S) ;  
    if(buffer_full)  
        wait();  
    put_msg(msg);  
    V(S) ;  
}
```

## Consumer

```
while(true) {  
    P(S) ;  
    if(buffer_empty)  
        wait();  
    msg=get_msg();  
    V(S) ;  
    consume(msg);  
}
```

# P/C Solution with Semaphores

- The solution cannot work
  - If the buffer is **empty** the consumer waits and the producer cannot add any items
  - If the buffer is **full** the producer waits and the consumer cannot remove any items
- These situations are called **deadlock**

# P/C Solution with Counting Semaphores

- We add two more semaphores to the code:
  - One to count the number of **empty** slots in the buffer
  - One to count the number of **full** slots in the buffer

# P/C Solution with Counting Semaphores

- Semaphores:

```
int N = buffer_size;  
semaphore S(1, NULL);  
semaphore full_s(0, NULL);  
semaphore empty_s(N, NULL);
```



## Producer

```
while(true) {  
    msg=produce();  
  
    P(empty_s);  
    P(S);  
    put_msg(msg);  
    V(S);  
    V(full_s);  
}
```

## Consumer

```
while(true) {  
    P(full_s);  
    P(S);  
    msg=get_msg();  
    V(S);  
    V(empty_s);  
  
    consume_msg(msg);  
}
```

# Questions

- True or false?
  - The value of a semaphore cannot be modified by an operation other than P or V
  - The initial value of a semaphore cannot be zero
- Is the order of Ps important? And the order of Vs?

Monitors

# The Problem with Semaphores

- Semaphores are nice but they have an important problem
  - Using more than one semaphore can lead to **deadlocks**, if order of Ps not correctly set
  - It is difficult to program with semaphores

# Monitors

- Monitors are higher-level sync primitives
  - A **monitor** is a collection of procedures, variables and data grouped together in a special kind of structure
  - Their aim is to avoid the problems that can arise when protecting critical sections with semaphores

# Example Monitor

```
monitor example {  
    int i;           (Internal data)  
    condition c;     (Condition variable)  
  
    void p() {       (Monitor procedure)  
        ...  
    }  
}
```

# Monitor Rules

- Processes may call monitor procedures whenever they wish
- Processes can't access internal monitor data (variables, etc.) using external procedures

# Monitor Rules

- Only one process can be active inside a monitor
- Mutual exclusion inside the monitor is guaranteed by the compiler



# Blocking in Monitors

- Mutual Exclusion is implemented using internal **condition** variables
- Conditions have two operations
  - Wait
  - Signal

## wait(**condition**)

- This is executed when the monitor discovers that a process cannot proceed
  - The monitor **blocks** a process calling wait() and makes it wait on condition
  - Another process can then be allowed into the monitor
  - A process that calls wait() is **always blocked**

`signal(condition)`

- This is executed to wake up a process that is waiting on **condition**
- After executing **signal()** the calling process must exit the monitor immediately
  - This guarantees mutual exclusion

# Monitor for Producer/Consumer

- Requires two conditions:
  - Buffer is full
  - Buffer is empty
- Requires some internal data so we know when it is empty or full

# Monitor for Producer/Consumer

- Requires two operations
  - Put a message in the buffer
  - Remove a message from the buffer

# Monitor for Producer/Consumer

```
monitor pr_co {  
    int count;  
    condition full_s, empty_s;  
    void put(msg) {  
        if(count == N)  
            wait(empty_s);  
        put_msg(msg);  
        count++;  
        if(count == 1)  
            signal(full_s);  
    }  
}
```

# Monitor for Producer/Consumer

```
msg get () {  
    if (count==0)  
        wait (full_s) ;  
    msg=get_msg () ;  
    count-- ;  
    if (count==N-1)  
        signal (empty_s) ;  
}  
}
```

# Producer and Consumer Code

## Producer

```
while(true) {  
    msg=produce();  
    pr_co.put(msg);  
}
```

## Consumer

```
while(true) {  
    msg=pr_co.get();  
    consume_msg(msg);  
}
```



# Producer and Consumer Code

- Simplest possible consumer and producer code
- Monitor takes care of any issues, not producer or consumer

Messages

# Messages System


- Synchronisation with both semaphores and monitors relies on **shared memory**
- The message system is a mechanism for inter-process communication (IPC), that can also be relied upon for synchronisation without having to share memory

# Message System

- Elements in the message system:
  - **message**: information that can be exchanged between two (or more) processes or threads
  - **mailbox**: a place where messages are stored between the time they are sent and the time they are received

# Message System Scheme

- Process A sends messages to process B using a mailbox



# Message System Scheme

- One process or the other owns the data
  - Never two at the same time
- If processes need to send messages both ways, we need two mailboxes

# Message System Operations

- Two basic operations
  - **send(mailbox, message)**
    - Put a message in the mailbox
  - **receive(mailbox, message)**
    - Remove a message from the mailbox
- In practice we also need the system calls **create(mailbox)** and **delete(mailbox)**

# Message System Addressing

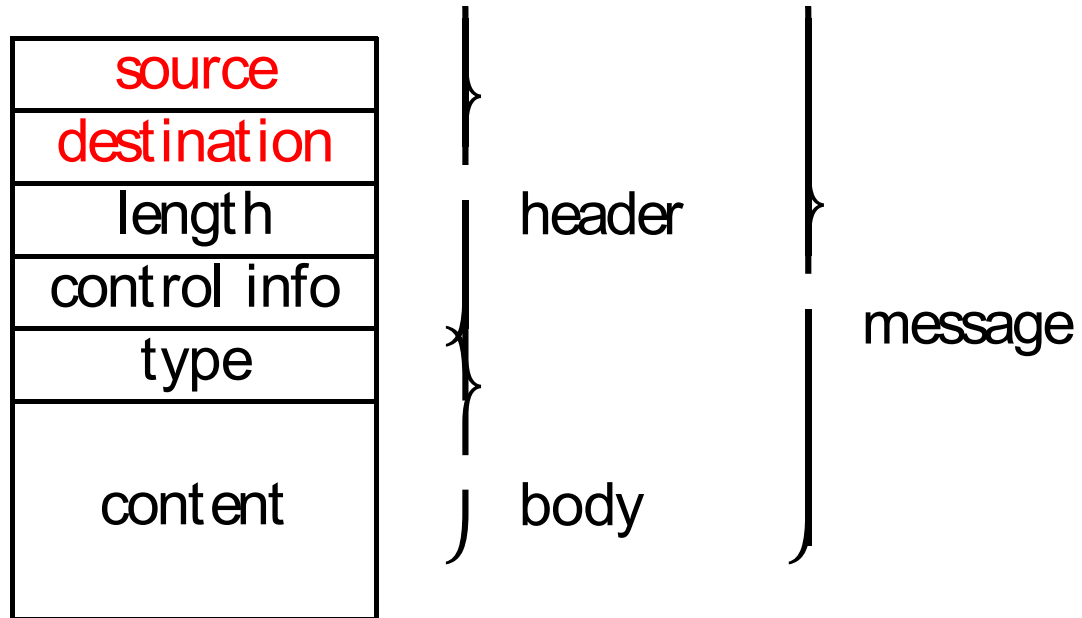
- The message identifies **source** and **destination** processes
  - Through their PIDs
- There may be multiple destinations, or destination may be unspecified



# Message Format

- Depends on the objectives of the system
  - **fixed-length**: minimise processing and storage overhead
  - **variable-length**: more flexible approach for operating systems

# Message Format



# Why Use Messages?

- Communicating processes sometimes need to be separate
  - They do not trust each other
  - The programs were written at different times by different programmers who knew nothing about each other
  - They run on different processors
  - Less error-prone without shared memory

# Mutual Exclusion with Messages

- Mutual exclusion must be guaranteed solely relying on the mailbox
- Two levels of synchronisation in a message system
  - Producer/Consumer Constraints
  - Execution flow constraints

# Producer/Consumer Constraints

- `receive()`
  - Calling thread waits when the mailbox is empty until a message is added
- `send()`
  - Calling thread waits when mailbox is full until there is space for another message

# Execution Flow Constraints

- Send and receive: could be blocking or non-blocking
  - *Blocking send*: when a process sends a message it **blocks** until the message is received at the destination.
  - *Non-blocking send*: After sending a message the sender **continues** with its processing without waiting for it to reach the destination

# Execution Flow Constraints

- Send and receive: could be blocking or non-blocking
  - *Blocking receive*: When a process executes a receive it waits **blocked** until the receive is completed and the required message is received
  - *Non-blocking receive*: The process executing the receive proceeds **without waiting** for the message(!).

# Execution Flow Constraints

- The most common combination is that of
  - *Blocking Receive*
  - *Non-blocking send*



# Mutual Exclusion using Messages

- Mutual exclusion can be achieved using a **single shared mailbox**
- Processes can **enter** the critical section whenever they **receive** a message
- Whenever they **leave** the critical section they must **send** a new message to the mailbox

## Parent process

```
create_mailbox(mbox) ;  
send(mbox, NULL) ;  
parallel {  
    P(1) ;  
    P(2) ;  
    . . .  
    P(N) ;  
}
```

- Creates N processes
- Empty message required to access critical section

## Child Process

```
while (true) {  
    receive (mbox, msg) ;
```

Critical Section

```
    send (mbox, NULL) ;  
}
```

## P/C Algorithm Using Messages

- Not only we need to ensure mutual exclusion but also to keep a count
- For this two mailboxes are needed
  - one mailbox for consumers: data produced and consumed
  - one mailbox for producers: accountancy purposes

# P/C Algorithm Using Messages

- Producers generate data and send it to the consumers's mailbox
  - Only when the producer's mailbox indicates there is at least one free slot in the buffer
- Consumers take messages from the consumer's mailbox when available
  - Then send empty messages to the producer's mailbox to signal new free slot

## Parent Process

```
create mailbox(cons_mbox);  
create mailbox(prod_mbox);  
for(i=0;i<N;i++)  
    send(prod_mbox, NULL);  
  
parallel {  
    producers();  
    consumers();  
}
```

## Producer Process

```
while (true) {  
    msg=produce() ;  
    receive(prod_mbox, token) ;  
    send(cons_mbox, msg) ;  
}
```

## Consumer Process

```
while (true) {  
    receive (cons_mbox, msg) ;  
    consume_msg (msg) ;  
    send (prod_mbox, NULL) ;  
}
```



# Next Class

- Next Lecture:
  - Deadlock & Starvation
- Review Chapter 5