



University College Dublin
An Coláiste Ollscoile, Baile Átha Cliath

Operating Systems Virtual Memory

Dr. Vivek Nallur (vivek.nallur@ucd.ie)

Virtual Memory

Virtual Memory: Characteristics of paging and segmentation:

- the address space of a process may be broken up into pieces (pages or segments)
- memory references within a process are logical addresses, dynamically translated into physical addresses

Virtual Memory: Characteristics of paging and segmentation:

- by virtue of the page/segment table, pages/segments of a process need not be contiguously located in main memory
- last but not least: it is not necessary that all pages/segments of a process be simultaneously in main memory during execution: part of the process may be disk

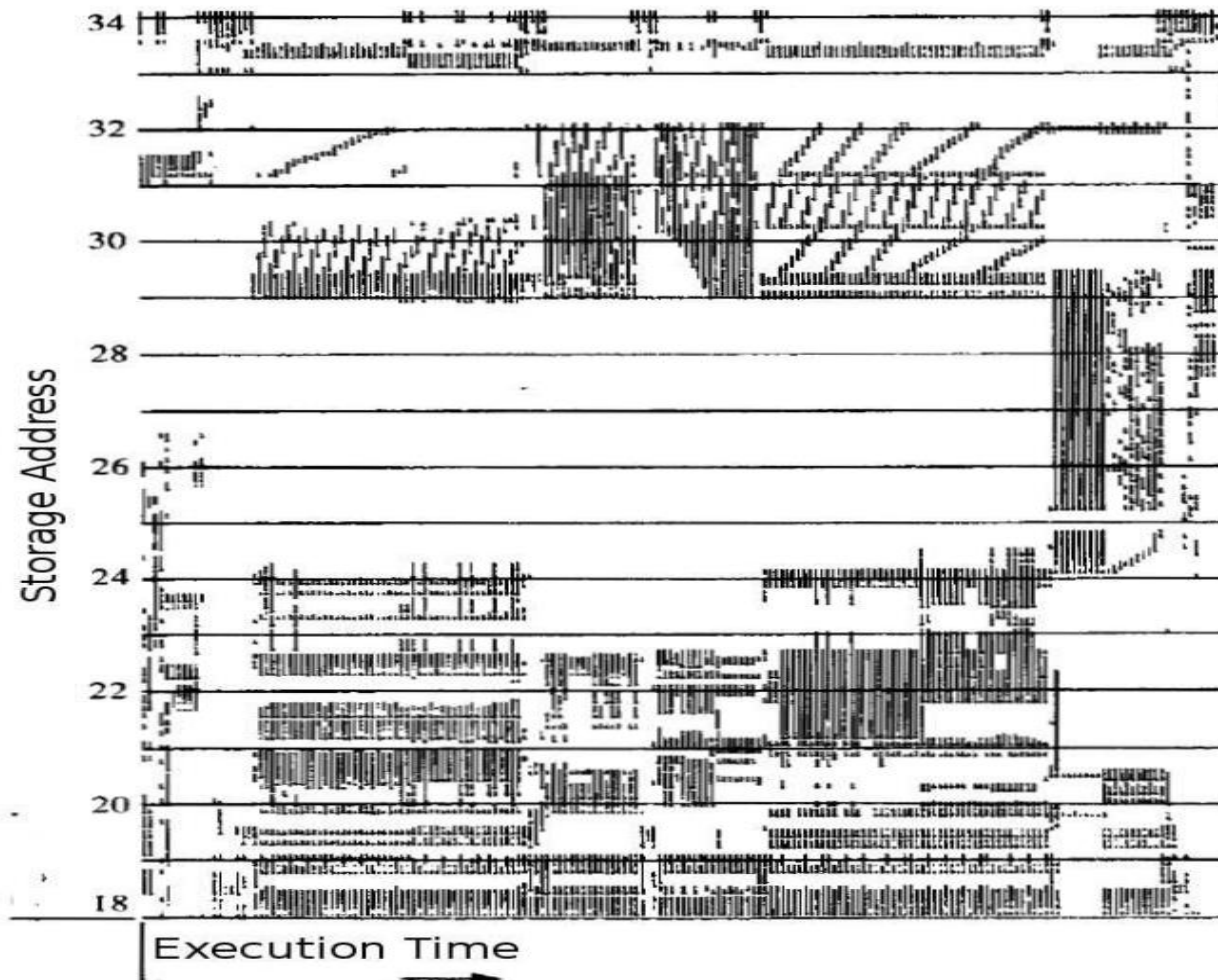
Virtual Memory

- **Virtual memory (VM):** is an illusion supported by system hardware and software that a process has a vast and linear expanse of available memory (much bigger than the main memory unit)
 - with VM, the main memory can be seen as a cache for the disk
 - logical addresses are also called **virtual addresses** in the context of VM

Real Programs and VM

- The VM scheme is appealing: is it practical as well?
- An examination of a real program shows that it can usually be divided into two parts:
 1. parts frequently needed
 2. parts rarely or never needed:
 - code handling unusual error conditions
 - code handling certain options and features rarely used
 - allocation of more memory than strictly needed (arrays, tables. . .)
- Also, memory references within a program tend to be clustered (principle of **locality of references**)

Example: Locality of Memory References



Real Programs and VM

- Therefore, it is possible to make intelligent guesses about which pieces of the address space of the process will most likely be needed in the near future
 - if guesses are good, then a system using VM will perform efficiently
- If guesses are poor, the system will suffer from thrashing
 - it will spend too much time swapping processes' pieces between memory and disk, rather than executing instructions
 - disk accesses are much slower than memory accesses

Virtual Memory Features

- VM is commonly implemented by demand paging
 - when a program is loaded, the OS brings into main memory only a few pages of it (including its starting point)
 - further pages are then brought to memory or swapped to disk as needed
 - the **resident set** is the portion of the process that is in main memory at a given time

Virtual Memory Features

- VM system supported by hardware: paging mechanism, which generates page faults
 - when pages in disk are referenced
 - software: page swapping management (OS algorithm)

Page Table with Demand Paging

- Additions to page table entries:
- **present bit**, set to one if the piece is in main memory
 - an interrupt (**page fault**) is generated when a reference is made to a piece not present in main memory
- **modified bit**, set to one if the piece has been altered since it was last loaded into main memory
 - if not modified: it does not need to be written to disk when it has to be swapped out
- **control bits**:
 - r/w page: a read-only page can be shared by two or more processes
 - kernel/user page: kernel pages can be forced to be in memory at all times

Page Table with Demand Paging

- main entry:
- if page in memory: frame number
- if page in disk: address in disk, or index to a table
 - referenced in the PCB— used to obtain its address

VM Advantages

- Programs **not constrained by the physical memory space**
- they can be as large as the virtual address space will allow :
 - 1970s: 32-bit virtual addresses → 4 GB of memory
 - 2000s: 64-bit virtual addresses → 16 EB of memory
- in both periods, that amount of real memory would cost millions of euro.
- Also in size, 16 EB in modern 4 TB hard drives could be placed end to end 467.2 km) and cover the whole of Beijing Subway system with 2.2 km to spare.

VM Advantages

- **Better multiprogramming**: more processes can be maintained in memory at any given time
 - more likely that one of these processes will be in ready state (remember $1 - p^n$ CPU utilisation formula)
- **Less I/O** is needed to load a program or to swap it
 - increase in CPU utilisation and throughput

Page Fault Procedure

- The page-fault procedure requires three major steps:
 1. service the page-fault interrupt
 - trap to the OS
 - save registers and process state
 - check if the page reference is legal and determine its location in the disk
 2. read in the missing page
 3. restart the process

Page Fault Procedure

- The page-fault procedure requires three major steps:
 1. service the page-fault interrupt
 2. read in the missing page
 - issue a read of the missing page from the disk to a free frame
 - since this is an I/O operation (slow), another ready process is dispatched to the CPU
 - eventually, an interrupt from the disk is issued (I/O completed)
 3. restart the process

Page Fault Procedure

- The page-fault procedure requires three major steps:
 1. service the page-fault interrupt
 2. read in the missing page
 3. **restart the process**
 - an interrupt from the disk indicates that the page is in memory
 - the process's page table is updated to show that the desired page is now in memory; the process can now be switched to the ready state
 - eventually the process is scheduled again, and resumes execution at the instruction interrupted through the page fault

Page Faults and Performance

- It is important to keep page faults to a minimum: why?
- If the probability of a page fault (or page fault rate) is p ($0 \leq p \leq 1$), then the **effective memory access time** (T_e) is

$$T_e = p \times T_f + (1 - p) \times T_a$$

T_f : page fault time (time to process a page fault)

T_a : memory access time

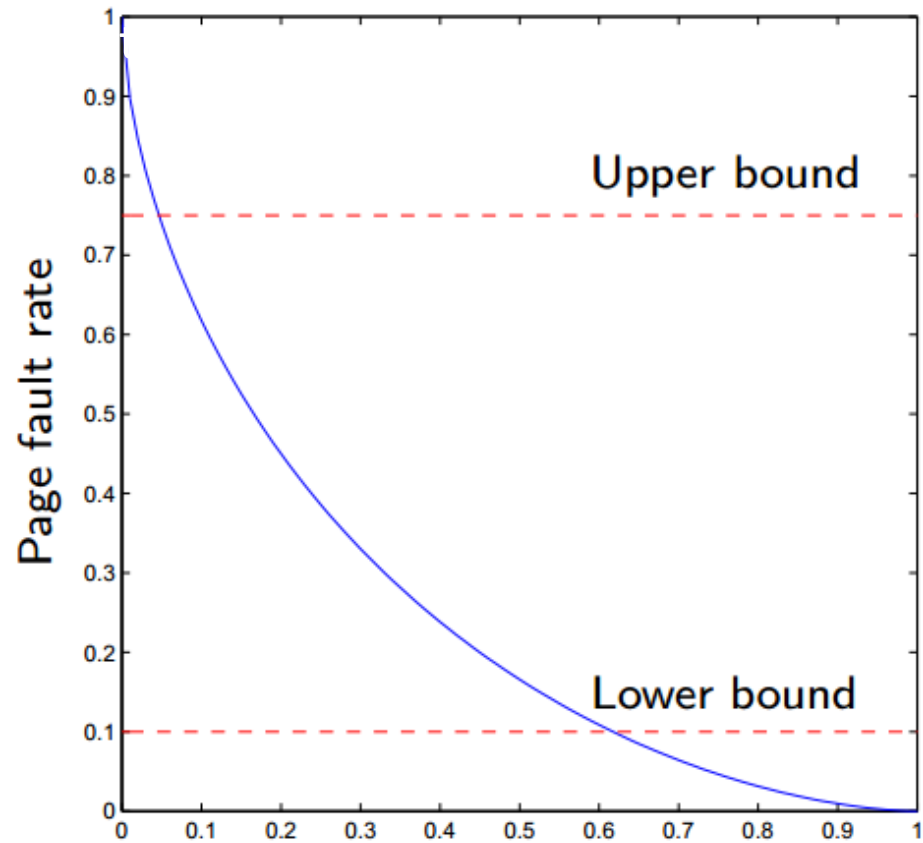
$100 \times p$: % of memory access that raise a page fault.

Page Faults and Performance

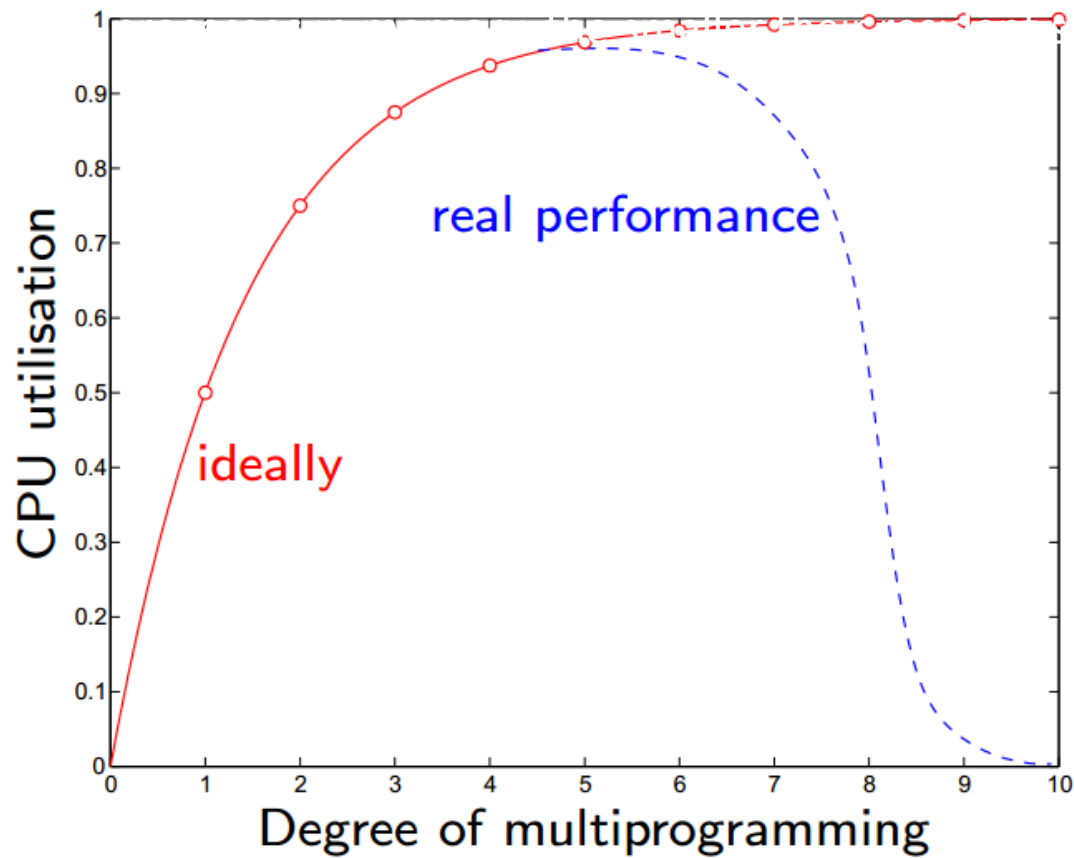
$$T_e = p \times T_f + (1 - p) \times T_a$$

- Typically, T_f can be 5 orders of magnitude higher than T_a (e.g. **10 ms** vs **100 ns**)
 - therefore, even low probabilities of page fault may raise T_{ae} noticeably

Page Faults and Thrashing



Thrashing and Multiprogramming



Thrashing and Multiprogramming

- To avoid this, **load control** measures are needed. Examples:
 - only execute processes whose resident set is large enough
 - adjust effective degree of multiprogramming so that mean time between faults equals target (suspend processes as needed)

Memory Management in VM

- In a system with VM, the OS must mainly deal with two memory management issues:
 1. Replacement policy
 2. Resident set management

Memory Management in VM

1. Replacement policy

- what happens when a page fault occurs and there is no free frame to swap a page in?
- the OS must select a frame for replacement (to be swapped out to disk) when a new page must be brought in
- restrictions are usually placed on the page replacement policy:
 - much of the kernel and control structures of the OS are held on locked frames (cannot be replaced)

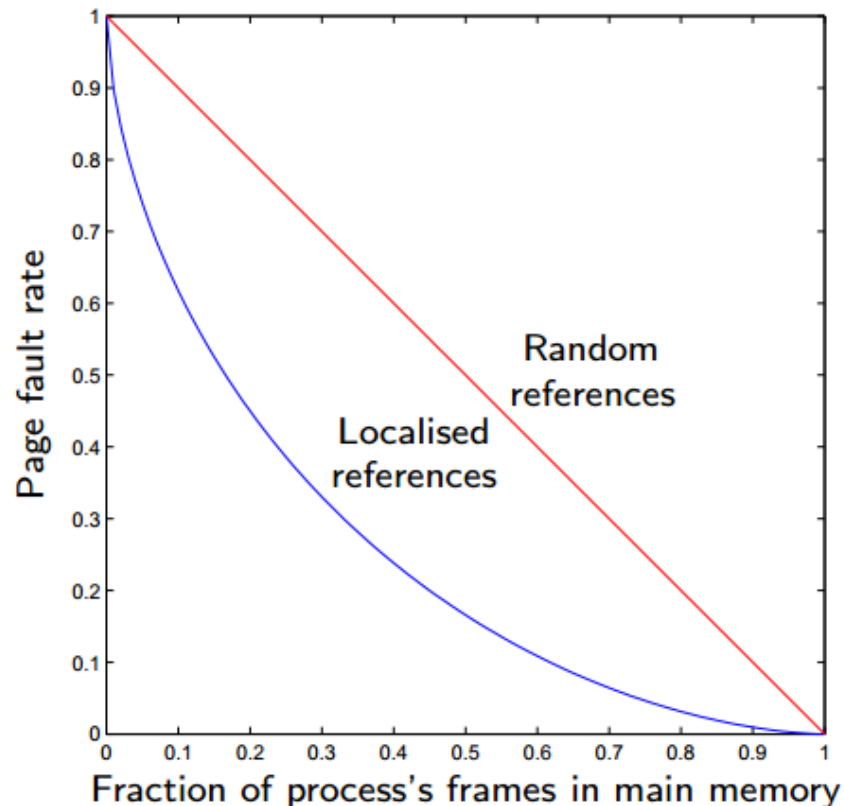
Memory Management in VM

2. Resident set management

- choice of dynamic or static number of frames for each active process
- choice of replacement type allowed; examples:
 - limited to pages of the process that caused the page fault
 - encompassing any frame in main memory

Replacement Policy and Locality

- A good replacement policy should exploit the principle of locality of references
- If memory references were random rather than localised, we would not be able to pin down the working set efficiently:



Replacement Algorithms

1. Optimal
2. LRU (least recently used)
3. FIFO
4. Clock policy

Replacement Algorithms

Optimal

- Optimal
 - replaces the page that **will not be referenced** for the longest period of time
 - minimum number of page faults, but **impossible to implement** (knowledge of future events required)
 - standard yardstick used to gauge other algorithms
 - toy example: consider the page address stream 2 3 2 1 5 2 4 5 3 2 5 2 by a process with static resident set of just 3 frames

Replacement Algorithms

Optimal

2	3	2	1	5	2	4	5	3	2	5	2
2	2	2	2	2	2	4	4	4	2	2	2
	3	3	3	3	3	3	3	3	3	3	3
			1	5	5	5	5	5	5	5	5

Replacement Algorithms

LRU (least recently used)

- replaces the page that has **not been referenced** for the longest period of time
- by the principle of locality: it is likely that this page will not be referenced in the near future either
- almost as good as the optimal policy, but difficult to implement (overheads associated to time keeping)

2

3

2

1

5

2

4

5

3

2

5

2

2

2

2

2

2

2

2

2

3

3

3

3

3

3

3

5

5

5

5

5

5

5

5

1

1

1

4

4

4

2

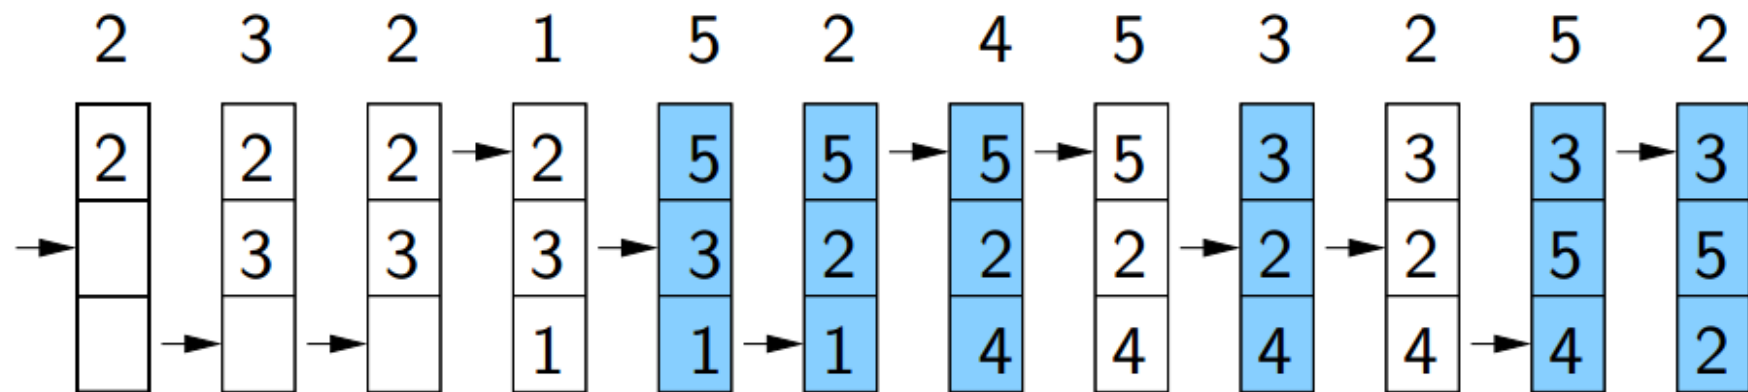
2

2

Replacement Algorithms

FIFO

- frames traversed as a circular buffer, triggered by replacements
- pages are removed in round-robin style using a pointer (→)
- rationale: a page fetched long ago may be now out of use (when main memory is composed by many frames)
- simple to implement, but some replacements will not be good



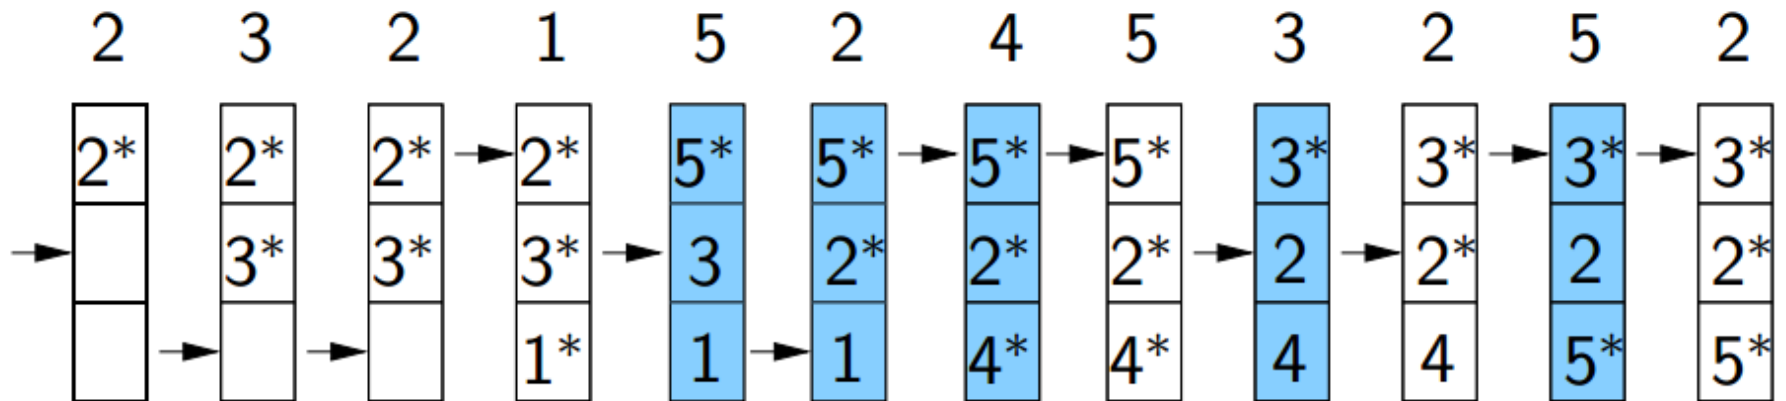
Replacement Algorithms

Clock policy

- variant of FIFO that accounts for page use
- hardware keeps a **use bit (reference bit)** per page
 - when page referenced or first loaded, use bit is set to 1
- replacement is made by OS:
 - frames traversed as a circular buffer as in FIFO policy
 - first frame encountered with use bit set to 0 is replaced (logic: it has not been used for while)
 - while traversing, frames with use bit 1 get it set to 0

Replacement Algorithms

Clock policy



Page Size Considerations

- Page size is invariably a power of 2, but how do we select it?
- No single best answer:
 - small page size leads to large page table
 - example: for a VM of 4 MB (222) we would have 4096 pages of 1024 bytes, but only 512 pages of 8192 bytes
 - memory is better used with smaller page sizes
 - on average, half of the last page is wasted
 - I/O transfer time is small compared to seek and latency, which favours larger pages
 - however locality improves with smaller page size (higher resolution)
- The historical trend is towards larger pages
 - 1990: typically around 4096 bytes
 - 2000's: OS's developed concept of Huge Page size
 - Typically in consumer level 2MB or 4MB
 - 2010's : This concept in enterprise systems can go up to 256 MB

Page Table Structure

- Issue: page tables can take a big portion of memory
 - 2^{m-n} entries required for an m -bit long virtual address with n page offset bits
 - example: with $m = 32$ and 4 KB page size $\rightarrow 2^{20} \approx 10^6$ entries
 - this problem is even more severe with modern 64-bit addresses
- Solutions:
 - hierarchical paging
 - inverted page table

Hierarchical Paging

- Two-level page table: the $m - n$ page number bits are divided into two sections
 1. the first points at an entry in the outer page table, which gives a frame corresponding to a page of the page table proper
 2. the second points at an entry within that page of the page table

Hierarchical Paging

- Example (using same values as in previous slide): $m - n = 20$ can be divided into two 10-bit numbers
- virtual address:

- the

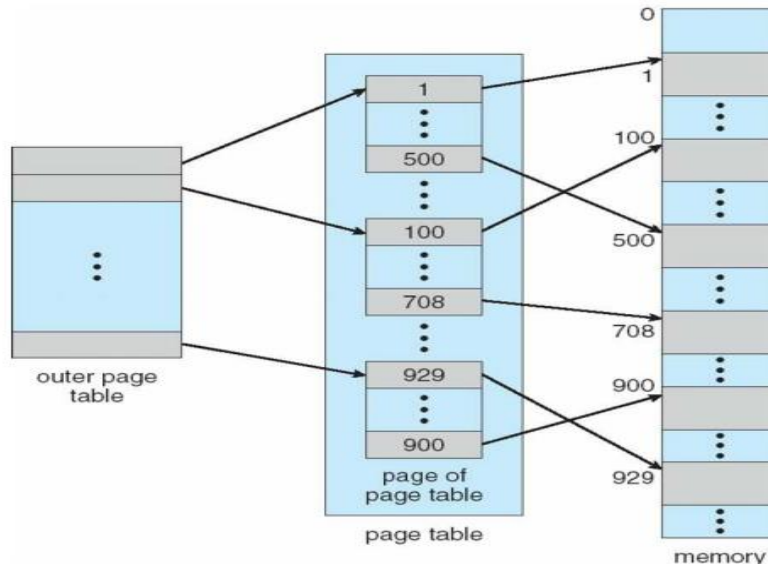
page number		page offset
p_1	p_2	n

 in one 4 KB page
- with

10	10	12
----	----	----

 page table, and enough frames to reference the pages actually used by the process

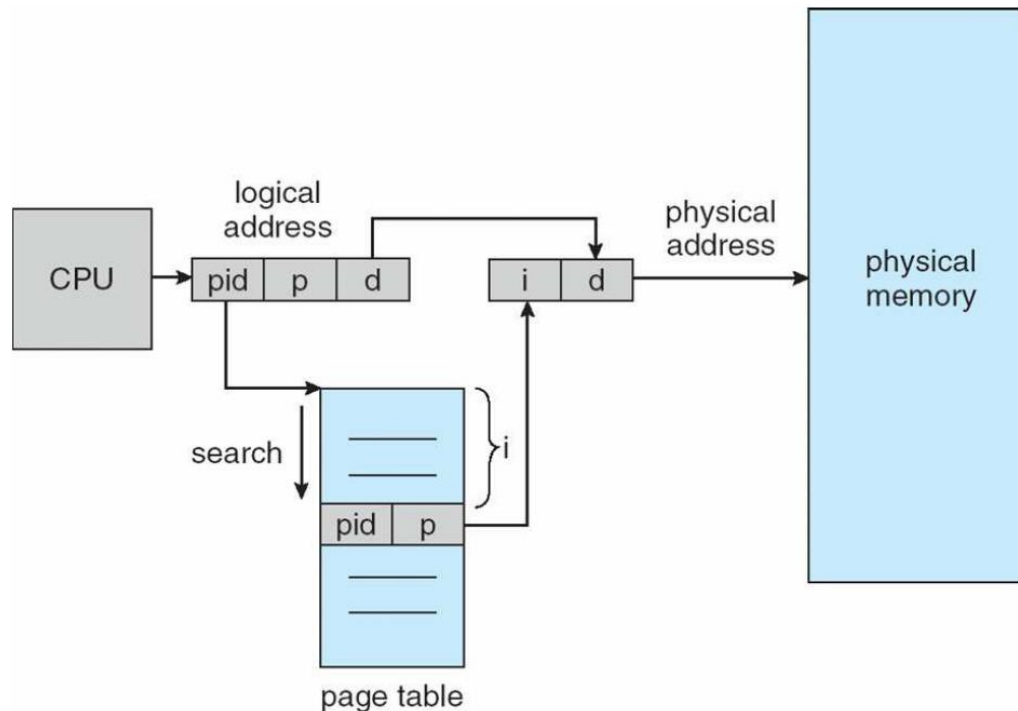
Hierarchical Paging (Two Levels)



- Drawback: the time to convert a virtual address into a physical one is longer
- This problem is even worse if m is large and we need more than two levels

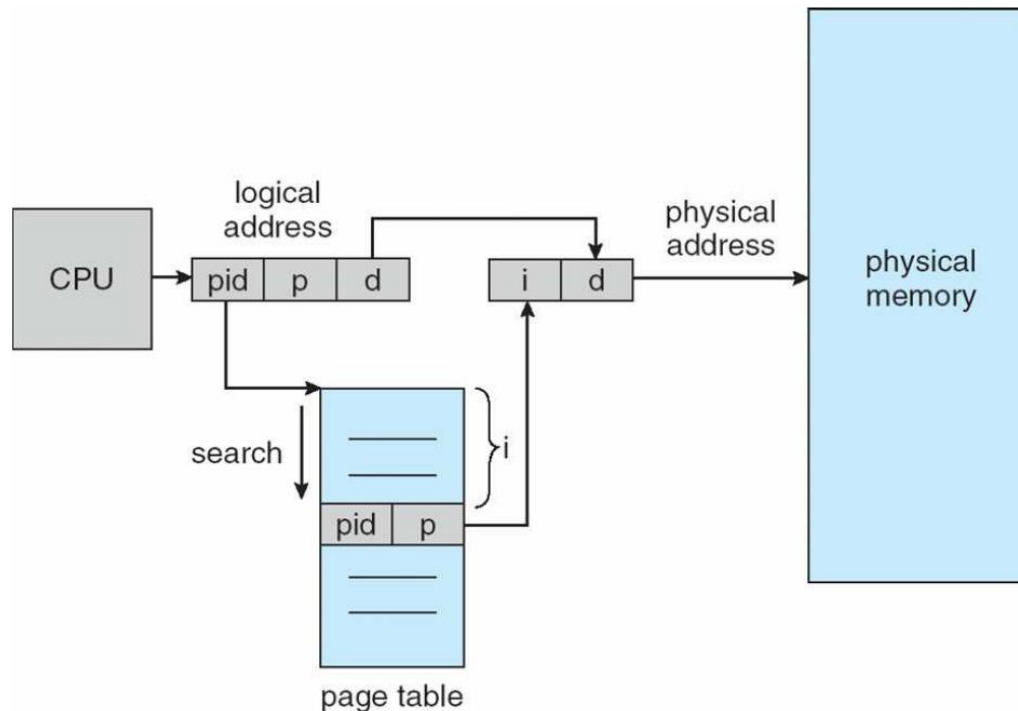
Inverted Page Table

- For large m a good solution is an **inverted page table**
 - the inverted page table has **one entry per frame of physical memory** (so there is only one)
 - each entry includes a PID and page number



Inverted Page Table

- Issue: sequential search (linear inverted page table)
 - for this reason it is usually **hashed to speed up access**, using linked lists for collisions
- If no match found: page fault



Next week

- Remember class time change
- Study Time
 - Review Chapter 9