

# Performance of Computer System

## Workload Selection and Characteristics

Dr. Lina Xu

`lina.xu@ucd.ie`

School of Computer Science,  
University College Dublin

October 11, 2019

- Different types of workloads
- What workloads are commonly used
- How to select appropriate workload types
  - ▶ Right level of details
  - ▶ Representativeness
  - ▶ Timeliness

# Workload Types

- Real workload
  - ▶ One observed during normal system operations
  - ▶ Non-repeatable
- Synthetic workload
  - ▶ Approximation of real workload
  - ▶ Can be applied repeatedly in a controlled manner
  - ▶ No large data files, no sensitive data
  - ▶ Easily modified and ported
  - ▶ Easily measured
- Test workload
  - ▶ Any workload used in performance studies
  - ▶ Real or synthetic

# Workload Selection

Workload selection is important

- “Would you tell me, please, which way I ought to go from here?”
- “That depends a good deal on where you want to get to,” said the Cat.
- “I don’t much care where—” said Alice.
- “Then it doesn’t matter which way you go,” said the Cat.
- From Alice’s Adventures in Wonderland.

# Services Exercised

- System Under Test (SUT)
- Component Under Study (CUS)
- Metrics chosen should reflect system level performance
- Workload chosen should reflect SUT, not CUS

# Level of Detail

## Most frequent request

- valid if one service is requested much more often than others
- examples: add instruction, kernels

## Frequency of request types

- example: instruction mix
- context sensitive services - must use a set (e.g. caching)

## Time stamped sequence of requests (trace)

- too much detail for analytical modeling
- may require exact reproduction of component behaviour for timing

# Level of Detail

## Average resource demand

- analytical models use request rate rather than requests
- group similar services in classes; use avg. demand per class

## Distribution of resource demands, used if

- variance in resource demands is large
- distribution impacts performance

# Representativeness

## Arrival rate

- Should be the same or proportional to that of real application

## Resource demands

- Total demand should be = or proportional to that of real application

## Resource usage profile

- Amount and sequence in which resources are consumed
- Especially important when forming a composite workload



# Timeliness

Users change usage pattern based on

- New services available
- Changes in system performance: users optimise demand

Anticipate changes

- Monitor user behaviour on ongoing basis
- Future may be different than past or present

# Other Considerations

## Load level

- full capacity (best case)
- beyond capacity (worst case)
- real workload (average case)
- for procurement, consider typical case
- for design, consider best  $< \text{---} >$  worst cases

## Impact of external components

- don't use workload that makes external component bottleneck
- e.g. if studying CPU performance, don't use data so large that system is paging
- otherwise, all alternatives will give equally good performance

## Repeatability

- want to be able to repeat results without excessive variance
- highly random resource demands should be avoided

# Workload Characteristics

## Motivation

- Motivation
  - ▶ Since a real-user environment is generally not repeatable, it is necessary to study the real-user environments, observe the key characteristics, and develop a workload model that can be used repeatedly. This process is called workload characterization.
- Different approaches for characterising workload

# Workload Characterisation

## Why

- Observe key performance characteristics of a workload
- Develop a model that can be used for further study

## Workload parameters


- Measured quantities that depend on workload not system
- Types
  - ▶ Service request
  - ▶ Resource demands
- Examples
  - ▶ Transaction types
  - ▶ Instructions,
  - ▶ Packet types destinations

# Techniques for Workload Characterisation

- Averaging
- Specifying Dispersion
- Single-parameter histograms
- Multiparameter histograms
- Principal components analysis
- Markov models
- Clustering

# Averaging

- Arithmetic mean of values  $x_1, x_2, \dots, x_n$ :  $\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i$
- Caution: arithmetic mean is not always appropriate “index of central tendency”
- Median = 50th percentile value
- Mode = most frequent



# Specifying Dispersion


- Averaging is insufficient if there is large variability in data values
- Variability of  $\{x_1, x_2, \dots, x_n\}$  is commonly specified by variance
  - ▶  $s^2 = \frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})^2$
- Sample standard deviation,  $s = \text{sqrt}(\text{sample variance})$ 
  - ▶ often more meaningful: same units as the mean
- Alternatives for summarising variability
  - ▶ range: maximum - minimum
  - ▶ 10 and 90 percentiles
  - ▶ semi-interquartile range (SIQR)
  - ▶ mean absolute deviation
- The ratio of the standard deviation to the mean is called the Coefficient Of Variation (C.O.V.).  $= s/\text{mean}$

# Specifying Dispersion

$$\text{range} = \text{max} - \text{min} \quad \text{IQR} = Q_3 - Q_1$$

16, 24, 26, 26, 26, 27, 28    23, 25, 28, 28, 32, 33, 35

$$\text{range} = 28 - 16 \quad \text{IQR} = 33 - 25$$





# Specifying Dispersion

- The resource demands of various programs executed on six university sites were measured for 6 months.
- The average demand by each program is shown.
- Notice that the C.O.V. of the measured values are rather high, indicating that combining all programs into one class is not a good idea.
- Programs should be divided into several classes.
- For all editors in the same data. The C.O.V. are now much lower.

TABLE 6.1 Workload Characterization Using Average Values

| Data                    | Average      | Coefficient of Variation |
|-------------------------|--------------|--------------------------|
| CPU time (VAX-11/780)   | 2.19 seconds | 40.23                    |
| Number of direct writes | 8.20         | 53.59                    |
| Direct-write bytes      | 10.21 kbytes | 82.41                    |
| Number of direct reads  | 22.64        | 25.65                    |
| Direct-read bytes       | 49.70 kbytes | 21.01                    |

TABLE 6.2 Characteristics of an Average Editing Session

| Data                    | Average      | Coefficient of Variation |
|-------------------------|--------------|--------------------------|
| CPU time (VAX-11/780)   | 2.57 seconds | 3.54                     |
| Number of direct writes | 19.74        | 4.33                     |
| Direct-write bytes      | 13.46 kbytes | 3.87                     |
| Number of direct reads  | 37.77        | 3.73                     |
| Direct-read bytes       | 36.93 kbytes | 3.16                     |


# Single-parameter Histograms

SPH: relative frequencies of various values of a parameter



- divide complete range into buckets
- count observations that fall in each

Uses

- simulation: generate test workload matching distribution
- analytical model: validate probability distribution used in model



# Single-parameter Histograms Examples



# Single-parameter Histograms Examples




# Single-parameter Histograms: Disadvantages

- much data:  $n$  buckets,  $m$  parameters/ component,  $k$  components
  - ▶ should only be used if variance is high and averages are inappropriate
- SPH ignore correlation among parameters


# Multiparameter histograms

MPH: use when significant correlation between parameters



# Multiparameter histograms

MPH: use when significant correlation between parameters



# Multiparameter histograms: Disadvantages


- rendering more than 2-parameter histograms is problematic
- much detail; uncommon to use them



# Weighted Sum of Parameters

- Classify workload components by weighted sum of their parameter values
  - ▶  $y = \sum_{j=1}^n a_j x_j$
- Use  $y$  to classify components into categories, e.g. high, low
- Problem: choosing appropriate weights for parameters.
  - ▶ Bad choice of weights may group dissimilar components

# The average consumption of 17 types of food in grams per person per week for every country in the UK



# Principal components analysis

## Problem

- Find weights so that weighted sums provide maximum discrimination among components
- $y_i = \sum_{j=1}^n a_{ij}x_j$

For each component  $i$ ,

- $y_i$  is a linear combination of parameter values  $x_j$
- $a_{ij}$  is loading of  $x_j$  on  $y_i$



Choose weights so that  $y$ 's form an orthogonal set, namely

- $\langle y_i, y_j \rangle = \sum_1^k a_{ik}a_{kj} = 0$

Properties:  $y$ 's form an ordered set such that

- $y_1$  explains highest percentage of variance in resource demands
- Successive  $y_i$  explain increasingly lower percentages

# Principal components analysis

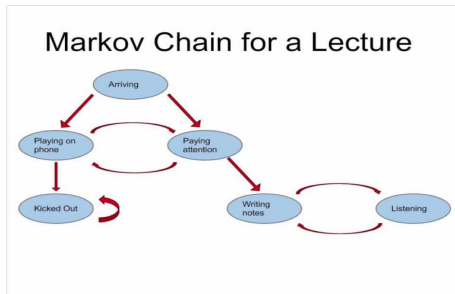


# Markov models

## One of the Queueing Models

- This section is about being able to describe the behavior of queues. Queues are certainly a prevalent object in computer systems, and our goal here is to write the equations that describe them. The language we'll use here is mathematical, but nothing really more complicated than algebra.
- Sometimes, not only the relative frequency but order of service requests is important!

# Markov models



# First Order Markov Models

- Probabilistic models that can generate sequences of states (here, representing service requests)
- Next state depends only upon current state
- Completely characterised by a transition probability matrix
- Transition probability matrix properties
  - ▶  $a_{ij} = P(\text{system will enter state } j \text{ — system is in state } i)$
  - ▶  $0 \leq a_{ij} \leq 1$
  - ▶  $\sum a_{ij} = 1$

# First Order Markov Models Example I

Modeling packets on network

- small packets = 87.5%; large packets = 12.5% (1 in 8 is large)
- large packet always followed by a small packet

**Pairwise percentages**

|                | next packet |       |
|----------------|-------------|-------|
|                | small       | large |
| current packet |             |       |
| small          | 75%         | 12.5% |
| large          | 12.5%       | 0%    |




# First Order Markov Models Example I

Pairwise percentages

| current packet | next packet |       |
|----------------|-------------|-------|
|                | small       | large |
| small          | 75%         | 12.5% |
| large          | 12.5%       | 0%    |

Transition Matrix

| current packet | next packet        |                      |
|----------------|--------------------|----------------------|
|                | small              | large                |
| small          | $75/87.5 = .85714$ | $12.5/87.5 = .14286$ |
| large          | 1                  | 0                    |



# First Order Markov Models Example II

Modeling packets on network

- small packets = 87.5%; large packets = 12.5% (1 in 8 is large)
- large packet followed by a small packet half the time


# First Order Markov Models Example II

Pairwise percentages

| current packet | next packet |       |
|----------------|-------------|-------|
|                | small       | large |
| small          | 81.25%      | 6.25% |
| large          | 6.25%       | 6.25% |

Transition Matrix

| current packet | next packet                  |                             |
|----------------|------------------------------|-----------------------------|
|                | small                        | large                       |
| small          | $81.25/87.5$<br>=<br>.928571 | $6.25/87.5$<br>=<br>.071429 |
| large          | $6.25/12.5$<br>=<br>.5       | $6.25/12.5$<br>=<br>.5      |




# First Order Markov Models Properties

- Finite: The model consists of a finite number of states
- Memory-less: The next state depends only upon the present state, not past states
- Absorbing state: (enter, but never leave) any state with a 1 on the main diagonal in the transition probability matrix
- Time independent: Transition matrix probabilities do not vary over time


# Clustering

- Workload often consists of large number of components
- Want to classify components into small number of clusters whose members are similar



# How to Cluster

- Take a sample, that is, a subset of workload components.
- Select workload parameters
- Transform parameters, if necessary.
- Remove outliers.
- Scale observations (e.g. normalise to 0 mean, unit variance)
- Select distance metric, e.g. Euclidean, Manhattan distance
- Select clustering algorithm
- Perform clustering



# How to Cluster

Select workload parameters

- Select those that (1) impact performance, and (2) vary significantly

Transform parameters, if necessary.

- If the distribution of a parameter is highly skewed, one should consider the possibility of replacing the parameter by a transformation or function of the parameter.

# How to Cluster

## Remove outliers

- Outliers affect normalisation and thus can affect clustering
- Remove if they do not consume significant fraction of system resources

## Scale observations (e.g. normalise to 0 mean, unit variance)

- Normalize to Zero Mean and Unit Variance:

- ▶  $x'_{ik} = \frac{x_{ik} - \bar{x}_k}{s_k}$

- Weights

- ▶  $x'_{ik} = W_k x_{ik}$



# How to Cluster

Scale observations (e.g. normalise to 0 mean, unit variance)

- Range Normalisation: The range is changed from  $[x_{min}, k, x_{max}, k]$  to  $[0, 1]$ .
- $x'_{ik} = \frac{x_{ik} - x_{min,k}}{x_{max,k} - x_{min,k}}$

# How to Cluster

Select distance metric

- Euclidean
- Manhattan distance
- Weighted Euclidean Distance

# How to Cluster

Select clustering algorithm

- Nearest neighbour: min distance between objects in two clusters
- Furthest neighbour: max distance between objects in two clusters

# How to Cluster

## Perform clustering

- Start with  $n$  clusters, using minimum spanning tree to merge
- Represent with dendrogram

# Minimum spanning tree


## Minimum Spanning Tree Problem

- Undirected graph  $G$  with vertices for each of  $n$  objects
- Weights  $d(u, v)$  on the edges giving the distance  $u$  and  $v$ ,
- Find the subgraph  $T$  that connects all vertices and minimises

$T$  will be a tree!

# Minimum spanning tree – clustering

- Agglomerative approach
- Divide the  $n$  items up into  $k$  groups so that the minimum distance between items in different groups is maximized



# Minimum spanning tree – clustering

- 1 Start with  $k = n$  clusters.
- 2 Find the centroid of the  $i_{th}$  cluster,  $i = 1, 2, \dots, k$ . The centroid has parameter values equal to the average of all points in the cluster.
- 3 Compute the inter-cluster distance matrix. Its  $(i, j)_{th}$  element is the distance between the centroids of cluster  $i$  and  $j$ .
- 4 Find the smallest nonzero element of the distance matrix. Let the distance between clusters  $l$  and  $m$ , be the smallest. Merge clusters  $l$  and  $m$ . Also merge any other cluster pairs that have the same distance.
- 5 Repeat steps 2 to 4 until all components are part of one cluster.

# Clustering Example

